

AAiT logo
placeholder
(see README)

ADDIS ABABA UNIVERSITY
ADDIS ABABA INSTITUTE OF TECHNOLOGY
SCHOOL OF INFORMATION TECHNOLOGY AND ENGINEERING
Department of [Name]

Video Saliency Prediction with a Frozen Self-Supervised Video Encoder

A thesis submitted to the School of Information Technology and Engineering, Addis Ababa Institute of Technology, Addis Ababa University, in partial fulfillment of the requirements for the degree of Master of Science in [Program Name]

By
Full Name of the Student

Advisor: Dr. Full Name of the Advisor
Co-advisor: Dr. Full Name of the Co-advisor

Addis Ababa, Ethiopia
November 2026

ADDIS ABABA UNIVERSITY

ADDIS ABABA INSTITUTE OF TECHNOLOGY

SCHOOL OF INFORMATION TECHNOLOGY AND ENGINEERING

Video Saliency Prediction with a Frozen Self-Supervised Video Encoder

By

Full Name of the Student

Approval by Board of Examiners

This is to certify that the thesis prepared by Full Name of the Student, entitled "Video Saliency Prediction with a Frozen Self-Supervised Video Encoder", and submitted in partial fulfillment of the requirements for the degree of Master of Science in [Program Name], complies with the regulations of the University and meets the accepted standards with respect to originality and quality.

	Name	**Signature**	**Date**
Advisor	Dr. Full Name of the Advisor	_____	_____
Co-advisor	Dr. Full Name of the Co-advisor	_____	_____
Internal Examiner	_____	_____	_____
External Examiner	_____	_____	_____
Chairperson	_____	_____	_____
School Dean	_____	_____	_____

DECLARATION

I, the undersigned, declare that this thesis is my original work, that it has not been presented for a degree in this or any other university, and that all sources of material used for the thesis have been duly acknowledged.

	Name	**Signature**	**Date**
Student	Full Name of the Student	_____	_____

This thesis has been submitted for examination with my approval as university advisor.

	Name	**Signature**	**Date**
Advisor	Dr. Full Name of the Advisor	_____	_____
Co-advisor	Dr. Full Name of the Co-advisor	_____	_____

Place: Addis Ababa, Ethiopia
Date of submission: November 2026

To ...

ABSTRACT

Predicting where people look in videos supports applications from compression to content design. Recent video saliency models draw their accuracy from large pretrained encoders, which they fine-tune on small eye-tracking datasets and run once for every output frame. This thesis asks whether a self-supervised video encoder trained by feature prediction, V-JEPA 2, suffices without adaptation, and whether the error of its pretrained predictor, a learned measure of surprise, indicates where people look.

The proposed model, SURGE-D, combines the frozen encoder with a decoder of 7.4 million parameters that builds a multi-scale pyramid from four intermediate layers, decodes it with factorised spatio-temporal convolutions and predicts saliency maps for sixteen frames per forward pass. Design choices were decided by rules written in advance, and each model was scored once per test set by a validated evaluator. The evaluation covers DHF1K, UCF Sports, Hollywood-2 with openly licensed gaze recordings, and six audio-visual datasets.

On the DHF1K validation set, a three-seed ensemble with the ViT-g encoder reached a correlation coefficient of 0.603, against 0.564 for SalFoM, the closest foundation-model approach, scored by the same evaluator. It was better on all six metrics, with bootstrap intervals excluding zero, and remained better when SalFoM's maps were given their best blur. A causal variant that, like SalFoM, sees no future frames was as accurate as the offline model and beat SalFoM on the same frames by 0.029 in correlation. With an encoder of SalFoM's size, only the ensemble was reliably better, offline and causally. Offline, the model needs about a ninth of SalFoM's encoder computation per output frame and less than a twentieth of its measured time; the causal variant needs nearly twice the computation but less time. After brief fine-tuning, it was competitive on UCF Sports and exceeded the best verified peer-reviewed results on four of six audio-visual datasets, by margins within the spread of published evaluations, but failed on conversations, where gaze follows the speaker. The predictor's error was anti-correlated with fixations; combined with simple cues, it added a little information about where people look, but with a negative weight. Encoder pre-training and scale produced the only large gains; over three seeds, a learned centre prior and deep supervision were equivalent to their absence within 0.003 in correlation, and no other addition tested, from encoder adaptation to audio, brought a gain that could be established. The thesis contributes a frozen-encoder saliency model, a negative result on predictive surprise, controlled evidence on design choices and a transparent evaluation protocol.

Keywords: video saliency prediction, eye fixations, self-supervised learning, joint-embedding predictive architecture, frozen encoder, predictive surprise

ACKNOWLEDGMENT

Write your acknowledgment here.

TABLE OF CONTENTS

Approval Page ii
Declaration iii
Abstract v
Acknowledgment vi
List of Figures x
List of Tables xi
List of Acronyms xvi
1. Introduction 1
1.1. Background 1
1.2. Motivation 1
1.3. Statement of the Problem 2
1.4. Research Questions 2
1.5. Objectives 3
1.5.1. General Objective 3
1.5.2. Specific Objectives 3
1.6. Research Methodology 4
1.6.1. Literature Review 4
1.6.2. Data Collection 4
1.6.3. Tools and Techniques 4
1.6.4. Evaluation 4
1.7. Scope and Limitations 4
1.8. Contributions 5
1.9. Organization of the Thesis 6
2. Literature Review 7
2.1. Visual Attention and Saliency 7
2.1.1. Eye Movements and Fixation Data 7
2.1.2. Bottom-Up and Top-Down Attention 8
2.1.3. Attention in Dynamic Scenes 9
2.1.4. Surprise and Prediction 10
2.2. Deep Saliency Prediction for Images 10
2.2.1. From Hand-Crafted to Learned Features 10
2.2.2. Saliency as a Probability Distribution 11
2.2.3. Fixation Prediction and Salient Object Detection 11

2.3. Learning Representations of Video . . . 11
2.3.1. Spatio-Temporal Convolutional Networks . . . 11
2.3.2. Video Transformers and Dense Prediction . . . 12
2.3.3. Self-Supervised Representation Learning . . . 12
2.3.4. Joint-Embedding Predictive Architectures . . . 13
2.3.5. Transferring and Adapting Pretrained Models . . . 14
2.4. Eye-Tracking Datasets for Video . . . 15
2.5. Evaluation Metrics . . . 16
2.6. Summary . . . 17
3. Related Work . . . 18
3.1. Recurrent and Two-Stream Models . . . 18
3.2. Three-Dimensional Convolutional Encoder–Decoders . . . 19
3.3. Transformer-Based Models . . . 19
3.4. Video Foundation Models for Saliency . . . 20
3.5. Audio-Visual Saliency Models . . . 21
3.6. Prediction Error as an Attention Signal . . . 21
3.7. Summary and Research Gap . . . 22
4. Methodology . . . 24
4.1. Research Design . . . 24
4.2. Problem Formulation . . . 25
4.3. Datasets . . . 27
4.3.1. Data Collection . . . 27
4.3.2. Preprocessing . . . 29
4.4. Proposed Architecture . . . 30
4.4.1. Frozen Video Encoder . . . 30
4.4.2. Three-Dimensional Simple Feature Pyramid . . . 31
4.4.3. Top-Down (2+1)D Decoder . . . 32
4.4.4. Temporal Upsampling, Centre Prior and Output . . . 32
4.5. Components Evaluated but Not Retained . . . 33
4.5.1. Predictive Surprise from the V-JEPA 2 Predictor . . . 33
4.5.2. Low-Rank Adaptation of the Encoder . . . 35
4.5.3. Audio Fusion . . . 36
4.5.4. Static-Image Co-Training . . . 36
4.6. Training Procedure . . . 36
4.7. Inference and Post-Processing . . . 38
4.8. Evaluation Metrics and Statistical Analysis . . . 39
4.8.1. Metrics . . . 39
4.8.2. Statistical Analysis . . . 40

4.9. Ethical Considerations . . . 40
5. Experiments, Results and Discussion . . . 42
5.1. Experimental Setup . . . 42
5.1.1. Implementation Details . . . 42
5.2. Validation of the Evaluator . . . 43
5.3. Selection of the Video Encoder (Gate G1) . . . 44
5.4. Predictive Surprise as a Fixation Cue (Gate G2) . . . 45
5.5. Results . . . 49
5.5.1. Quantitative Results on DHF1K . . . 49
5.5.2. Causal Prediction on DHF1K . . . 52
5.5.3. Generalisation to Other Datasets . . . 53
5.5.4. Qualitative Results . . . 55
5.5.5. Computational Cost . . . 57
5.6. Ablation Study . . . 58
5.7. Discussion . . . 63
5.8. Threats to Validity . . . 66
6. Conclusion and Future Work . . . 68
6.1. Conclusion . . . 68
6.2. Contributions . . . 69
6.3. Future Work . . . 70
References . . . 72
Appendix A: Additional Results . . . 85
A.1. Paired Bootstrap on the DHF1K Validation Set . . . 85
A.2. Checks Added After the First Draft . . . 85
A.3. All Metrics on the Audio-Visual Datasets . . . 86
A.4. Encoder Comparison on the Audio-Visual Datasets . . . 92
Appendix B: Reproducibility and Sample Code . . . 93
B.1. Organisation of the Code . . . 93
B.2. Corrections and Deviations . . . 93
B.3. Commands for the Main Results . . . 95
B.4. Forward Pass of the Decoder . . . 96

LIST OF FIGURES

Figure 4.1: Architecture of SURGE-D. A clip of 16 frames passes through the frozen V-JEPA 2 encoder; four intermediate layers (blocks 6/12/18/24 of ViT-L or 10/20/30/40 of ViT-g) feed a three-dimensional simple feature pyramid, which a top-down (2+1)D decoder turns into logits at 64×64 resolution for eight tubelet steps. After temporal interpolation to 16 frames and the addition of a learned centre prior, a softmax over pixels yields one saliency distribution per frame; the central eight frames are kept. Tensor sizes are given as time $\times$ height $\times$ width 30

Figure 5.1: Learning curves of gate G1. (a) Holdout CC after each epoch; circles mark the best epoch. (b) Mean CC on the training clips of each epoch. Holdout accuracy peaks early and then declines while training accuracy keeps rising, for all three encoders 45

Figure 5.2: Per-video CC difference between the SURGE-D ViT-g ensemble and SalFoM on the 100 DHF1K validation videos, sorted. The dashed line is the mean difference; the interval is the paired bootstrap 95 % interval 50

Figure 5.3: Qualitative examples on DHF1K validation. Rows: videos at the 10th, 50th and 90th percentile of SURGE-D's per-video CC, and the videos with the largest gain (0696) and loss (0663) relative to SalFoM; columns: middle frame with the observers' fixations, ground-truth density, SalFoM and the SURGE-D ViT-g ensemble, overlaid on the frame. Numbers give the CC of the shown frame. Video-level CC (SURGE-D / SalFoM): 0650 0.455/0.376, 0632 0.613/0.639, 0685 0.741/0.715, 0696 0.664/0.437, 0663 0.388/0.462 56

Figure 5.4: Holdout CC of the frozen ViT-L model against the number of training videos. Thin lines are the three seeds, whose subsets are nested, and the thick line with the printed values is their mean 61

Figure 5.5: Holdout CC as a function of the Gaussian blur applied to the output, for single models and three-seed ensembles. The selected values are 12 (ViT-L single) and 8 (others) 63

LIST OF TABLES

Table 2.1: Video eye-tracking datasets relevant to this thesis 15

Table 3.1: Summary of the closest related works 22

Table 4.1: Decision rules, each written before the experiment that it governs, except the one marked * (see text). Rules marked † were written in a follow-up after the first draft of this thesis and rules marked ‡ in a second follow-up, in each case before the runs they govern. Holdout CC is measured on the internal 60-video DHF1K holdout unless stated otherwise; outcomes are reported in Chapter 5 26

Table 4.2: Eye-tracking datasets used in this thesis. "Holdout" is the internal selection set drawn from the training part; the test part is scored once per model 28

Table 4.3: Training and inference settings of SURGE-D 38

Table 5.1: Experimental environment 42

Table 5.2: Gate G1: frozen ViT-L encoders with the same decoder, internal DHF1K holdout (60 videos, every fifth frame, no blur), best of ten epochs, seed 0 44

Table 5.3: Gate G2: zero-shot fixation prediction by prediction-error maps and reference maps on the DHF1K holdout (60 videos), at about 4 and 15 sampled frames per second. The masked-tube error is not part of the gate. Intervals are given in Tables A.9 and A.10 46

Table 5.4: Incremental information U of the prediction-error maps in bits per fixation, with 95 % bootstrap intervals over the 60 holdout videos, and the fitted weight of the error map in the two halves of the cross-fitting. An interval above zero is read as incremental information (Table 4.1); the two primary comparisons are marked ⋆. The last block gives, for scale, the loss when one simple cue is removed from the model of the first row, and the gain of the two simple cues together over the centre prior 47

Table 5.5: Results on the DHF1K validation set (100 videos, 60,273 frames). The upper block was scored by the evaluator of this thesis; the lower block is copied from the literature and was computed with other evaluation code. The second SalFoM row applies the blur that is best for SalFoM on this set, an upper bound for the baseline (Section 5.5.1). Best value per column in bold 49
Table 5.6: Paired bootstrap differences on DHF1K validation (100 videos, 10,000 resamples), with 95 % intervals; differences whose interval excludes zero are in bold. The intervals for all six metrics are given in Table A.1 50
Table 5.7: Checks of the comparison with SalFoM on DHF1K validation: CC and paired bootstrap differences (100 videos, 95 % intervals). "Best blur": SalFoM's maps blurred with $\sigma = 16$, the value of the grid that is best for SalFoM on this set, an upper bound for the baseline. Differences whose interval excludes zero are in bold 51
Table 5.8: Causal prediction on DHF1K validation, scored on every tenth frame (100 videos, 6,076 frames). The causal models and SalFoM use only the predicted frame and earlier frames; the offline models, for reference, also use later frames. Paired bootstrap differences in CC with 95 % intervals (10,000 resamples); differences whose interval excludes zero are in bold. All six metrics are compared in Tables A.5 and A.6 . . 53
Table 5.9: Results on the UCF Sports test set (47 videos, 2,929 frames). Published results were computed with other evaluation code 54
Table 5.10: CC on the six audio-visual datasets (fold means; DIEM has one split). The lower block lists the best peer-reviewed results whose tables could be verified, computed with other evaluation code. Visual-only: ViNet-E, or ViNet-A on Coutrot2 [45]; on SumMe, which those models do not report, the visual-only variant of the relevance-guided fusion model AVRSP of Yu et al. [116]. Audio-visual: AVRSP [116] on DIEM and Coutrot2, DiffSal [91] on the other four 55
Table 5.11: Results on Hollywood-2 with the gaze of Vig et al. (884 test clips). Not comparable with published Hollywood-2 results, which use other ground truth 56
Table 5.12: Parameters and operations of SURGE-D. TFLOPs count multiply–add pairs as two operations 57
Table 5.13: Measured inference time on one RTX A5000, in milliseconds per forward pass and per output frame (median of 50 passes; third of three runs). The causal rows time the same forward pass as the offline rows but keep one map per pass 58

Table 5.14: Ablations on the DHF1K holdout (60 videos) unless marked "val". ΔCC is relative to the reference in each block; for the components, it is the mean over three seeds of the difference to the reference of the same seed, with its 95 % interval from a bootstrap over seeds and videos 59
Table 5.15: Learning curve over the number of training videos (frozen ViT-L, DHF1K holdout, seeds 0 / 1 / 2). The last four columns of the upper part are means over the seeds. The differences in the lower part are paired by seed and video, their intervals come from the bootstrap over seeds and videos, and the outcomes follow the three-seed rule of Table 4.1 with its margin of 0.003 62
Table 5.16: Audio fusion on the audio-visual datasets (ViT-L, fold-mean CC). The rule requires audio − visual ≥ +0.010 and matched − rotated ≥ +0.005 63

Table A.1: Paired bootstrap differences between SURGE-D and SalFoM on DHF1K validation, all six metrics 85
Table A.2: Paired bootstrap differences between the ViT-g and ViT-L encoders on DHF1K validation, all six metrics 86
Table A.3: Paired bootstrap differences between SURGE-D and SalFoM with its best blur ($\sigma = 16$) on DHF1K validation, all six metrics (100 videos, 10,000 resamples) 86
Table A.4: Every use of the DHF1K validation set, the test set of this thesis. No reported model or setting was selected on it 87
Table A.5: Paired bootstrap differences between causal SURGE-D and SalFoM with its best blur ($\sigma = 16$) on DHF1K validation, every tenth frame, all six metrics (100 videos, 10,000 resamples) 88
Table A.6: Paired bootstrap differences between causal and offline SURGE-D on the same DHF1K validation frames (every tenth frame), all six metrics (100 videos, 10,000 resamples). For KLD, negative differences favour the causal model 88
Table A.7: Training on all 600 DHF1K training videos instead of 540 (ViT-g): paired bootstrap differences on DHF1K validation, all six metrics (100 videos, 10,000 resamples) 89
Table A.8: Three-seed ensemble minus its seed-0 member on DHF1K validation: paired bootstrap differences, all six metrics (100 videos, 10,000 resamples). For KLD, negative differences favour the ensemble . . . 89

Table A.9: Gate G2 with the untrained mask token, as in its first run, with 95 % bootstrap intervals over the 60 holdout videos, computed from the stored per-video scores. The masked-tube error was added in the second follow-up. The reference maps do not depend on the mask token. The lower block gives paired differences 90
Table A.10: Gate G2 re-scored with the trained mask token, with 95 % bootstrap intervals over the 60 holdout videos. The reference maps are those of Table A.9. The lower block gives paired differences 90
Table A.11: Information gain of each cue on its own over the fitted centre prior, in bits per fixation, cross-fitted as in Equation (4.11), with 95 % bootstrap intervals over the 60 holdout videos (descriptive). A positive gain can come from a negative weight. Maps without a token named are computed with the untrained token 91
Table A.12: SURGE-D ViT-g three-seed ensemble on the audio-visual datasets (fold means; DIEM has one split) 91
Table A.13: CC difference between the ViT-g and ViT-L three-seed ensembles on the audio-visual datasets (pooled test videos, 10,000 resamples) . . . 92

Table B.1: Main files of the implementation 94

LIST OF ACRONYMS

AAiT	Addis Ababa Institute of Technology
AAU	Addis Ababa University
AST	Audio Spectrogram Transformer
AUC	Area Under the ROC Curve
AUC-J	Area Under the ROC Curve, Judd variant
AV	Audio-Visual
CC	Linear Correlation Coefficient
CI	Confidence Interval
CNN	Convolutional Neural Network
ConvLSTM	Convolutional Long Short-Term Memory
CPU	Central Processing Unit
CUDA	Compute Unified Device Architecture
FiLM	Feature-wise Linear Modulation
FLOPs	Floating-Point Operations
FPN	Feature Pyramid Network
GELU	Gaussian Error Linear Unit
GN	Group Normalisation
GPU	Graphics Processing Unit
JEPA	Joint-Embedding Predictive Architecture
KL	Kullback–Leibler
KLD	Kullback–Leibler Divergence
LoRA	Low-Rank Adaptation
LSTM	Long Short-Term Memory
MAE	Masked Autoencoder
mIoU	Mean Intersection over Union
NSS	Normalised Scanpath Saliency
RAM	Random-Access Memory
RNN	Recurrent Neural Network
ROC	Receiver Operating Characteristic
sAUC	Shuffled Area Under the ROC Curve
SIM	Similarity (histogram intersection)
SiTE	School of Information Technology and Engineering
TFLOPs	Tera Floating-Point Operations (10^{12})
UMT	Unmasked Teacher

V-JEPA	Video Joint-Embedding Predictive Architecture
ViT	Vision Transformer
VSP	Video Saliency Prediction

CHAPTER ONE

INTRODUCTION

1.1. Background

When people watch a video, they do not look everywhere. Their eyes move several times per second, and different observers tend to look at the same few regions: moving objects, faces and, in conversations, the person who is speaking [1], [2]. Where people look is shaped by the image itself, by the viewer's task and by strong regularities such as the preference for the centre of the screen [3], [4], [5]. *Video saliency prediction* is the task of predicting, for every frame of a video, the distribution of human fixations. Models of this kind support applications in which the allocation of attention matters, from compression and summarisation to the design and evaluation of visual content [6], and they are also scientific models of human attention whose successes and failures say something about what drives gaze.

The field has progressed in step with the representations it builds on. Early models computed saliency from the local contrast of hand-crafted features such as intensity, colour, orientation and motion [7]. The introduction of large eye-tracking benchmarks, most importantly DHF1K [8], [9], enabled deep networks to be trained for the task, and the accuracy of these networks has risen with each generation of their encoders: image networks with recurrent layers, three-dimensional convolutional networks pretrained for action recognition, video transformers and, most recently, video foundation models trained with self-supervision on large collections of video [10].

Self-supervised video models differ in what they learn to predict. One family, the joint-embedding predictive architectures, learns by predicting the *representations* of hidden parts of a video from the visible parts [11], [12]. Such models have two properties that make them interesting for saliency prediction. Their features are reported to work well without adapting the encoder to a new task [12], which suggests that a frozen encoder with a small task-specific decoder might suffice. And they contain a trained predictor, whose errors indicate what the model failed to anticipate, a learned analogue of the "surprise" that has been proposed as a driver of human attention [13].

1.2. Motivation

Three considerations motivated this work. The first is practical. SalFoM [10], the first video saliency model built on a video foundation encoder, which reported the best DHF1K test results among the methods it compared with, fine-tunes a 300-million-

parameter encoder on a few hundred training videos and runs it once for every output frame. Fine-tuning so large a model on so little data risks overfitting, which its authors offer as an explanation for its weaker results on a small dataset, and the per-frame design makes inference expensive. A frozen encoder with a decoder that predicts many frames at once would avoid both problems, if its features are good enough.

The second consideration is scientific. The idea that attention is drawn to the unexpected has strong support from computational models built on hand-crafted features [13], but it had not been tested with the learned predictor of a modern video model. If such a predictor's errors marked where people look, they would provide a principled attention cue learned without any gaze data; if they did not, understanding why would clarify what these models actually learn.

The third consideration is methodological. Published comparisons in video saliency rely mostly on numbers copied from other papers, computed with different evaluation code and sometimes on different data splits, and they rarely report uncertainty. Small reported gains are therefore hard to interpret. A study that re-scores its main baseline with the same validated evaluator, quantifies uncertainty and fixes its decision rules in advance can give answers that are more reliable, including negative answers.

1.3. Statement of the Problem

Recent video saliency models obtain their accuracy from large pretrained encoders, which they adapt to small saliency datasets, at a high computational cost per output frame. It is not known whether this adaptation is necessary, or whether the representations of a self-supervised video encoder trained by feature prediction are sufficient on their own, with only a light decoder trained on gaze data. Nor is it known whether the prediction errors of such an encoder's predictor, which offer a learned measure of surprise, carry information about where people look. Finally, because comparisons in the field are rarely made under a common evaluation protocol with uncertainty estimates, it is unclear which design choices, beyond the encoder, still matter once a strong encoder is used. This thesis addresses these questions with a frozen V-JEPA 2 encoder [14], a purpose-built decoder, and an evaluation on nine eye-tracking test sets whose decision rules were fixed in advance in a dated log.

1.4. Research Questions

This study answers the following research questions:

RQ1. Can a frozen self-supervised video encoder of the V-JEPA family, combined with a lightweight spatio-temporal decoder, predict human fixations in video more accurately than the closest foundation-model approach, SalFoM, when both are

evaluated under an identical protocol?

RQ2. Does the prediction error of the encoder's pretrained latent predictor ("predictive surprise") carry information about where people look, beyond that carried by simple appearance and motion cues?

RQ3. Which design factors determine accuracy when the encoder is frozen: the pretraining and scale of the encoder, parameter-efficient adaptation, the learned centre prior, deep supervision, static-image co-training, temporal sampling, the amount of training data and ensembling?

RQ4. How well does the approach transfer to eye-tracking datasets with other content and viewing tasks, and does adding the soundtrack improve it on audio-visual datasets?

1.5. Objectives

1.5.1. General Objective

The general objective of this study is to design, implement and rigorously evaluate a video saliency prediction model built on a frozen self-supervised video encoder, and to test whether the encoder's own prediction error is a useful cue for human attention.

1.5.2. Specific Objectives

To achieve the general objective, the study will:

- review the literature on visual attention, video saliency prediction and self-supervised video representation learning;
- prepare the eye-tracking datasets and an evaluation protocol whose metric implementation is validated against reference code, including a new ground truth for Hollywood-2 from openly licensed gaze recordings;
- design and implement a decoder that turns the intermediate features of a frozen V-JEPA 2 encoder into saliency maps for many frames per forward pass;
- test the predictive surprise of the V-JEPA 2 predictor as a fixation cue against simple reference cues;
- compare encoders and design choices under decision rules fixed before each experiment;
- evaluate the final models on DHF1K and eight further test sets, with paired statistical comparisons where the data allow; and
- measure the computational cost of the approach.

1.6. Research Methodology

1.6.1. Literature Review

The literature was searched in publisher indexes and through the reference lists of key papers, covering visual attention, image and video saliency prediction, video representation learning and evaluation methodology. Only peer-reviewed publications are cited. The venue and bibliographic details of every reference were verified against the publishers' records or, for proceedings without digital object identifiers, against the indexes of those proceedings; works available only as preprints were excluded.

1.6.2. Data Collection

No new gaze data were recorded. The study uses public eye-tracking datasets: DHF1K, UCF Sports, Hollywood-2 with the gaze recordings of Vig et al., the six audio-visual datasets of the STAViS benchmark, and SALICON for one ablation (Section 4.3).

1.6.3. Tools and Techniques

The model was implemented in Python with PyTorch [15]; the pretrained encoders were loaded with the Transformers library [16]. All experiments ran on a single workstation with one NVIDIA RTX A5000 graphics card (Section 5.1).

1.6.4. Evaluation

Models were evaluated with six standard saliency metrics [17], computed by an evaluator validated against reference code. Differences between models were tested with a paired bootstrap over videos, and every component that was tested was governed by a rule written before the corresponding experiment, with one stated exception (Section 4.1).

1.7. Scope and Limitations

Scope. The study covers offline (non-causal) saliency prediction for video, in which the map of a frame may depend on the frames around it; a causal variant, which sees only past frames, is evaluated on DHF1K as a control. The main model is visual; the soundtrack is examined in one separate experiment. The encoders are frozen V-JEPA 2 models in two sizes; V-JEPA 2.1 and UMT-L are included in the encoder comparison. Evaluation covers nine test sets: DHF1K, UCF Sports, Hollywood-2 with the Vig gaze data, and six audio-visual datasets.

Limitations. The DHF1K test-set annotations are withheld by the dataset authors; predictions for the test set were prepared, but their scores were not available when the

thesis was written, so DHF1K results are reported on the official validation set. The standard Hollywood-2 ground truth could not be obtained, so the Hollywood-2 results use other gaze data and are not comparable with published numbers. Full fine-tuning of the encoder, other foundation encoders and efficient causal (streaming) prediction were outside the scope of the study; the causal variant was evaluated only on DHF1K, on every tenth frame, and the encoder comparison, the clip-length test and the audio experiment used a single training seed because of the available computing resources.

1.8. Contributions

The main contributions of this thesis are:

- **SURGE-D**, a video saliency model that combines a frozen V-JEPA 2 encoder with a light feature-pyramid decoder of factorised spatio-temporal convolutions and predicts maps for sixteen frames per forward pass. With 7.4 million trained parameters, its three-seed ViT-g ensemble reaches a correlation coefficient of 0.603 on the DHF1K validation set, against 0.564 for SalFoM scored by the same evaluator, and is better on all six metrics with paired bootstrap intervals that exclude zero. It remains better when SalFoM's maps receive the blur that suits them best and when SURGE-D is restricted, like SalFoM, to past frames. With the ViT-L encoder, the size of SalFoM's, the ensemble remains better under both conditions, but single models are not reliably better. In offline prediction, SalFoM's per-frame design needs about nine times the encoder computation per output frame of SURGE-D with ViT-g and, as released, more than twenty times its measured time.
- **A negative result on predictive surprise**: the error of the pretrained V-JEPA 2 predictor is anti-correlated with fixations, both when it forecasts the second half of a clip and when it fills in masked tubes as in pre-training. Combined with frame difference, token atypicality and the centre bias, the error adds a modest amount of information about fixations, but with a negative weight: observers tend to look away from regions that the predictor finds hard to predict.
- **Controlled evidence on what matters** once the encoder is frozen: the encoder's pre-training and scale produce the only large gains, seed ensembles a small one. Over three seeds, the learned centre prior and deep supervision are equivalent to their absence within 0.003 CC, while low-rank adaptation, static-image co-training, longer clips, more training videos and a simple audio fusion produce no gain that could be established.
- **A transparent evaluation** across nine test sets with a validated evaluator, decision rules written in advance in a dated log, paired statistics and a new,

openly reproducible Hollywood-2 ground truth from the gaze recordings of Vig et al., with the code, splits and decision log provided with the thesis.

1.9. Organization of the Thesis

The rest of this thesis is organized as follows. Chapter 2 reviews the background concepts and techniques. Chapter 3 discusses related work and identifies the research gap. Chapter 4 presents the proposed approach and the evaluation protocol. Chapter 5 describes the experiments and discusses the results. Chapter 6 concludes the thesis and outlines future work.

CHAPTER TWO

LITERATURE REVIEW

This chapter introduces the concepts on which the thesis builds. It begins with human visual attention: how eye movements are recorded and summarised as fixation data, what draws attention in static and dynamic scenes, and the notion of a saliency map (Section 2.1). It then follows the development of computational saliency models from hand-crafted features to deep networks that predict saliency as a probability distribution (Section 2.2), and reviews how video representations are learned, from spatio-temporal convolutions to the self-supervised and predictive video encoders that this thesis uses, together with what is known about transferring such representations to new tasks (Section 2.3). The last two sections describe the eye-tracking datasets and the evaluation metrics of the field (Sections 2.4 and 2.5). Specific video saliency models are reviewed in Chapter 3.

2.1. Visual Attention and Saliency

2.1.1. Eye Movements and Fixation Data

The human visual system cannot process the whole visual field at full resolution. High acuity is limited to the fovea, so the eyes move several times per second to bring regions of interest into central vision, and the sequence of fixations reveals what the observer attends to. An eye tracker samples the direction of gaze at a fixed rate, often several hundred times per second, and the resulting record alternates between *fixations*, periods during which gaze rests on one location and visual information is taken in, and *saccades*, the rapid movements between them. Raw gaze samples must therefore be segmented before they can be analysed. Salvucci and Goldberg [18] organised the algorithms used for this purpose by the criterion they apply, chiefly the velocity of the gaze or the spatial dispersion of consecutive samples, and compared how they differ in accuracy and robustness. The construction of the Hollywood-2 ground truth in this thesis (Section 4.3) relies on such a segmentation, provided with the recordings.

Saliency datasets summarise the fixations of many observers on each image or video frame in two forms [19]. A *fixation map* marks the fixated pixels, and a continuous *fixation density map* is obtained by convolving the fixation map with a Gaussian kernel whose width approximates the extent of foveal vision. The first form is used to evaluate a prediction at the fixated locations, the second to compare distributions, and both are used in this thesis.

Observers do not all look at the same places, so no model can be expected to predict the fixations of a new observer perfectly. Wilming et al. [20] formalised this by deriving a lower bound on achievable performance from image-independent regularities of fixation data, such as the preference for the centre, and an upper bound from the consistency between observers. In the same spirit, Kümmerer et al. [21] expressed the quality of a model as the fraction of the explainable information about fixations, between a centre-bias baseline and a gold standard derived from the observers themselves, that the model captures. Both views imply that differences between models must be judged relative to the room that the data leave for improvement.

2.1.2. Bottom-Up and Top-Down Attention

Two broad influences on the selection of fixations are usually distinguished: *bottom-up* factors, which depend on the properties of the image, such as local contrast in colour, orientation or motion; and *top-down* factors, which depend on the observer's goals, knowledge and task [3], [4].

The feature-integration theory of Treisman and Gelade [22] proposed that simple features are registered in parallel across the visual field, whereas attention is required to bind them into objects. Building on this view, Koch and Ullman [23] introduced the *saliency map*: a topographic map that combines the conspicuity of locations across feature dimensions, from which a winner-take-all mechanism selects the next location to attend, while inhibition of return prevents attention from returning immediately to the same place. Itti, Koch and Niebur [7] turned this idea into an influential computational model, in which centre–surround differences of intensity, colour and orientation at several scales are combined into a single saliency map. Their model, and the many that followed it, defined saliency in terms of local feature contrast [4]. Graph-based visual saliency, for example, computes activation maps from feature maps and normalises them through Markov chains [24], while the SUN model defines bottom-up saliency as the self-information of features relative to the statistics of natural images [25].

Top-down influences are harder to model because they depend on what the observer is trying to do. Judd et al. [26] showed that a model learned from eye-tracking data, combining low-level features with higher-level detectors for faces and people and with a centre prior, predicted fixations better than purely bottom-up models. This shift, from designing saliency features to learning them from human data, set the direction for the deep models that followed.

The observer's task changes where people look. Castelhano et al. [27] showed that the same scenes are explored differently under instructions to search and to memorise: the task changed which regions were fixated and how long gaze dwelt on particular objects, but not basic parameters such as the average saccade amplitude or the duration of individual fixations. Borji and Itti [28] revisited Yarbus's classic demonstration that

eye movements depend on the question put to the observer and showed that the task can be decoded from eye-movement features above chance. Eye-tracking data recorded under a task, such as the action-recognition task of two datasets used in this thesis (Section 2.4), therefore measure a different quantity from free-viewing data, and a model trained on one need not transfer to the other.

2.1.3. Attention in Dynamic Scenes

Video adds time, and with it motion, events and narrative. In dynamic scenes, gaze is strongly attracted by motion: in the DIEM recordings, the clustering of gaze across viewers was predicted by motion [1]. Carmi and Itti [29] distinguished features that cause attentional selection from features that merely correlate with it, using video clips with abrupt scene cuts, after which the influence of memory and expectation on gaze is briefly disrupted.

Gaze in video is also remarkably consistent across viewers. Dorr et al. [30] recorded 54 observers watching high-resolution videos of outdoor scenes. In some parts of the videos up to about 80 % of the observers looked at the same region, although the variability was usually much greater, and the authors used professionally edited Hollywood trailers to probe the upper limit of coherence that camera work and cuts can produce. Smith and Mital [31] called this clustering of gaze *attentional synchrony*. In their experiments, free viewing of dynamic scenes produced more synchrony, longer fixations and more gaze on people and on regions of high flicker than viewing static scenes, but a task that required identifying where a video was filmed reduced these differences. Gaze in video is thus strongly, but not exclusively, driven by the stimulus. This regularity is what makes video saliency prediction possible, and a task can override it.

Two further biases are prominent in video. The first is the *centre bias*. Observers tend to fixate the centre of the screen, and Tatler [32] showed that this bias cannot be explained by motor tendencies or by the distribution of image features alone. For dynamic scenes, Tseng et al. [5] found that the centre bias correlates strongly with *photographer bias*, the tendency to frame interesting content centrally, and is influenced by viewing strategy at scene onsets. The second is the strong attraction of *faces*, particularly of people who speak: in recordings of conversations, observers look mostly at faces, and especially at talking faces, and the original soundtrack makes them follow the turn-taking of the speakers more closely [2]. Sound thus influences gaze even when it carries no spatial information [33].

2.1.4. Surprise and Prediction

A distinct account holds that attention is drawn to events that violate expectations. Itti and Baldi [13] defined *Bayesian surprise* as the difference between an observer's posterior and prior beliefs after an observation. In their experiments, 72 % of the gaze shifts of observers watching television were directed towards locations more surprising than average, rising to 84 % for regions selected by all observers. Their model computed surprise from simple simulated early visual neurons, and earlier work had used it to detect surprising events in video [34]. The account connects to predictive-coding theories of the visual cortex, in which neural responses signal the error between the input and a prediction formed from context [35]. Whether the prediction errors of modern learned video models play a similar role is one of the questions examined in this thesis.

2.2. Deep Saliency Prediction for Images

2.2.1. From Hand-Crafted to Learned Features

The arrival of large eye-tracking and attention datasets enabled saliency models to be learned end to end. SALICON [36] collected mouse-contingent attention data, a scalable proxy for eye tracking, on images from MS COCO, and thereby provided a dataset large enough to train deep networks for saliency. Pan et al. [37] trained convolutional networks directly on such data, while Huang et al. [38] fine-tuned networks pretrained for object recognition and argued that their semantic features close much of the gap between low-level saliency and human fixations.

Two architectural themes of these models recur in video saliency. The first is the modelling of the centre bias. Fully convolutional networks respond in the same way at every position and therefore cannot represent location-dependent patterns; DeepFix [39] added location-biased convolutional layers for this purpose, and Cornia et al. [40], who refined the prediction iteratively with an attentive convolutional LSTM, learned a set of Gaussian prior maps, a design that recurs in later models, including the one proposed in this thesis. The second theme is global context: TranSalNet [41] integrated transformer components into a convolutional saliency model to capture long-range contextual information, which convolutions with local receptive fields encode poorly.

The DeepGaze line of work clarified *why* deep models succeed. Kümmerer et al. [42] separated the contributions of low- and high-level features to fixation prediction, and later showed that transfer learning from pretrained backbones had been the key driver of progress, with a twist: replacing the VGG19 backbone of DeepGaze II by ResNet50 features raised performance considerably, but better ImageNet classifiers, such as EfficientNet, brought no further gain [43]. The quality of the pretrained representation for saliency is therefore not simply its quality for classification. Borji [6] reviewed

the deep-learning era of saliency prediction for images and videos, and identified the remaining gap to human performance, the failure cases of current models and the difficulty of fair model comparison as open issues.

2.2.2. Saliency as a Probability Distribution

Early deep models were trained with losses borrowed from regression or classification, such as the per-pixel squared error [37]. Jetley et al. [44] argued that a saliency map should instead be treated as a probability distribution over image locations. They formulated the map as a generalised Bernoulli distribution, produced it with a softmax over pixels, and trained it with losses that measure distances between probability distributions, which outperformed standard losses on four benchmarks. The probabilistic view also clarifies evaluation: a model that outputs a distribution can be compared with others by the likelihood it assigns to the observed fixations [21]. Many recent video saliency models follow this view, training with the Kullback–Leibler divergence and the correlation coefficient between the predicted and the ground-truth distributions [10], [45], and the model of this thesis does the same (Section 4.6).

2.2.3. Fixation Prediction and Salient Object Detection

A related but distinct task is *salient object detection*, which segments the full extent of the objects that stand out in an image rather than predicting the distribution of fixations [46]. The outputs differ in kind: a fixation map is a smooth density concentrated on the few points that people actually look at, such as the eyes of a face, whereas a salient-object mask is a binary segmentation of whole objects. Fan et al. [47] connected the two in video. Their DAVSOD dataset of 226 videos uses human eye-fixation data to decide which objects are salient, and it emphasises the *saliency shift*, the change of the salient object over time. Some recent models address both tasks with one architecture [48]. This thesis is concerned with fixation prediction only.

2.3. Learning Representations of Video

2.3.1. Spatio-Temporal Convolutional Networks

Early deep video models extended image convolutions into time. C3D learned spatio-temporal features with three-dimensional convolutions on video clips [49]. Carreira and Zisserman inflated two-dimensional networks pretrained on ImageNet into three dimensions (I3D) and introduced Kinetics, a large dataset of human actions that became the standard source of video pre-training [50]. Xie et al. [51] showed that many of the costly three-dimensional convolutions can be replaced by separable spatial and temporal ones, producing the S3D network that many video saliency models of the following

years use as their encoder. Tran et al. [52] analysed such factorisations systematically; their R(2+1)D block, a spatial convolution followed by a temporal one, is the basis of the decoder blocks in this thesis.

Other designs separate appearance from motion explicitly. The two-stream networks of Simonyan and Zisserman [53] combine a spatial network that sees still frames with a temporal network that sees stacked optical flow, and showed that a network trained on dense optical flow performs well despite limited training data; two-stream designs were also among the first deep video saliency models (Section 3.1). SlowFast networks [54] use two pathways at different frame rates: a slow pathway that captures spatial semantics and a lightweight fast pathway that captures motion at fine temporal resolution. A SlowFast encoder trained for spatio-temporal action localisation underlies one of the strongest recent saliency models on action-oriented datasets [45].

2.3.2. Video Transformers and Dense Prediction

The Transformer [55] replaced recurrence and convolution by self-attention, which relates every element of a sequence to every other. The Vision Transformer (ViT) [56] applied it to images by splitting them into patches of 16×16 pixels that are embedded as tokens. For video, ViViT [57] and TimeSformer [58] extract spatio-temporal tokens and factorise attention over space and time to cope with the long token sequences, while Video Swin [59] restricts attention to local windows, reintroducing an inductive bias of locality.

A plain ViT produces features at a single resolution, whereas dense prediction tasks need several. Feature pyramid networks [60] build multi-scale representations from convolutional backbones with a top-down pathway and lateral connections. For transformers, DPT [61] assembles tokens from several stages into image-like maps at different resolutions and fuses them with a convolutional decoder, and Li et al. [62] showed that a simple pyramid built by up- and down-sampling the features of a plain ViT suffices for object detection. The decoder of this thesis follows these designs (Section 4.4).

2.3.3. Self-Supervised Representation Learning

Supervised pre-training ties a representation to a fixed set of labels. Self-supervised learning instead derives the training signal from the data itself, which allows much larger and more varied training sets. Several families of methods can be distinguished by the signal they use.

Contrastive methods learn by making the representations of two augmented views of the same image similar and those of different images dissimilar. SimCLR [63] identified the composition of data augmentations, a learnable non-linear projection between

the representation and the loss, and large batches as the keys to this approach, and MoCo [64] maintained a large and consistent set of negative examples with a queue of encoded samples and a slowly updated momentum encoder. *Non-contrastive* methods dispense with negative examples. In BYOL [65], an online network is trained to predict the representation that a target network, updated as a slow-moving average of the online network, produces for another view of the same image. DINO [66] interpreted a similar scheme as self-distillation without labels and found that the features of self-supervised Vision Transformers contain explicit information about the segmentation of objects, which does not emerge as clearly with supervised training. *Language supervision* offers yet another signal: CLIP [67] learned from 400 million image–text pairs to predict which caption belongs to which image, which allows zero-shot transfer to new tasks.

Masked modelling learns by reconstructing hidden parts of the input. Masked autoencoders (MAE) [68] mask a large fraction of image patches and train the network to reconstruct their pixels, and VideoMAE [69] applied the same principle to video with very high masking ratios. UMT [70] instead aligns the unmasked tokens of a video student with the features of an image foundation model, which speeds up training and gives semantically rich features; for the UMT-L encoder on which SalFoM builds (Section 3.4), that teacher is CLIP [10]. Such video foundation models have since been scaled further: InternVideo2 [71] combines masked video modelling, cross-modal contrastive learning and next-token prediction and scales the video encoder to six billion parameters.

2.3.4. Joint-Embedding Predictive Architectures

Joint-embedding predictive architectures (JEPA) take a further route. Rather than reconstructing pixels or contrasting views, they predict the *representations* of masked parts of the input from the visible parts. I-JEPA [11] predicts the representations of several target blocks of an image from a single context block, and V-JEPA [12] extends this to video. With x the visible (context) part of a clip, y the full clip, E_θ the encoder, $\bar{E}$ an exponential-moving-average copy of it, P_ϕ the predictor and $\mathcal{M}$ the set of masked tokens, the objective can be written as [14]

$$\mathcal{L}_{\text{JEPA}} = \frac{1}{|\mathcal{M}|} \sum_{i \in \mathcal{M}} \left\| P_\phi\big(E_\theta(x), \Delta_y\big)_i - \text{sg}\big(\bar{E}(y)_i\big) \right\|_1, \tag{2.1}$$

where Δ_y are learnable tokens that mark the positions to be predicted and sg stops the gradient. The moving-average target and the stop-gradient, which prevent the trivial solution in which all representations collapse to a constant, follow the non-contrastive methods described above [14], [65]. Because the target is a representation rather than pixels, the encoder can ignore unpredictable detail. Bardes et al. [12] report that the resulting features perform well on both appearance- and motion-oriented tasks without

adapting the encoder, that is, with a frozen backbone. A related observation for images is that self-supervised features can serve many tasks without fine-tuning [72].

Since the loss in Equation (2.1) is applied only to the masked tokens, the visible tokens receive no direct training signal. Mur-Labadia et al. [14] found that the features of the V-JEPA 2 models consequently carry little local structure (22.2 mIoU in linear-probe semantic segmentation on ADE20K), and proposed V-JEPA 2.1, which adds a loss on the visible tokens and self-supervision at intermediate layers to obtain denser features. Both model families are used in this thesis.

A JEPA model also contains its predictor P_ϕ. Models that learn to predict how their environment evolves are often called *world models*. Ha and Schmidhuber [73], for example, trained a generative recurrent model of reinforcement-learning environments and used its compressed spatio-temporal representation for control. A JEPA video model resembles a world model in that its predictor anticipates unseen content, but it predicts representations rather than pixels, and the content it predicts need not lie in the future: masks may cover any part of a clip. After training, the predictor's error on a new video, the difference between what it predicted for a region and what the encoder actually computed there, is a learned measure of how unexpected that region is. This is the quantity examined as a possible attention cue in Section 5.4.

2.3.5. Transferring and Adapting Pretrained Models

How well pretrained features transfer depends on the depth at which they are taken and on how different the new task is. Yosinski et al. [74] showed that the features of a deep network change from general to task-specific with depth and that their transferability decreases as the source and target tasks diverge. Intermediate layers, which dense prediction models already combine because they offer different resolutions [60], [61], are therefore also likely to carry features that are less specific to the pre-training task than those of the last layer.

A pretrained encoder can be used frozen, with only a task-specific head trained on top, or adapted to the task. Full fine-tuning updates all parameters and can overfit small datasets. It can also damage good features: Kumar et al. [75] showed that fine-tuning distorts pretrained features while the head is still being learned. Across ten benchmarks with distribution shift, it was on average about 2 % more accurate than training only a linear head on the training distribution but about 7 % less accurate on shifted data. Parameter-efficient methods adapt a model with few trainable parameters: adapters insert small bottleneck layers into a frozen network [76], and LoRA [77] adds trainable low-rank updates to the weight matrices of a frozen transformer, matching full fine-tuning on language tasks while training a small fraction of the parameters. Conditioning information, such as an audio signal, can be injected into a network through feature-wise linear modulation (FiLM), which scales and shifts intermediate

features as a function of the conditioning input [78]. Finally, averaging the predictions of several independently trained models, a deep ensemble, is a simple way of improving accuracy and calibration [79]. Each of these techniques is evaluated in this thesis.

2.4. Eye-Tracking Datasets for Video

Table 2.1 lists the video eye-tracking datasets relevant to this thesis. They differ in size, content and, importantly, viewing task.

Table 2.1: Video eye-tracking datasets relevant to this thesis

Dataset	Videos	Viewers	Content and task
DHF1K [8], [9]	1,000	17	varied YouTube videos, free viewing; test labels withheld
LEDOV [80]	538	32	varied videos, free viewing
Hollywood-2 [81], [82]	1,707	19	movie clips; most viewers performed action recognition
UCF Sports [82], [83]	150	19	sports videos; same protocol as Hollywood-2
Hollywood-2, Vig et al. [84]	1,707	4 + 1 pooled	independent recording during action recognition
DIEM [1]	84	42	advertisements, trailers, documentaries; audio on
Coutrot1 [2]	60	72	four visual categories: one or several moving objects, landscapes, faces
Coutrot2 [85]	15	40	four people in a meeting
ETMD [86], [87]	12	10	clips from Hollywood movies
AVAD [88]	45	16	short audio-visual scenes
SumMe [87], [89]	25	10	user-made videos

DHF1K [8], [9] is the largest free-viewing benchmark: 1,000 videos annotated by 17 observers, split into 600 training, 100 validation and 300 test videos whose annotations are withheld and evaluated by the dataset authors on request. Hollywood-2 and UCF Sports are action recognition datasets [81], [83] for which Mathe and Sminchisescu [82] recorded the gaze of 19 observers, most of whom were asked to recognise the action; their ground truth is therefore task-driven and concentrated on the actors. Vig et al. [84] independently recorded gaze on Hollywood-2 during action recognition and released the raw recordings openly. The six audio-visual datasets of the lower part of Table 2.1 were assembled into a common benchmark by Tsiami et al. [90]; the viewer counts in the table follow their description and that of later work [91]. Fixation maps in all these datasets are converted into continuous density maps by Gaussian smoothing, and the binary fixation maps are kept for location-based metrics.

The differences in task are central to interpreting results across datasets. Free viewing (DHF1K, DIEM) captures where people look spontaneously; the action recognition task (Hollywood-2, UCF Sports) concentrates gaze on the actors and their movements; and in conversations (Coutrot2) gaze follows the speaker. Because the viewing task changes

where people look [27], [31], a model trained on one condition need not transfer to another, which is why results on these datasets are usually obtained after fine-tuning on each.

2.5. Evaluation Metrics

No single metric captures all aspects of a saliency map. Bylinskii et al. [17] analysed the common metrics and grouped them into *location-based* metrics, which evaluate the predicted map at the fixated locations, and *distribution-based* metrics, which compare the predicted map with the continuous ground-truth density. The location-based family includes the normalised scanpath saliency (NSS) [92], the mean of the standardised prediction at fixated pixels, and the area under the ROC curve in the variant of Judd et al. (AUC-J) [26], which separates fixated from non-fixated pixels. The shuffled AUC (sAUC) draws its negatives from the fixations made on other images or videos, so that a model is not rewarded merely for predicting the centre bias [25], [93]. The distribution-based family includes the linear correlation coefficient (CC), the histogram intersection (SIM) and the Kullback–Leibler divergence (KLD). Le Meur and Baccino [19] discuss the strengths and weaknesses of these comparisons, and Wilming et al. [20] concluded that no single measure has all the properties one would want, although the AUC and the KL divergence combine many of them. Section 4.8 gives the exact definitions used in this thesis.

The metrics respond differently to the same map [17]. KLD heavily penalises predictions that assign almost no probability to regions that observers fixate; SIM measures the overlap of two distributions and is sensitive to ground-truth mass that the prediction misses; the AUC metrics depend only on the ranking of saliency values; and sAUC penalises centre bias. Kümmerer et al. [21] proposed to compare models in terms of information gain over a baseline, which unifies several metrics in a probabilistic framework, and later showed that the ranking of models depends on how a probabilistic prediction is turned into a map for each metric. They argued that each metric should be given the map that is optimal for it under the model's prediction [94]; in practice, video saliency benchmarks score one map per model with all metrics. In video saliency practice, the metrics are computed per frame, averaged over the frames of each video and then over videos, usually with evaluation code derived from the MIT saliency benchmark [17]; small differences between implementations make numbers from different papers only approximately comparable.

2.6. Summary

Visual attention is measured through eye movements that are segmented into fixations and summarised as fixation and density maps, whose consistency across observers bounds what any model can achieve. Attention is shaped by image features, by the observer's task and by strong regularities such as the centre bias and the attraction of faces, and in video additionally by motion, events and sound; free viewing of video produces a marked synchrony of gaze across observers, which a task can reduce. Computational models moved from hand-crafted feature contrast to deep networks that predict saliency as a probability distribution and whose success rests largely on pretrained representations. Video representation learning followed a parallel path from three-dimensional convolutions and two-stream designs to transformers and, most recently, to self-supervised encoders trained on large video collections with contrastive, masked or predictive objectives. Among these, the joint-embedding predictive models are distinctive in learning by predicting representations and in containing a predictor whose errors could indicate what is unexpected. Whether such encoders should be adapted or kept frozen depends on the quality of their features and on the size of the target data. The datasets and metrics of the field are diverse in task and in what they reward, which makes careful and consistent evaluation essential. The next chapter reviews how these ingredients have been combined into video saliency models.

CHAPTER THREE

RELATED WORK

Deep video saliency models differ mainly in two respects: how they represent the temporal dimension of a video, and what their encoder has been trained on before it meets eye-tracking data. This chapter follows that development, from recurrent models built on image features (Section 3.1), through three-dimensional convolutional encoders (Section 3.2) and transformers (Section 3.3), to models that build on video foundation encoders (Section 3.4). Audio-visual models, which add the soundtrack, are reviewed separately (Section 3.5), as is the small body of work that treats prediction error as a cue for attention (Section 3.6). The chapter closes with a summary and the research gap that this thesis addresses (Section 3.7).

3.1. Recurrent and Two-Stream Models

The first deep video saliency models extended image saliency networks with mechanisms for motion and memory. Bak et al. [95] used two-stream networks, one stream on the frames and one on optical flow, and compared ways of fusing them. Lai et al. [96] coupled appearance and motion streams more tightly through dense residual cross connections and added a convolutional gated recurrent unit for long-term dynamics. Gorji and Clark [97] took the opposite route and augmented a strong static saliency model with recurrent "attentional push" pathways, which model how the photographer and the actors in a scene direct the viewer's attention.

Recurrence on top of per-frame convolutional features became the dominant design. DeepVS [80] combined an object-to-motion network with a saliency-structured convolutional LSTM, motivated by the observation, made on its own large eye-tracking dataset, that moving objects attract gaze. ACLNet [8], [9], introduced together with DHF1K, added a supervised attention module that encodes static saliency to a CNN-LSTM, so that the recurrent part can concentrate on temporal change; it was trained with static images from SALICON alongside the videos. Linardos et al. [98] showed that a simple exponential moving average of features can perform on a par with a convolutional LSTM, which suggests that much of the temporal modelling in these networks amounts to smoothing. UNISAL [99] pursued a lightweight encoder–RNN–decoder design trained jointly on image and video saliency data, with domain-adaptive priors, fusion and smoothing to cope with the differences between datasets, and reached strong results at a fraction of the size of its competitors.

These models share a limitation: their spatial features come from networks pretrained on images, and their temporal modelling is confined to a recurrent state on top of those features. Their reliance on static saliency data, which helped them markedly [9], is a sign that the video data alone were not sufficient to train them.

3.2. Three-Dimensional Convolutional Encoder–Decoders

A second generation replaced per-frame features by spatio-temporal features from three-dimensional convolutional networks pretrained for action recognition, typically S3D [51] trained on Kinetics [50]. TASED-Net [100] encodes a clip of consecutive frames and decodes the features spatially while aggregating the temporal dimension, producing one map per clip; maps for a whole video are obtained by sliding the window, under the assumption that a limited number of past frames suffices to predict the current one. HD2S [101] supervises intermediate "conspicuity" maps at several levels of abstraction and adds domain adaptation through gradient reversal and dataset-specific priors. ViNet [102] decodes features from several levels of an action-recognition encoder with trilinear upsampling and three-dimensional convolutions; it is causal, runs in real time and, despite using no audio, outperformed audio-visual models on six audio-visual datasets.

Later work in this family focused on efficiency and on the choice of encoder. TinyHD [103] combines several simple heterogeneous decoders with hierarchical map distillation from a larger teacher, and predicts several output frames per forward pass to reduce cost. MSFF-Net [104] fuses multi-scale spatio-temporal features through a bi-directional pyramid with attention-guided fusion. Most recently, ViNet-S and ViNet-A [45] paired a lightweight decoder with either an S3D encoder or a SlowFast encoder pretrained for spatio-temporal action localisation; their ensemble, ViNet-E, reported the best results among the compared methods on UCF Sports and Hollywood-2. The strength of the action-localisation encoder on these two datasets, whose gaze was recorded during action recognition [82], illustrates how much the pre-training of the encoder matters.

Three-dimensional convolutions capture motion well, but their receptive field in time is limited by the kernel sizes, and their encoders were trained for classification on a fixed set of action categories.

3.3. Transformer-Based Models

Self-attention relaxes the locality of convolutions. STSANet [105] inserts spatio-temporal self-attention modules at several levels of a three-dimensional convolutional backbone to capture long-range relations between time steps, and fuses the levels with an attentional

multi-scale module. VSFT [106] was the first purely transformer-based model in the field and extended the task to forecasting the saliency of future frames. TMFI-Net [107] builds on multi-level features of a video transformer, enriches them top-down with high-level semantics in a semantic-guided encoder, and decodes them hierarchically with attention along the channel, temporal and spatial dimensions. THTD-Net [108] argues that the temporal dimension of transformer features should be kept high throughout decoding, and reaches results comparable to multi-branch decoders with a single-branch one. UniST [48] uses one saliency-aware transformer for both video saliency prediction and video salient object detection, and TM2SP [109] combines a Video Swin encoder [59] with a multi-level spatio-temporal feature pyramid guided by a pretrained image saliency encoder.

These models raised the published DHF1K test correlation from about 0.52 to about 0.55 [10]. They remain, however, trained end to end, or with an encoder pretrained for supervised action recognition, on the few hundred videos of a saliency dataset, and they are large: TMFI-Net has about 61 million parameters and STSANet about 169 million [45].

3.4. Video Foundation Models for Saliency

Foundation models trained with self-supervision on large video collections offer a different starting point. SalFoM [10] was the first video saliency model built on such an encoder. It fine-tunes the UMT-L encoder [70], a ViT-L trained by masked video modelling guided by an image–text teacher, together with a heterogeneous decoder that includes a locality-aware spatio-temporal transformer. Given 16 frames at 224×224 pixels, it predicts the saliency map of the last frame, and maps for a whole video are produced with a sliding window. SalFoM reported the best DHF1K test results among the methods it compared with (CC 0.569, NSS 3.353), but it was weaker than earlier models on UCF Sports, which its authors attribute to the task-driven gaze of that dataset and to overfitting on its small training set [10].

The same group has since combined foundation encoders with explicit long- and short-term temporal modelling [110] and distilled them into smaller models to reduce their cost [111]. A different direction couples a multimodal large language model with a saliency decoder to produce saliency maps together with textual explanations, trained on a new dataset of gaze annotated with reasons [112].

Two properties of these approaches are relevant here. First, the encoder is adapted to the saliency data, either by full fine-tuning or by distillation; whether a frozen foundation encoder is sufficient, and how it compares with an adapted one, was not examined in the studies whose full texts could be consulted. Second, the sliding-window design runs the large encoder once for every output frame, which makes inference

expensive.

3.5. Audio-Visual Saliency Models

Because most videos have sound, and sound influences where people look [2], [33], a separate line of work fuses the soundtrack with the frames. STAViS [90] combined a spatio-temporal visual saliency network with auditory features, learned to localise sound sources, and fused the two saliencies; it also introduced the protocol, with three folds per dataset, under which six audio-visual datasets are now evaluated. CASP-Net [113] addresses the temporal inconsistency between the two modalities with a consistency-aware predictive coding module, and DiffSal [91] formulates prediction as conditional generation with a diffusion model conditioned on audio-visual features. Models based on implicit neural representations map space–time coordinates to saliency values and fuse the modalities parametrically [114], [115], and the relevance-guided fusion of Yu et al. [116] adjusts how much audio is retained according to its semantic relevance to the visual content.

Whether these models actually use the audio is disputed. Jain et al. [102] found that an audio-augmented variant of ViNet became agnostic to its audio input after training and observed the same behaviour in an earlier audio-visual model. Agrawal et al. [117] examined a range of audio modules and fusion techniques on ten datasets and found that none improved on visual-only models. On several audio-visual datasets, visual-only models such as ViNet-E report results that equal or exceed those of audio-visual models [45].

3.6. Prediction Error as an Attention Signal

The idea that attention is drawn to what violates expectation has a long history. Itti and Baldi formalised "surprise" as the divergence between an observer's prior and posterior beliefs, detected surprising events in video with low-level feature models [34], and showed that this measure predicted human gaze shifts better than the other metrics they tested [13]. Predictive-coding accounts of the visual cortex propose that neural responses signal the error of predictions made from context [35], and deep networks built on this principle have been trained to predict future video frames [118].

Modern self-supervised video models contain an explicit predictor. In the joint-embedding predictive architectures I-JEPA [11] and V-JEPA [12], a predictor network learns to predict the representations of masked regions from the visible ones. Such a predictor offers, in principle, a learned and high-level analogue of surprise: its error indicates which parts of a scene the model could not anticipate. In the literature reviewed for this thesis, the error of such a learned latent predictor has not been

evaluated against human fixations. CASP-Net uses predictive coding in a different sense, to align audio and visual features [113].

3.7. Summary and Research Gap

Table 3.1 summarises the most relevant works. Three observations emerge from the review.

Table 3.1: Summary of the closest related works

Work	Approach	Evaluation	Limitation relevant to this thesis
ACLNet [9]	CNN-LSTM with static-saliency attention	DHF1K, Hollywood-2, UCF	image features; needs static data
UNISAL [99]	lightweight encoder–RNN–decoder, image and video jointly	DHF1K, Hollywood-2, UCF, SALICON	recurrent temporal model on image features
TASED-Net [100]	S3D encoder, temporally aggregating decoder	DHF1K, Hollywood-2, UCF	one map per clip; action-recognition encoder
ViNet-S/A/E [45]	S3D or action-localisation encoder, light decoder	DHF1K val, Hollywood-2, UCF, six AV sets	supervised encoders; one map per clip
TinyHD [103]	heterogeneous decoders, distillation, multi-output	DHF1K, Hollywood-2, UCF	efficiency focus; CNN encoder
STSANet [105]	self-attention on a 3D CNN	DHF1K, Hollywood-2, UCF, DIEM	large (about 169 M parameters)
TMFI-Net [107]	video transformer, multi-scale integration	four video datasets	trained on few videos; large
SalFoM [10]	fine-tuned UMT-L, heterogeneous decoder	DHF1K test, Hollywood-2, UCF	full fine-tuning; one encoder pass per output frame
STAViS [90]	visual saliency with sound localisation and fusion	six AV datasets	audio benefit later disputed [117]
DiffSal [91]	audio-visual conditional diffusion	six AV datasets	iterative generation; audio-dependent design
Itti and Baldi [13]	Bayesian surprise of low-level features	gaze in video	hand-crafted features, no learned predictor

First, progress in video saliency prediction has tracked progress in the encoder: from image networks with recurrence, through action-recognition three-dimensional networks, to video transformers and, most recently, video foundation models. The decoders, by contrast, have converged on variations of multi-scale fusion. Yet the encoders used for saliency have all been trained with supervision or with objectives that reconstruct or align tokens, and they have been adapted to the saliency data. Encoders trained with a feature-prediction objective, such as V-JEPA [12], whose authors report that its features perform well without adapting the model's parameters, have not been evaluated for

this task in the peer-reviewed studies whose full texts could be consulted.

Second, the evaluation practice of the field makes small differences hard to interpret. DHF1K results are reported on the validation or the test set, published numbers for the same model differ between papers by up to 0.03 CC [45], [48], and comparisons are almost always made against numbers copied from other papers rather than against re-scored predictions with an uncertainty estimate.

Third, the idea that motivated this work, that a learned video model's prediction error marks where people look, has support in classical models of surprise [13] but had not been tested with the predictor of a modern self-supervised video model.

This thesis addresses these gaps. It evaluates a frozen V-JEPA 2 encoder with a light decoder that predicts many frames per pass, compares it with SalFoM, the closest foundation-model approach, under one validated evaluator with paired statistics, tests the encoder's own predictor as a source of attention cues, and examines which design choices still matter once the encoder is fixed. It resembles SalFoM in building on a video foundation encoder and shares its multi-scale decoder design with much of the field; it differs in keeping the encoder frozen, in the objective with which that encoder was trained, in the multi-frame output, and in an experimental protocol whose decision rules were written in advance.

CHAPTER FOUR

METHODOLOGY

This chapter describes how the research questions of Section 1.4 were investigated. It first sets out the research design, which fixes every decision rule before the experiment that it governs (Section 4.1), and formalises the prediction task (Section 4.2). It then describes the eye-tracking data and their preparation (Section 4.3), the proposed model, SURGE-D, as it is implemented (Section 4.4), the components that were implemented and tested but not retained (Section 4.5), the training and inference procedures (Sections 4.6 and 4.7), and the evaluation metrics and statistical analysis (Section 4.8). The chapter closes with the ethical considerations of the work (Section 4.9). Every equation in this chapter corresponds to code in the accompanying repository; where the implementation departs from the original research proposal, the departure is stated.

4.1. Research Design

The study follows a design-science approach in which an artefact, a video saliency model, is built and then evaluated through controlled experiments. Two features of the design deserve emphasis, because they determine how the results in Chapter 5 may be read.

Pre-registered decision rules. The original proposal contained several components whose value was uncertain: the choice of video encoder, a "surprise" signal derived from the encoder's predictor, parameter-efficient adaptation of the encoder, static-image co-training, and others. Rather than include all of them and report the final configuration, each component was treated as a hypothesis with an explicit test. Before each experiment was run, its decision rule (the quantity to be measured, the data on which it is measured, and the threshold that decides the outcome) was written into a dated research log, and the rule was applied unchanged afterwards. Table 4.1 lists these rules. This procedure guards against two familiar sources of optimism in deep-learning research: choosing among variants after seeing test results, and adjusting thresholds after seeing which side of them a result falls on. The log is kept in the repository of the thesis, so the rules are specified in advance in an internal, dated record rather than registered with an external service. Choices that were fixed without a test, such as the decoder architecture, the tapped layers, the loss weights, the squashing of the input and the learning rates (Table 4.3), were not governed by a rule.

Separation of selection and testing. All model selection (the best training epoch, the output blur, every gate in Table 4.1) was performed on internal holdout sets carved out of the training data. Each test set was scored once per reported model, after all of that model's settings had been fixed. For DHF1K, whose official test annotations are withheld by the dataset authors, the 100-video validation set serves as the test set. No reported model or setting was selected on it. It did inform two decisions about the predictions prepared for the official test set; these are listed with every other use of the set in Table A.4.

The rule for published results (marked *) differs from the others in one respect. The rule written before the first audio-visual experiment required only a higher fold-mean CC than the published value, and noted that cross-paper numbers disagree by a few hundredths. The tie margin of 0.01 was introduced when the first audio-visual results were interpreted and was then applied unchanged to every later comparison with published numbers.

Unit of analysis. Saliency metrics are computed per frame and averaged within each video; the video is the unit of analysis, and uncertainty is estimated by resampling videos (Section 4.8.2). The margins of the early rules (0.003 for single-seed ablations and 0.005 for gate G1) were set by judgement before any seed variation had been measured. The seed-to-seed spread later observed for the main model, about 0.003 CC, informed the margins of the later gates for ViT-g and for 32-frame clips.

4.2. Problem Formulation

Let a video be a sequence of frames $I_1, \dots, I_N$ with $I_n \in \mathbb{R}^{H \times W \times 3}$. For each frame, an eye-tracking dataset provides a binary fixation map $F_n \in \{0,1\}^{H \times W}$, marking the pixels fixated by at least one observer, and a continuous fixation density Y_n, obtained by blurring the fixations with a Gaussian kernel. Video saliency prediction is the task of estimating, for every frame, a probability distribution $\hat{S}_n$ over pixel locations,

$$\hat{S}_n(p) \geq 0, \qquad \sum_p \hat{S}_n(p) = 1, \tag{4.1}$$

that approximates where observers look in frame n. Treating saliency as a distribution, rather than as an unnormalised score map, is the probabilistic view advocated by Kümmerer et al. [21], [94]; it makes the training objective and the distribution-based metrics well defined.

The model sees a temporal window rather than a single frame. A clip of $T = 16$ frames is sampled with a stride of $s = 2$, so that the sampled indices are $n_k = n_0 + ks$ for $k = 0, \dots, T-1$; at the 30 frames-per-second rate of DHF1K the clip spans 31

Table 4.1: Decision rules, each written before the experiment that it governs, except the one marked * (see text). Rules marked † were written in a follow-up after the first draft of this thesis and rules marked ‡ in a second follow-up, in each case before the runs they govern. Holdout CC is measured on the internal 60-video DHF1K holdout unless stated otherwise; outcomes are reported in Chapter 5

Decision	Rule
Evaluator validity	Re-scoring the same maps with an independent evaluator must agree within 0.002 (CC, SIM) and 0.01 (NSS, KLD)
Backbone (gate G1)	Highest holdout CC wins; the proposal's choice, V-JEPA 2.1, is kept if within 0.005 CC of the winner
Surprise (gate G2)	Enabled only if a surprise map beats both token atypicality and frame difference on NSS and has positive information gain over the centre prior
LoRA adaptation	Adopted if the three-seed mean holdout gain is at least +0.003 CC and LoRA wins in at least two of three seeds
Single-seed ablations	A difference is reported only if $\lvert\Delta\text{CC}\rvert \geq 0.003$ on the holdout
Claim over a baseline	The 95 % bootstrap interval of ΔCC over test videos must lie above zero
ViT-g backbone	Adopted if its holdout CC is at least 0.5833 (the ViT-L three-seed mean plus 0.005)
Training on 600 videos	Replaces the 540-video models if the ensemble's ΔCC interval is above zero and it is worse on at most two of six metrics
32-frame clips	Pursued if the holdout CC is at least 0.5904 (the 16-frame three-seed mean plus 0.005)
Audio fusion	"Helps" on a dataset only if the fold-mean ΔCC is at least +0.010 and matched audio beats rotated audio by at least 0.005
Published results*	Comparisons with numbers from other evaluation code treat margins below 0.01 CC as ties
Baseline blur†	A claim over SalFoM stands only if it also holds against SalFoM's maps with the blur of the same grid that is best for SalFoM on the test set
Causal control†	SURGE-D is better than SalFoM under an identical protocol only if the 95 % interval of ΔCC between its causal three-seed ensemble and SalFoM at its best blur lies above zero
Ablations, three seeds†	An effect if the 95 % interval of the seed-averaged ΔCC excludes zero and $\lvert\Delta\text{CC}\rvert \geq 0.003$; equivalence if the 90 % interval lies within ±0.003; otherwise inconclusive
Training-set size‡	With 135, 270 and 540 training videos and three seeds, accuracy at 540 videos is limited by the amount of data if halving it is an effect with a negative sign, and is not if halving it is equivalent, both as defined for the three-seed ablations; otherwise inconclusive
Incremental information‡	A prediction-error map adds information to the centre prior, frame difference and token atypicality if the 95 % interval of its cross-fitted information gain lies above zero; the gain is negligible if the 90 % interval lies within ±0.01 bits per fixation; otherwise inconclusive
Measured time‡	The statement that offline SURGE-D costs less per output frame than SalFoM is withdrawn if SalFoM's measured time per output frame is less than twice that of SURGE-D with ViT-g at a batch size of one or of eight

original frames, or about one second. The stride is defined in frames, so on datasets with other frame rates the clip covers a correspondingly different duration. Indices that fall outside the video are clamped to its first or last frame. The model predicts maps for all T sampled frames but only the central $T/2$ positions,

$$\mathcal{O} = \{T/4, \ldots, 3T/4 - 1\} = \{4, \ldots, 11\}, \tag{4.2}$$

are used for training and inference. Each prediction therefore has at least eight original frames (about 0.27 s) of context on either side. The formulation is thus *non-causal*: a frame's map depends on frames that follow it. This suits offline applications such as video compression or summarisation but not streaming use, and it differs from methods that predict the last frame of a clip from past frames only, such as SalFoM [10]. The difference is returned to when results are compared in Chapter 5.

Causal variant. To measure how much of the accuracy of SURGE-D depends on frames after the predicted one, a causal variant was trained as a control (Section 5.5.2). It predicts the map of frame n from the window $n - 30, n - 28, \ldots, n$ (sixteen frames with a stride of two; indices below zero are clamped to the first frame) and uses the output at the last sampled position. During training, each clip is supervised on its last sampled frame only. The causal variant sees no frame after the predicted one, as in SalFoM, but it needs one forward pass for every output frame.

4.3. Datasets

4.3.1. Data Collection

No new eye-tracking data were collected. The study uses ten public datasets, summarised in Table 4.2. All were read from their original distributions; the splits and holdout lists used here are provided with the code.

DHF1K [8], [9] is the primary benchmark. Its 600 training videos were split once, with a fixed random seed, into 540 training videos and a 60-video holdout used for all model selection. The official validation set (videos 601–700, 60,273 frames) is the test set of this thesis. Predictions for the 300 official test videos were generated with the final model and packaged in the format required by the dataset authors; their scores were not available when this thesis was written and are not reported.

UCF Sports [83] and Hollywood-2 [81] are action-recognition datasets for which Mathe and Sminchisescu [82] recorded the gaze of observers performing action recognition. For UCF Sports the standard 103/47 split was used, with ten training videos held out. The Mathe and Sminchisescu ground truth for Hollywood-2 could not be obtained during the project; instead, the independently recorded and openly licensed gaze data

Table 4.2: Eye-tracking datasets used in this thesis. "Holdout" is the internal selection set drawn from the training part; the test part is scored once per model

Dataset	Content and viewing condition	Split used (train / holdout / test)
DHF1K [8], [9]	1,000 YouTube videos, 640 × 360 at 30 fps, free viewing, 17 observers; official split 600 / 100 / 300 (test labels withheld)	540 / 60 / 100 (official validation set as test)
UCF Sports [82], [83]	150 sports videos; gaze recorded during action recognition; standard split 103 / 47	93 / 10 / 47
Hollywood-2 [81] with the gaze of Vig et al. [84]	1,707 movie clips (823 / 884); four observers who watched every clip plus a pooled fifth observer, recorded during action recognition	741 / 82 / 884
DIEM [1]	84 videos of many genres; audio present	one standard split, 64 training and 17 test videos
Coutrot1 [2]	60 clips in four visual categories	three folds each, every
Coutrot2 [85]	15 clips of four people in a meeting	listed test video tested
ETMD [86], [87]	12 clips from six Hollywood movies	once [90] (SumMe: 23 of 25); holdout = 10 % of each
AVAD [88]	45 short audio-visual clips	training fold (1–4 videos)
SumMe [87], [89]	25 user-made videos	
SALICON [36]	10,000 static training images	co-training ablation only

of Vig, Dorr and Cox [84] were used[1]. These data come from fewer observers and cover whole clips rather than the shot-segmented clips used in the saliency literature, so results on them are reported under the separate name "Hollywood-2 (Vig gaze)" and are not compared with published Hollywood-2 numbers.

The six audio-visual datasets (DIEM [1], Coutrot1 [2], Coutrot2 [85], ETMD [86], AVAD [88] and SumMe [89]) were used in the preprocessed form and with the fold lists released by Tsiami et al. [90], whose protocol has become the standard for these data: DIEM has one fixed split, which uses 81 of its 84 videos, and each of the other five datasets has three folds in which every listed test video appears in exactly one test set. In SumMe, 23 of the 25 videos are test videos; the other two appear only in training lists. The gaze data for ETMD and SumMe were recorded by Tsiami et al. [87]. SURGE-D is a visual model; the soundtracks were used only in the audio experiment of Section 4.5.3. SALICON [36], a dataset of static images with mouse-tracking saliency, was used only in the static co-training ablation.

[1] Figshare item 93951, released under the CC BY licence; every file was checked against its published checksum.

4.3.2. Preprocessing

Frames. Frames were read as distributed (DHF1K, UCF Sports and the STAViS releases) or decoded once with OpenCV (Hollywood-2, 487,556 frames). For the model, every frame is resized to 256×256 pixels with area interpolation, which changes the aspect ratio of 16:9 video, and normalised with the ImageNet channel means and standard deviations [119]. Squashing was chosen over cropping because cropping removes image borders and alters the centre-bias statistics of the fixation maps; its effect on accuracy was not ablated.

Training targets. For each target frame, the ground-truth density is resized with area interpolation to the 64×64 output grid and normalised to sum to one, and the fixated pixels of the binary fixation map are counted on the same grid. The resulting pair (Y_t, F_t) is used by the loss (Section 4.6). Evaluation, in contrast, always takes place at the native resolution of the ground truth (Section 4.8).

Frames without fixations. In the STAViS releases and in the Hollywood-2 (Vig gaze) maps, some frames have no recorded fixation, for example during blinks. These frames are excluded from scoring. DHF1K and UCF Sports contain no such frames.

Hollywood-2 (Vig gaze) maps. The Vig gaze recordings provide raw 1000 Hz samples with confidence values and saccade annotations. A frame f of a clip with frame rate r was taken to cover the interval $[(f-1)/r,\ f/r)$. For each observer and frame, the samples with confidence 1.0 that do not fall inside an annotated saccade were collected. If any remained, the observer contributed one fixation point, the median of those samples, rounded and clipped to the frame. The binary fixation map thus has at most five points per frame. The density map is this binary map blurred with a Gaussian of $\sigma = 20$ pixels, following the DHF1K recipe [8]. Of the 487,556 frames, 99.8 % received at least one point, with 4.53 points per frame on average. For model selection, 82 of the 823 training clips were held out with a fixed seed.

Audio features. For the audio experiment only, each soundtrack of the STAViS release was mixed to mono and resampled to 16 kHz, and the frozen Audio Spectrogram Transformer [120] fine-tuned on AudioSet [121] was applied to consecutive 10.24 s chunks[2]. Averaging the patch tokens over the twelve frequency rows of the patch grid gives one 768-dimensional vector per time column, that is, roughly one every 0.1 s. The feature of video frame k (centred at time $(k-\frac{1}{2})/r$) is the mean of the columns whose centres lie within ± 0.5 s of that time. Videos without a soundtrack receive zero features.

[2]Checkpoint `MIT/ast-finetuned-audioset-10-10-0.4593`, BSD-3-Clause licence.

4.4. Proposed Architecture

SURGE-D consists of a frozen self-supervised video encoder and a light trainable decoder (Figure 4.1). The decoder reads four intermediate layers of the encoder, rebuilds a multi-scale feature pyramid, decodes it top-down with factorised spatio-temporal convolutions, and adds a learned centre prior before normalising each output frame into a probability distribution. Only the decoder and the prior are trained: 6.99 million parameters with the ViT-L encoder and 7.39 million with ViT-g. The name derives from the original design of the project, a *SURprise-Gated Encoder*–decoder in which the predictor's error was to modulate the decoder; the suffix D marks the reduced design that remained after the tests of Section 4.5, in which the surprise branch did not pass its gate.

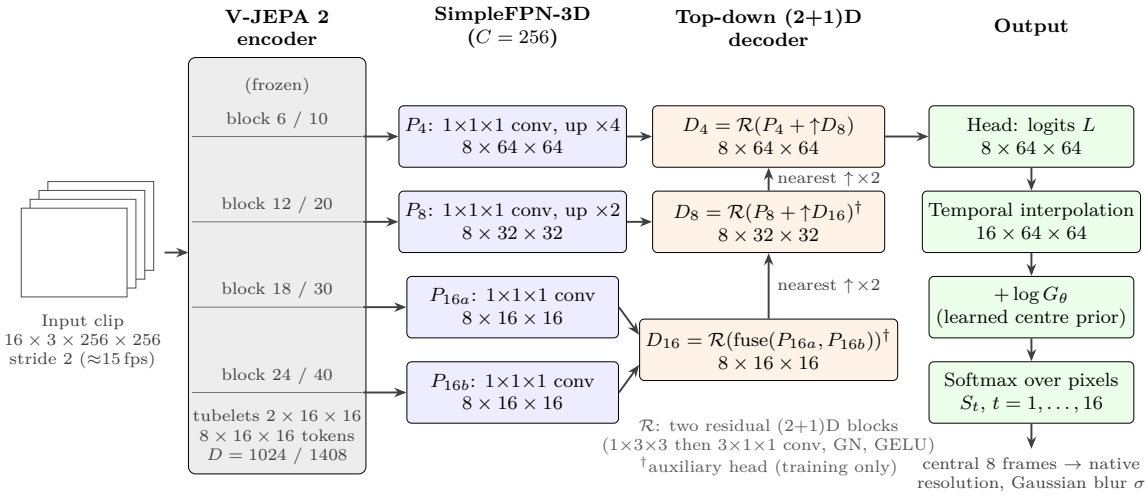


Figure 4.1: Architecture of SURGE-D. A clip of 16 frames passes through the frozen V-JEPA 2 encoder; four intermediate layers (blocks 6/12/18/24 of ViT-L or 10/20/30/40 of ViT-g) feed a three-dimensional simple feature pyramid, which a top-down (2+1)D decoder turns into logits at 64 × 64 resolution for eight tubelet steps. After temporal interpolation to 16 frames and the addition of a learned centre prior, a softmax over pixels yields one saliency distribution per frame; the central eight frames are kept. Tensor sizes are given as time × height × width

4.4.1. Frozen Video Encoder

The encoder is a Vision Transformer [56] trained with the video joint-embedding predictive architecture (V-JEPA) objective of Bardes et al. [12]. In this objective, which extends the image model I-JEPA [11] to video, parts of a clip are masked and a predictor network is trained to predict the *representations* of the masked regions, produced by an exponential-moving-average target encoder, from the representations of the visible ones. Because the loss is computed in representation space rather than pixel space, the encoder is not required to model unpredictable pixel detail, which Bardes et al. argue leads to features that transfer well with a frozen backbone [12].

The checkpoints used are the publicly released V-JEPA 2 models[3], which follow this objective: as Mur-Labadia et al. describe it, an L_1 loss between the predictor's outputs and the target encoder's representations is applied to the masked tokens only [14]. The same study reports that this choice leaves the visible tokens without a direct training signal and that frozen V-JEPA 2 features consequently carry little local structure, reaching only 22.2 mIoU in linear-probe semantic segmentation on ADE20K. Whether such features suit a dense task such as saliency prediction was therefore an open question, which gate G1 addresses (Section 5.3). Two sizes were used. ViT-L/16 has 24 transformer blocks of width $D = 1024$ and 303.9 million encoder parameters; ViT-g/16 has 40 blocks of width $D = 1408$ and 1,012.2 million parameters (both counts measured on the loaded checkpoints, excluding the predictor). Both embed the video in *tubelets* of $2 \times 16 \times 16$ pixels, so a $16 \times 256 \times 256$ clip becomes $8 \times 16 \times 16 = 2{,}048$ tokens.

Let $h^{(\ell)} \in \mathbb{R}^{2048 \times D}$ denote the output of block ℓ. Four blocks at equal fractions of the depth are tapped, $\ell \in \{6, 12, 18, 24\}$ for ViT-L and $\ell \in \{10, 20, 30, 40\}$ for ViT-g, and each tap is reshaped into a feature volume of size $D \times 8 \times 16 \times 16$. Intermediate layers are used because dense prediction benefits from features at several levels of abstraction [60], [61]. The encoder is frozen: its weights are not updated and no gradient is propagated into it. Freezing was the outcome of an experiment rather than an assumption; the alternative, low-rank adaptation of the encoder, is described in Section 4.5.2 and evaluated in Chapter 5.

4.4.2. Three-Dimensional Simple Feature Pyramid

A plain Vision Transformer produces features at a single resolution. Li et al. [62] showed that a feature pyramid for dense prediction can be built from such a backbone by simple up- and down-sampling of its features. SURGE-D uses a three-dimensional variant built from the four taps rather than from the last layer only. With $\phi(\cdot)$ denoting group normalisation [122] followed by a GELU non-linearity, W_ℓ a $1 \times 1 \times 1$ convolution from D to $C = 256$ channels, and $U(\cdot)$ a spatial $\times 2$ transposed convolution (kernel and stride $1 \times 2 \times 2$) followed by ϕ, the pyramid levels are

$$\begin{aligned} P_4 &= U\Big(U\Big(\phi(W_1 h^{(\ell_1)})\Big)\Big) \in \mathbb{R}^{C \times 8 \times 64 \times 64}, & P_8 &= U\Big(\phi(W_2 h^{(\ell_2)})\Big) \in \mathbb{R}^{C \times 8 \times 32 \times 32}, \\ P_{16a} &= \phi(W_3 h^{(\ell_3)}) \in \mathbb{R}^{C \times 8 \times 16 \times 16}, & P_{16b} &= \phi(W_4 h^{(\ell_4)}) \in \mathbb{R}^{C \times 8 \times 16 \times 16}. \end{aligned} \tag{4.3}$$

The shallowest tap feeds the finest level, as in standard feature pyramids, and the two deepest taps both remain at the token resolution of 1/16. The temporal resolution of

[3] `facebook/vjepa2-vitl-fpc64-256` and `facebook/vjepa2-vitg-fpc64-256`, released under the MIT licence through the Hugging Face hub and loaded with the Transformers library [16]. The technical report that introduced them had not appeared in a peer-reviewed venue when this thesis was written and is therefore not cited; their properties are taken from the checkpoints themselves and from the peer-reviewed follow-up study of Mur-Labadia et al. [14].

eight tubelet steps is preserved throughout.

4.4.3. Top-Down (2+1)D Decoder

The decoder refines the pyramid from coarse to fine. Its building block is a residual block that factorises a three-dimensional convolution into a spatial and a temporal part, in the manner of R(2+1)D [52],

$$\mathcal{B}(x) = \mathrm{GELU}\Big(x + \mathrm{GN}\Big(\mathrm{Conv}_{3\times1\times1}\Big(\mathrm{GELU}(\mathrm{GN}(\mathrm{Conv}_{1\times3\times3}(x)))\Big)\Big)\Big), \tag{4.4}$$

where the subscripts give the kernel size as time × height × width and GN is group normalisation. Let $\mathcal{R}$ denote two such blocks in sequence and $\uparrow$ nearest-neighbour spatial upsampling by a factor of two. The decoder computes

$$\begin{aligned} D_{16} &= \mathcal{R}\Big(\phi(\mathrm{Conv}_{1\times1\times1}[P_{16a}; P_{16b}])\Big), \\ D_8 &= \mathcal{R}\Big(P_8 + \uparrow D_{16}\Big), \\ D_4 &= \mathcal{R}\Big(P_4 + \uparrow D_8\Big), \end{aligned} \tag{4.5}$$

where $[\cdot;\cdot]$ concatenates channels. A head with a $1 \times 3 \times 3$ convolution to 128 channels, a GELU and a $1 \times 1 \times 1$ convolution to one channel turns D_4 into logits $L \in \mathbb{R}^{8\times64\times64}$. During training, two auxiliary heads ($1 \times 1 \times 1$ convolutions) produce additional logits from D_{16} and D_8; they are upsampled to 64×64 and supervised with the same targets, a form of deep supervision [123]. The auxiliary heads are not used at inference.

4.4.4. Temporal Upsampling, Centre Prior and Output

The logits have one time step per tubelet. They are interpolated trilinearly to the $T = 16$ sampled frames, giving L_t for $t = 1, \ldots, T$.

Human fixations in videos are biased towards the centre of the screen, partly because of how scenes are framed and partly because of viewing strategy [5], [32]. SURGE-D models this bias with a learned, input-independent prior. With pixel positions $p \in [0,1]^2$ in normalised coordinates, the prior is a mixture of $K = 8$ Gaussians with diagonal covariance,

$$\log G_\theta(p) = \alpha \left[\log \sum_{k=1}^{K} \pi_k \, \mathcal{N}\Big(p; \mu_k, \mathrm{diag}(\sigma_k^2)\Big) - \log \sum_{p'} \sum_{k=1}^{K} \pi_k \, \mathcal{N}\Big(p'; \mu_k, \mathrm{diag}(\sigma_k^2)\Big) \right], \tag{4.6}$$

where the mixture weights $\pi = \mathrm{softmax}(w)$, the means μ_k, the scales σ_k and the strength α are learned (41 parameters in total). The means are initialised on a grid in the central $[0.35, 0.65]^2$ region, the scales at 0.25 and the strength at one. The prediction for frame

t is

$$S_t(p) = \frac{\exp\Big(L_t(p) + \log G_\theta(p)\Big)}{\sum_{p'} \exp\Big(L_t(p') + \log G_\theta(p')\Big)}, \tag{4.7}$$

a softmax over the 64×64 positions, so that every output frame is a probability distribution as required by Equation (4.1). The prior is also added to the auxiliary logits during training.

The design of the decoder follows from three considerations. First, the encoder carries most of the representational capacity, so the decoder can be small; its 7 million parameters are about 2 % of the ViT-L encoder. Second, the (2+1)D factorisation adds temporal mixing at every scale at a lower cost than full three-dimensional kernels [52]. Third, predicting maps for many frames in one forward pass amortises the cost of the encoder: a single pass over a 16-frame clip yields usable maps for 16 original frames (Section 4.7). The measured cost is reported in Section 5.5.5.

4.5. Components Evaluated but Not Retained

Four further components were implemented, tested under the rules of Table 4.1, and not retained. They are described here because their negative results are part of the contribution of the thesis.

4.5.1. Predictive Surprise from the V-JEPA 2 Predictor

The original proposal was built around the hypothesis that the prediction error of a self-supervised video model indicates where something unexpected happens, and therefore where people look, in the spirit of the Bayesian surprise account of attention [13]. V-JEPA 2 provides the means to test this, because its released checkpoint includes the predictor used during pre-training.

For a window of 16 sampled frames, the first eight frames (four tubelets) are the context and the last eight the target. The context is encoded on its own, so that no information from the target frames leaks into it, and the predictor maps the context tokens to predicted target tokens $\hat{z}_p$. The target tokens z_p are taken from an encoding of the full window. With LN a parameter-free layer normalisation over the channel dimension, the *raw surprise* of target token p is

$$e(p) = \frac{1}{D} \sum_{c=1}^{D} \Big| \mathrm{LN}(\hat{z}_p)_c - \mathrm{LN}(z_p)_c \Big|. \tag{4.8}$$

A *contextual* variant subtracts the error obtained when the context is replaced by the context of a different video, $e(p) - e_{\text{wrong}}(p)$, so that only the part of the error that depends on the true past remains. Two reference maps tested whether surprise adds

anything beyond simpler cues: *token atypicality*, the mean absolute deviation of each target token from the mean token of its frame, and *frame difference*, the absolute grey-level difference between consecutive frames pooled to the 16×16 token grid. All maps were upsampled to 640×360, blurred with a Gaussian of $\sigma = 16$ pixels and scored on the first frame of each target tubelet of the 60 holdout videos, at about 4 and 15 frames per second. A decoder module that injects the surprise map through zero-initialised feature-wise modulation was implemented but, following the outcome of the gate, never trained.

The mask token of the predictor. The predictor marks every position that it has to predict with a learned *mask token*, the Δ_y of Equation (2.1). The released checkpoints store ten mask tokens, of which only the first was trained; the other nine are zero vectors[4]. The predictor of the model-loading library used here reads the second token by default[5], so gate G2 was originally run with an untrained mask token. This was found after the first draft of the thesis, while the readout described in the next paragraph was being prepared. Every prediction-error map was then computed with both tokens, and the gate was scored again with the trained token, by its own rule. Section 5.4 reports both readouts. SURGE-D uses only the encoder, so none of its results is affected.

Masked-tube error. Predicting the second half of a clip from the first half is not the task on which the predictor was trained: the masks of pre-training are spatial blocks that extend over the whole clip in time (Section 5.4). A second readout therefore follows the geometry of those masks. The 16×16 token grid is divided into sixteen blocks of 4×4 tokens. In each of four passes, the four blocks with the same row and column parity are masked in all eight tubelets, which hides 25 % of the tokens and leaves visible neighbours around every masked block. The encoder processes the visible tokens only, the predictor predicts the masked ones, and the error of Equation (4.8) is computed against the encoding of the full clip. The four passes together give one error for every token. Before use, the encoder restricted to visible tokens was checked in single precision against the encoder of the library: it reproduced the encoding of the full clip when all tokens were visible, and the encoding of the first eight frames when only their tokens were visible (tolerance 10^{-3}). This readout leaves more of the clip visible than pre-training did, so it matches pre-training in the shape of the mask and not in its extent.

[4] Parameter `predictor.embeddings.mask_tokens` of both checkpoints. The released pre-training code selects the token by the index of a group of clips with the same number of frames, and the released configurations define a single group, so only the first token was ever used.

[5] `mask_index = 1` in `VJEPA2PredictorEmbeddings.forward` of Transformers 5.22.

Incremental information. The gate compares stand-alone maps. A map can be a poor predictor on its own and still add information to other cues, so this was tested separately with a log-linear model of the fixation distribution, scored by its information gain over the centre prior [21]. The frame is divided into 45×80 cells of 8×8 pixels. For a set $\mathcal{K}$ of cue maps, the probability that a fixation falls into cell c is modelled as

$$p_{\mathcal{K}}(c) = \frac{\exp\left(w_0 \log g(c) + \sum_{k \in \mathcal{K}} w_k \, \tilde{z}_k(c)\right)}{\sum_{c'} \exp\left(w_0 \log g(c') + \sum_{k \in \mathcal{K}} w_k \, \tilde{z}_k(c')\right)}, \tag{4.9}$$

where g is the empirical centre prior of the 540 training videos and $\tilde{z}_k$ is cue k, upsampled and blurred as for the gate, averaged within each cell and standardised over the cells of the frame. The information gain is measured in bits per fixation,

$$\mathrm{IG}(\mathcal{K}) = \frac{1}{N} \sum_{i=1}^{N} \Big[\log_2 p_{\mathcal{K}}(c_i) - \log_2 g(c_i) \Big], \tag{4.10}$$

where $c_1, \ldots, c_N$ are the cells of the fixations of a frame; the gain is averaged over the scored frames of a video and then over videos. The weights are fitted by maximum likelihood, which is a concave problem, and they are cross-fitted: the 60 holdout videos are divided into two halves by alternating video number, and each half is scored with the weights fitted on the other. With the baseline set $\mathcal{B}$ of frame difference and token atypicality, the *incremental information* of a prediction-error map e is

$$U(e) = \mathrm{IG}(\mathcal{B} \cup \{e\}) - \mathrm{IG}(\mathcal{B}), \tag{4.11}$$

computed per video and reported with bootstrap intervals over the videos, the fitted weights being held fixed. Table 4.1 gives the rule by which U is read. The primary comparisons, fixed in advance, are the raw surprise and the masked-tube error at about 4 fps, both with the trained mask token.

4.5.2. Low-Rank Adaptation of the Encoder

Low-rank adaptation (LoRA) [77] fine-tunes a frozen linear layer $W \in \mathbb{R}^{d_{\mathrm{out}} \times d_{\mathrm{in}}}$ by adding a trainable low-rank update,

$$W'x = Wx + \frac{\alpha}{r} BAx, \qquad A \in \mathbb{R}^{r \times d_{\mathrm{in}}}, \; B \in \mathbb{R}^{d_{\mathrm{out}} \times r}, \tag{4.12}$$

with B initialised to zero so that training starts from the frozen model. LoRA with $r = \alpha = 16$ was applied to the query, key, value and output projections of the last twelve blocks of ViT-L (48 layers, 1.57 million parameters). Each LoRA model was initialised from the frozen model of the same seed and trained for three epochs with learning rates of 10^{-4} for the LoRA weights and 10^{-5} for the decoder.

4.5.3. Audio Fusion

To test whether the soundtrack helps on the audio-visual datasets, an audio fusion module was added to three pyramid levels (P_8, P_{16a}, P_{16b}). For each tubelet, the audio vector a (the mean of the features of its two frames) is embedded as $e = \mathrm{GELU}(W_p\,\mathrm{LN}(a)) \in \mathbb{R}^{256}$. A pyramid level F is then modulated as

$$\tilde{F} = F \odot (1+\gamma) + \beta + W_r\big(F \odot s\big), \qquad s(x) = \mathrm{sigmoid}\left(\frac{\langle W_q F(x),\, W_k e\rangle}{\sqrt{256}}\right), \tag{4.13}$$

where $(\gamma, \beta) = W_f e$ is a feature-wise linear modulation (FiLM) [78], $s(x)$ is an audio-visual correspondence map at each position x, and W_r is a $1\times1\times1$ convolution. Because W_f and W_r are initialised to zero, the module is exactly the visual model at the start of training. As a control, the trained model was also scored with the audio of each test video replaced by that of another test video ("rotated audio"); a model that uses sound should do better with the matched soundtrack.

4.5.4. Static-Image Co-Training

Training on static images helped early video saliency models [8]. To test whether it still helps on top of a video foundation encoder, the 10,000 SALICON training images [36] were added to each training epoch, each image repeated over the 16 input frames, with the same loss.

4.6. Training Procedure

Loss. The loss combines three standard saliency objectives: the Kullback–Leibler divergence of the prediction from the ground-truth density, the linear correlation coefficient, and the normalised scanpath saliency [17]. For a predicted distribution S, a target density Y and fixation counts F on the 64×64 grid,

$$\mathrm{KL}(Y \,\|\, S) = \sum_p Y(p)\Big[\log(Y(p)+\epsilon) - \log S(p)\Big], \tag{4.14}$$

$$\mathrm{CC}(S, Y) = \frac{\sum_p (S(p)-\bar{S})(Y(p)-\bar{Y})}{\big\|S-\bar{S}\big\|_2 \,\big\|Y-\bar{Y}\big\|_2 + \epsilon}, \tag{4.15}$$

$$\mathrm{NSS}(S, F) = \frac{1}{\sum_p F(p)} \sum_p F(p)\,\frac{S(p)-\bar{S}}{\mathrm{std}(S)}, \tag{4.16}$$

with $\epsilon = 10^{-8}$. The training loss over the central frames $\mathcal{O}$ of a clip is

$$\begin{aligned}\mathcal{L} = \frac{1}{|\mathcal{O}|}\sum_{t\in\mathcal{O}}\Big[\mathrm{KL}(Y_t\|S_t) + \big(1 - \mathrm{CC}(S_t, Y_t)\big) - 0.1\,\mathrm{NSS}(S_t, F_t)\Big] \\ + 0.25\sum_{a\in\{8,16\}}\frac{1}{|\mathcal{O}|}\sum_{t\in\mathcal{O}}\Big[\mathrm{KL}(Y_t\|S_t^a) + \big(1 - \mathrm{CC}(S_t^a, Y_t)\big)\Big],\end{aligned} \tag{4.17}$$

where S_t^a are the predictions of the auxiliary heads and the NSS term is averaged over frames with at least one fixation. The KL and CC terms follow the loss of SalFoM [10]; the small NSS term additionally rewards high predicted values at the fixated locations themselves, which the two distribution-based terms do not address directly. The loss is averaged over the clips of a batch.

Optimisation. The decoder was trained with AdamW [124] at a learning rate of 3×10^{-4} and weight decay 0.01; the prior used a learning rate of 10^{-3} without weight decay. The learning rate rose linearly during the first quarter epoch and then followed a cosine decay to zero [125]. Gradients were clipped to a norm of 1.0, and the forward pass ran in bfloat16 mixed precision. Each epoch drew eight random clips from every training video, and the batch size was eight. Clip start positions were drawn uniformly, allowing up to a quarter of the clip to extend beyond the start or end of the video, where frames are repeated. The only data augmentation was a horizontal flip of frames and targets with probability 0.5; the colour jitter mentioned in the original proposal was not implemented.

Schedule and model selection. The proposal foresaw 20 epochs. In the backbone comparison (Section 5.3), however, the holdout CC of every backbone peaked after the second or third epoch and then fell while the training CC kept rising. The schedule for all later runs was therefore fixed, before any of them started, at three epochs, with the checkpoint chosen by holdout CC after each epoch (computed on every fifth frame).

Fine-tuning on other datasets. For UCF Sports, the six audio-visual datasets and Hollywood-2 (Vig gaze), each model was initialised from the DHF1K model and trained for three further epochs with a learning rate of 10^{-4} (3×10^{-4} for the prior). The number of clips per training video was set so that an epoch contained about 1,500 clips (16 clips per video for UCF Sports). Each fold of the audio-visual datasets was trained and selected on its own holdout and its test videos were scored once. These holdouts are small: one video per fold for Coutrot2 and ETMD, two for SumMe, three for AVAD, four for Coutrot1 and six for the single DIEM split, so the selected epoch and blur are noisy on these datasets.

Seeds and ensembles. The main models were trained with three seeds (0, 1 and 2). A three-seed ensemble averages the members' probability maps before post-processing. On the fine-tuning datasets, ensemble member k is fine-tuned from the DHF1K model of seed k. Table 4.3 summarises the settings.

Table 4.3: Training and inference settings of SURGE-D

Setting	Value
Input	16 frames, stride 2, 256×256, ImageNet normalisation
Encoder	V-JEPA 2 ViT-L/16 or ViT-g/16, frozen
Decoder channels	256 (head: 128)
Output	64×64 per frame, central 8 of 16 sampled frames
Loss weights	KL 1, CC 1, NSS 0.1, auxiliary 0.25
Optimiser	AdamW, weight decay 0.01 (prior: 0)
Learning rate (DHF1K)	3×10^{-4} (prior 10^{-3}), warm-up 0.25 epoch, cosine
Learning rate (fine-tuning)	10^{-4} (prior 3×10^{-4})
Epochs, batch size	3, 8
Clips per video and epoch	8 (DHF1K); about 1,500 clips per epoch elsewhere
Augmentation	horizontal flip ($p = 0.5$)
Precision	bfloat16 autocast, gradient clipping at 1.0
Selection	best epoch by holdout CC
Output blur σ	tuned on the holdout over $\{0, 4, 8, 12, 16, 24\}$ pixels
Seeds	0, 1, 2; ensemble = mean of probability maps

4.7. Inference and Post-Processing

A whole video is processed with overlapping windows (Algorithm 4.1). Successive windows advance by $(T/2)\, s = 16$ original frames, so the central sampled positions of consecutive windows tile the video without gaps or overlaps. An original frame that lies between two sampled positions receives the linear interpolation of their maps. The 64×64 map is upsampled bilinearly to the native resolution, smoothed with a Gaussian of standard deviation σ (in native pixels) and normalised to sum to one. The value of σ was chosen on the holdout by CC for each model and each ensemble, and the same map is then used for all metrics, as is common practice in video saliency evaluation. Kümmerer et al. [94] argue instead that each metric should receive the map that is optimal for it under the model's probabilistic prediction; the single CC-tuned map treats all models alike but may understate the scores a model could reach on the other metrics (Section 5.8).

For an ensemble, the members' maps M are averaged before upsampling and blurring. Horizontal-flip test-time augmentation was implemented but not used in any reported result. The causal variant of Section 4.2 is run with one window per output frame, ending at that frame; because this costs one forward pass per frame, it is scored on every tenth frame (Section 5.5.2).

Algorithm 4.1 Whole-video inference with SURGE-D

Require: frames $I_0, \dots, I_{N-1}$; model f; clip length $T = 16$; stride $s = 2$; blur σ
Ensure: saliency maps $\hat{S}_0, \dots, \hat{S}_{N-1}$ at native resolution

$h \leftarrow (T/2)\, s$; $\quad c \leftarrow (T/4)\, s$; $\quad b \leftarrow -c$ ▷ hop, offset of first central frame, window start
while $b + c < N$ **do**
  $n_k \leftarrow \min(\max(b + ks,\, 0),\, N - 1)$ for $k = 0, \dots, T - 1$ ▷ clamp at the video ends
  $(S_0, \dots, S_{T-1}) \leftarrow f(I_{n_0}, \dots, I_{n_{T-1}})$ ▷ one forward pass, 64×64 maps
  for $m \leftarrow b + c$ **to** $\min(b + c + h,\, N) - 1$ **do**
    $q \leftarrow (m - b)/s$; $\quad j \leftarrow \lfloor q \rfloor$; $\quad w \leftarrow q - j$
    $M \leftarrow (1 - w)\, S_j + w\, S_{j+1}$ ▷ linear interpolation in time
    $\hat{S}_m \leftarrow \mathrm{normalise}\Big(\mathrm{Gauss}_\sigma(\mathrm{upsample}_{H \times W}(M))\Big)$
  end for
  $b \leftarrow b + h$
end while

4.8. Evaluation Metrics and Statistical Analysis

4.8.1. Metrics

Six metrics were computed, covering the two families distinguished by Bylinskii et al. [17]: distribution-based metrics, which compare the predicted map P with the ground-truth density Q, and location-based metrics, which evaluate P at the fixated pixels of F. All are computed at the native resolution of the ground truth. Let $z(P)$ denote P standardised to zero mean and unit standard deviation over the frame, $\hat{P} = P / \sum P$ and $\hat{Q} = Q / \sum Q$.

- **Correlation coefficient (CC)**, the Pearson correlation between P and Q, $\mathrm{CC} = \frac{1}{|\Omega|} \sum_p z(P)(p)\, z(Q)(p)$, where Ω is the set of pixels.
- **Normalised scanpath saliency (NSS)** [92], the mean standardised saliency at fixated pixels, $\mathrm{NSS} = \frac{1}{\sum_p F(p)} \sum_p F(p)\, z(P)(p)$.
- **Similarity (SIM)**, the histogram intersection $\sum_p \min(\hat{P}(p), \hat{Q}(p))$.
- **Kullback–Leibler divergence (KLD)**, $\sum_p \hat{Q}(p) \log\Big(\varepsilon + \hat{Q}(p) / (\hat{P}(p) + \varepsilon)\Big)$, with ε the machine epsilon of double precision as in the MIT benchmark code [17]; lower is better.
- **AUC-Judd (AUC-J)** [26], the area under the ROC curve obtained by thresholding P at the saliency values of the fixated pixels, with fixated pixels as positives and all other pixels as negatives. As in the reference implementation, a tiny random jitter breaks ties between quantised values.
- **Shuffled AUC (sAUC)** [25], [93], which takes its negatives from the fixations of other videos instead of from all pixels and therefore does not reward a centre bias. For each video, the fixations of one random frame of each of ten other videos

of the same test set are pooled; for each frame, 100 random draws of as many negatives as there are fixations are scored with thresholds in steps of 0.1 and the areas are averaged.

Each metric is averaged over the frames of a video and then over the videos of the test set; on the audio-visual datasets, the dataset score is the mean over its folds. The shuffled AUC depends on the sampling of negatives, so its values are compared only within this thesis.

Validation of the evaluator. The metrics were implemented independently of any existing saliency code. Before any model of this thesis was scored, the evaluator was checked against reference code by re-scoring a fixed set of saliency maps (Section 5.2).

4.8.2. Statistical Analysis

When two models are scored on the same test videos, their difference is assessed with a paired bootstrap over videos [126]. For a metric m and videos $v = 1, \ldots, V$, let $d_v = m_A(v) - m_B(v)$. The mean difference $\bar{d}$ is reported with a 95 % percentile interval obtained from 10,000 resamples of the V videos with replacement. A model is said to be better on a metric only if the interval excludes zero. The bootstrap treats videos as the independent units and therefore accounts for the strong correlation between frames of the same video. On the audio-visual datasets, the test videos of all folds are pooled for this purpose, which is possible because each video is tested exactly once. In the component ablations with three seeds (Table 4.1), each variant is compared on the holdout with the reference model of the same seed. The per-video differences are averaged over videos and seeds, and the intervals come from a two-level bootstrap that first resamples the three seeds and then the videos, drawing the same videos for every seed because the comparison is paired. The learning curve over the number of training videos (Section 5.6) is analysed in the same way. There, the training subsets are drawn separately for each seed, so the resampling of seeds also covers the choice of subset. With three seeds, the seed level of this bootstrap is coarse, and its intervals are approximate.

Published results of other methods were computed with other evaluation code and, sometimes, other protocol details. They cannot enter a paired test. They are reported side by side with the results of this thesis, and margins below 0.01 CC are treated as ties (Table 4.1).

4.9. Ethical Considerations

The thesis uses only eye-tracking data that were collected and released by other research groups under their own ethical procedures; no human participants were recruited, and

no personal data beyond the anonymised gaze recordings and the public video content of these datasets were processed. The licences of the data and models were respected: the Vig gaze data and the V-JEPA 2 checkpoints carry permissive licences, and the benchmark videos were used for research only, as their distributors require.

Saliency prediction has dual-use potential. Models of where people look can improve video compression, accessibility and the evaluation of user interfaces, but they can also be used to design content that captures attention against the viewer's interest. The models of this thesis predict average gaze over groups of observers and contain no information about individuals; the work does not, however, remove the broader risk, and the code accompanying the thesis is intended for research use.

Finally, the experiments consumed electricity on a single workstation graphics card. The design choices that keep the encoder frozen and amortise its cost over many output frames reduce this footprint compared with fully fine-tuned alternatives, and every result was obtained with a fixed number of runs decided in advance rather than by open-ended search.

CHAPTER FIVE

EXPERIMENTS, RESULTS AND DISCUSSION

This chapter reports the experiments in the order in which they were run, because later experiments were planned on the basis of earlier outcomes. After the setup (Section 5.1) and the validation of the evaluation code (Section 5.2), the two preliminary gates are reported: the choice of the video encoder (Section 5.3) and the test of predictive surprise as a fixation cue (Section 5.4). The main results on DHF1K, the transfer to eight further test sets, qualitative examples and the computational cost follow in Section 5.5, and the ablations in Section 5.6. The chapter ends with a discussion organised by research question (Section 5.7) and an assessment of the threats to validity (Section 5.8). Unless stated otherwise, every number in this chapter was produced by the code and evaluator described in Chapter 4, and every test set was scored once per reported model.

5.1. Experimental Setup

All experiments ran on a single workstation (Table 5.1), one job at a time. A monitoring process checked the GPU temperature, hardware events, disk space and memory every minute and could stop training cleanly; it was never triggered.

Table 5.1: Experimental environment

Item	Specification
Graphics Processing Unit (GPU)	NVIDIA RTX A5000, 24 GB
CPU / RAM	Intel Xeon w7-2475X (20 cores) / 256 GB
Operating system	Windows 11 Pro
Language	Python 3.12
Framework	PyTorch 2.14 with CUDA 13.0 [15]
Model loading	Transformers 5.22 [16]
Other libraries	NumPy [127], OpenCV, SciPy

5.1.1. Implementation Details

The settings of Table 4.3 were used throughout. Training one three-epoch DHF1K model took about 15 minutes of training time with the ViT-L encoder (4.7 minutes per epoch of 4,320 clips) and about 31 minutes with ViT-g (9.9 minutes per epoch), plus the holdout evaluation after each epoch. ViT-g fitted in 10 GB of GPU memory at a batch size of eight. In total, the experiments of this chapter comprise 40 DHF1K training runs and 125 fine-tuning runs on the other datasets (including the 16 runs with audio).

The DHF1K runs are 13 for the gates, the main program and its first ablations, 7 for the ViT-g, 600-video and 32-frame experiments, 14 for the first follow-up (eight further ablation seeds and six causal models) and 6 for the learning curve over the number of training videos. Holdout lists, configuration files and the log of every decision are part of the code accompanying the thesis.

5.2. Validation of the Evaluator

Because the metrics were implemented from scratch, the evaluator was validated before any SURGE-D model was scored (Table 4.1). The reference maps were the DHF1K validation predictions of SalFoM [10], produced on the same machine with the authors' released model code and weights, for which per-frame scores computed with the evaluation code released by the TMFI-Net authors [107] were available. These maps had been produced for an earlier project on the same machine, by a wrapper program that is not part of the code of this thesis. The wrapper follows the test procedure of TMFI-Net, whose scripts the SalFoM authors state they used for training and testing[1]. The map of frame t is predicted from the frames $t-15$ to t, the first frame being repeated before the start of a video. Frames are resized to 224×224 pixels with bilinear interpolation and normalised with the ImageNet statistics; the output is resized to the frame size, smoothed with an 11×11 Gaussian kernel and stored as an 8-bit image after min–max scaling. The wrapper ran SalFoM in bfloat16; on 16 test clips, the bfloat16 and single-precision maps agreed with a CC of at least 0.9999.

The evaluator reproduced the reference scores: CC 0.5639 against 0.5642 ($\Delta = -0.0003$) and NSS 3.3264 against 3.3268 ($\Delta = -0.0004$), well within the tolerance of the rule. The SIM, KLD and AUC-J values of the reference code use other normalisations and could not be compared directly. Against a second evaluator that follows the MIT benchmark conventions, written independently for the same earlier project and used here only as a reference, CC, NSS, SIM and KLD agreed per video to within 5×10^{-5} and AUC-J agreed in the mean; sAUC differed by 0.0023, because its negatives are drawn at random. Two deviations from the reference procedures were found and corrected during this check, before any model was scored: the first AUC-J implementation lacked the tie-breaking jitter of the reference code (which lowered AUC-J by 0.010 on quantised maps), and the first shuffled AUC pooled all negatives instead of drawing them per split. With the corrected evaluator, SalFoM's validation maps score CC 0.5639, NSS 3.326, SIM 0.4130, KLD 1.297, AUC-J 0.9243 and sAUC 0.7406. These are the SalFoM numbers used in all comparisons below. They are slightly lower than SalFoM's published scores because they are computed on the validation set rather than on the official test set, on which the authors report CC 0.569 and NSS 3.353 [10]. How the released weights

[1] README of the official SalFoM repository.

were selected is not documented. The training script of TMFI-Net keeps the checkpoint with the highest CC on the DHF1K validation set; if SalFoM's weights were chosen in the same way, its validation scores are optimistic, which would favour SalFoM in the comparisons of this chapter.

5.3. Selection of the Video Encoder (Gate G1)

The proposal had chosen V-JEPA 2.1 [14], which extends the V-JEPA 2 objective with a loss on the visible tokens and with self-supervision at intermediate layers, specifically to obtain denser features. Its ViT-L model is distilled from a two-billion-parameter teacher and was cooled down at 384×384 pixels [14][2]. Gate G1 compared three ViT-L encoders under identical conditions: V-JEPA 2, V-JEPA 2.1 and UMT-L [70], the encoder family used by SalFoM. Each was frozen and trained with the same decoder for ten epochs with seed 0, all V-JEPA models at 256×256 pixels. UMT-L received 224×224 input because its position table is fixed to that size and to 16 frames.

Table 5.2: Gate G1: frozen ViT-L encoders with the same decoder, internal DHF1K holdout (60 videos, every fifth frame, no blur), best of ten epochs, seed 0

Encoder	Best epoch	CC↑	NSS↑	SIM↑	KLD↓
V-JEPA 2 ViT-L (256^2)	3	**0.5782**	**3.314**	**0.4432**	**1.218**
V-JEPA 2.1 ViT-L (256^2)	2	0.5692	3.295	0.4377	1.245
UMT-L (224^2) [70]	2	0.5464	3.101	0.4157	1.323

Table 5.2 shows that V-JEPA 2 led on all four metrics. V-JEPA 2.1 trailed by 0.009 CC, outside the margin of 0.005 within which the proposal's choice would have been kept, so V-JEPA 2 was adopted for all later experiments. UMT-L trailed by 0.032 CC. With one seed per encoder, part of the gap between the two V-JEPA variants may be seed noise; the gap to UMT-L is several times larger than the seed-to-seed spread observed later (0.003 CC). The ranking of the two V-JEPA variants runs against the expectation that denser features help a dense task. Two explanations are compatible with the data and were not separated: the decoder reads four intermediate layers, which may already supply the local detail that the last layer of V-JEPA 2 lacks, and the V-JEPA 2.1 model was used at a lower resolution than the one at which it was cooled down.

The learning curves in Figure 5.1 were the more consequential outcome of this gate. For every encoder, holdout CC peaked in the second or third epoch and then declined steadily, while training CC kept rising. A frozen encoder with a 7-million-parameter decoder thus overfits the 540 training videos after a few passes. On the basis of these

[2]Checkpoint `vjepa2_1_vitl_dist_vitG_384` from the official repository.

curves, and before any further run, the schedule of the main program was shortened to three epochs (Section 4.6).

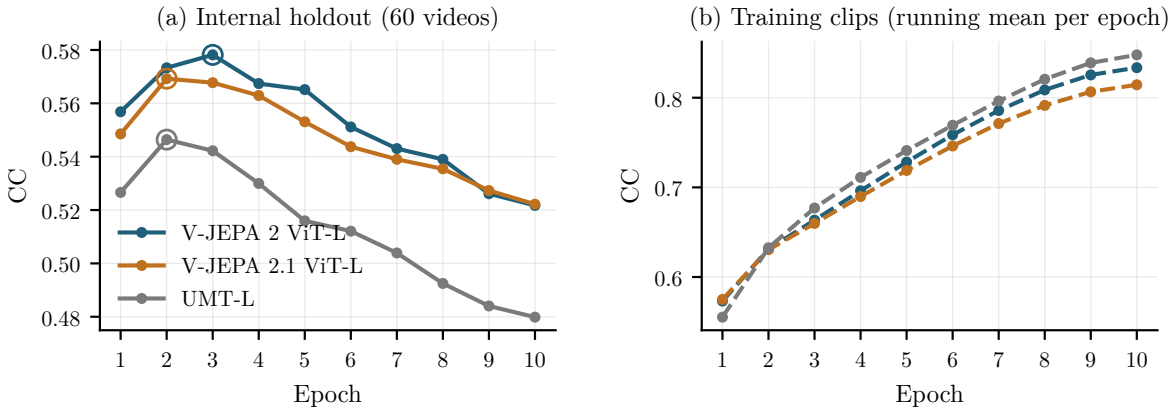


Figure 5.1: Learning curves of gate G1. (a) Holdout CC after each epoch; circles mark the best epoch. (b) Mean CC on the training clips of each epoch. Holdout accuracy peaks early and then declines while training accuracy keeps rising, for all three encoders

5.4. Predictive Surprise as a Fixation Cue (Gate G2)

Gate G2 tested the hypothesis that motivated the original proposal: that the error of V-JEPA 2's pretrained predictor indicates where people look (Section 4.5.1). No training is involved; the prediction-error maps and the reference maps were scored directly against the fixations of the 60 holdout videos, on 1,720 frames at about 4 sampled frames per second and on 7,904 frames at about 15. The gate was scored twice. Its first run, on which the decision about the surprise module was taken, used the untrained mask token that the model-loading library selects by default. After this error had been found, every map was recomputed with the trained token and the gate was scored again by its own rule (Section 4.5.1). Table 5.3 gives both readouts, together with the masked-tube error, which was added in the second follow-up and is not part of the gate.

The outcome does not depend on the mask token (Table 5.3). With the trained token, neither surprise variant beat token atypicality or frame difference at either frame rate, and both were *anti*-correlated with fixations. The NSS of the raw error is -0.665 [-0.763, -0.566] at 4 fps and -0.877 [-0.958, -0.796] at 15 fps, and that of the contextual variant is -0.239 [-0.373, -0.107] and -0.660 [-0.770, -0.548]. A negative NSS means that observers looked on average at regions that the predictor found *easier* than average to predict. The information gain of the gate, which compares an equal mixture of a map and the centre prior with the prior alone, was negative for every map. The gate therefore fails with the trained token as it had failed in its first run, and the decision not to train the surprise module stands. In the first run the raw error was also negative at both rates (-0.190 [-0.332, -0.045] and -0.502 [-0.620,

Table 5.3: Gate G2: zero-shot fixation prediction by prediction-error maps and reference maps on the DHF1K holdout (60 videos), at about 4 and 15 sampled frames per second. The masked-tube error is not part of the gate. Intervals are given in Tables A.9 and A.10

	NSS↑		CC↑	
Map	4 fps	15 fps	4 fps	15 fps
Prediction error, trained mask token				
Raw surprise, Equation (4.8)	−0.665	−0.877	−0.129	−0.168
Contextual surprise (true − other past)	−0.239	−0.660	−0.045	−0.127
Masked-tube error	−0.958	−0.942	−0.186	−0.182
Prediction error, untrained mask token (first run of the gate)				
Raw surprise	−0.190	−0.502	−0.043	−0.099
Contextual surprise	0.116	−0.366	0.018	−0.072
Masked-tube error	−0.858	−0.819	−0.170	−0.161
Reference maps				
Token atypicality	0.454	0.650	0.081	0.123
Frame difference	0.810	0.803	0.149	0.148
Empirical centre prior	1.464	1.479	0.289	0.292

−0.385]); the contextual variant was negative at 15 fps (−0.366 [−0.505, −0.225]) and could not be distinguished from zero at 4 fps (+0.116 [−0.043, +0.275]). For both tokens, every paired difference between a prediction-error map and frame difference or token atypicality lies below zero at both rates (Tables A.9 and A.10).

Correcting the mask token made the anti-correlation stronger, not weaker: every score of a prediction-error map in Table 5.3 is lower with the trained token. The masked-tube error, which follows the mask geometry of pre-training (Section 4.5.1), is the most negative map of the table, with an NSS of −0.958 [−1.046, −0.873] at 4 fps and −0.942 [−1.026, −0.858] at 15 fps for the trained token.

These are comparisons of stand-alone maps, and a stand-alone score can be negative merely because a map is low at the centre of the frame, where fixations concentrate. The second analysis therefore asks whether a prediction-error map adds information when it is combined with the centre prior, frame difference and token atypicality (Equation (4.11)). Table 5.4 reports it.

Both primary comparisons are read as incremental information by the rule of Table 4.1. At 4 fps and with the trained token, the raw surprise adds +0.0693 [+0.0346, +0.1065] bits per fixation to the three simple cues, and the masked-tube error adds +0.0619 [+0.0320, +0.0926]. In both, the fitted weight of the error map is negative in both halves of the cross-fitting: a cell whose error lies one standard deviation above the mean of its frame receives about 30 % less probability, other cues being equal. The prediction error therefore carries information about fixations beyond the centre bias, motion and atypicality, but of the opposite sign to the hypothesis: observers look less

Table 5.4: Incremental information U of the prediction-error maps in bits per fixation, with 95 % bootstrap intervals over the 60 holdout videos, and the fitted weight of the error map in the two halves of the cross-fitting. An interval above zero is read as incremental information (Table 4.1); the two primary comparisons are marked *. The last block gives, for scale, the loss when one simple cue is removed from the model of the first row, and the gain of the two simple cues together over the centre prior

	4 fps		**15 fps**	
Error map	U	**Weight**	U	**Weight**
Trained mask token				
Raw surprise*	+0.0693 [+0.0346, +0.1065]	−0.34 / −0.40	+0.0902 [+0.0559, +0.1263]	−0.36 / −0.45
Contextual surprise	+0.0125 [−0.0035, +0.0296]	−0.13 / −0.16	+0.0421 [+0.0184, +0.0675]	−0.24 / −0.27
Masked-tube error*	+0.0619 [+0.0320, +0.0926]	−0.34 / −0.41	+0.0561 [+0.0280, +0.0844]	−0.31 / −0.35
Untrained mask token				
Raw surprise	+0.0072 [−0.0027, +0.0176]	−0.13 / −0.12	+0.0470 [+0.0209, +0.0739]	−0.28 / −0.34
Contextual surprise	−0.0041 [−0.0103, +0.0016]	−0.02 / +0.08	+0.0197 [−0.0003, +0.0402]	−0.15 / −0.17
Masked-tube error	+0.0524 [+0.0220, +0.0848]	−0.35 / −0.43	+0.0447 [+0.0184, +0.0728]	−0.31 / −0.36
For scale				
Without frame difference	+0.1190 [+0.0769, +0.1675]		+0.0935 [+0.0575, +0.1335]	
Without token atypicality	+0.1845 [+0.1013, +0.2730]		+0.1552 [+0.0867, +0.2282]	
Simple cues over the prior	+0.3300 [+0.2210, +0.4479]		+0.3137 [+0.2149, +0.4211]	

often where the error is high. The amount is modest. The two simple cues together gain 0.3300 [0.2210, 0.4479] bits over the centre prior, nearly five times as much, and removing frame difference or token atypicality from the model that contains the raw surprise costs 0.1190 and 0.1845 bits.

The other comparisons of Table 5.4 carry no decision and received no correction for multiplicity. They agree in sign: wherever the interval of U lies above zero, the weight is negative in both halves. With the trained token, all three maps add information at 15 fps, and the contextual variant is inconclusive at 4 fps. With the untrained token, the masked-tube error adds information at both rates and the raw surprise at 15 fps; the raw surprise at 4 fps and the contextual variant at 15 fps are inconclusive, and the contextual variant at 4 fps is negligible (its 90 % interval, $[-0.0093, +0.0007]$, lies within ± 0.01 bits).

Two explanations of the first negative result can now be set aside. The first is the untrained mask token, since the trained token gives lower scores. The second was offered in the first draft of this thesis. The masks of V-JEPA 2's pre-training span the whole clip in time[3], so the predictor was trained to fill in masked scene content and never to extrapolate into the future. Predicting the second half of a clip from the first is a query that it never saw, and its error was therefore thought to mix surprise with the cost of an unfamiliar task. The masked-tube error poses a query of the kind seen in pre-training, and it is anti-correlated with fixations at least as strongly as the forecasting error. The form of the query thus does not explain the sign, with the reservation that this readout hides a smaller share of the clip than pre-training did.

A preliminary diagnostic, run before the gate on eleven clips without gaze data (ten Kinetics validation clips [50] and one animated film trailer), gives a partial picture of what the error measures. It used the untrained mask token and was not repeated with the trained one. When the true past of a clip was replaced by another clip's past, the mean prediction error rose in every clip (from 0.644 to 0.731 at about 4 fps), so the predictor does use the context. Yet the spatial pattern of the error correlated at $r \approx 0.60$ with token atypicality and at $r \approx 0.62$ with the error map computed from the *wrong* past, and negatively with a central Gaussian ($r \approx -0.11$). On these clips, much of the error thus reflected how unusual each target region is, whatever the past, and it was lower at the centre of the frame. The incremental analysis shows that this is not the whole account: with the centre prior and token atypicality in the model, the error still carries information about fixations, so it is more than a proxy for atypicality or for the centre bias as the model of Equation (4.9) represents them. What the error measures that draws gaze away was not established. Within these limits, the evidence of this section does not support reading the error of V-JEPA 2's predictor as

[3] `temporal_scale: [1.0, 1.0]` in the released pre-training configuration files of the official repository.

the expectation-violation signal that the Bayesian surprise account [13] associates with attended regions. Whether a predictor trained to forecast, or a surprise measured as a change of belief rather than as an error in feature space, would behave differently was not tested.

5.5. Results

5.5.1. Quantitative Results on DHF1K

The main program trained the frozen-encoder model (labelled B1 in the research log) with three seeds. The LoRA variant lost to it in all three seeds (Section 5.6), so the frozen model became the main model. After the ViT-g encoder passed its own gate, the same recipe was repeated with ViT-g. Table 5.5 reports all four configurations on the DHF1K validation set, next to SalFoM scored by the same evaluator and to published validation results of other methods.

Table 5.5: Results on the DHF1K validation set (100 videos, 60,273 frames). The upper block was scored by the evaluator of this thesis; the lower block is copied from the literature and was computed with other evaluation code. The second SalFoM row applies the blur that is best for SalFoM on this set, an upper bound for the baseline (Section 5.5.1). Best value per column in bold

Model	CC↑	NSS↑	SIM↑	KLD↓	AUC-J↑	sAUC↑
SalFoM [10] (released weights)	0.5639	3.326	0.4130	1.297	0.9243	0.7406
SalFoM, best blur ($\sigma = 16$)	0.5741	3.265	0.4009	1.280	0.9277	0.7557
SURGE-D ViT-L, single ($\sigma = 12$)	0.5836	3.349	0.4492	1.222	0.9283	0.7636
SURGE-D ViT-L, 3-seed ensemble ($\sigma = 8$)	0.5910	3.418	0.4486	1.187	0.9302	0.7614
SURGE-D ViT-g, single ($\sigma = 8$)	0.6003	3.457	0.4531	1.170	0.9308	**0.7709**
SURGE-D ViT-g, 3-seed ensemble ($\sigma = 8$)	**0.6029**	**3.486**	**0.4574**	**1.160**	**0.9314**	0.7680
Published, other evaluators (as reported in [45])						
TMFI-Net [107]	0.554	3.201	0.428	–	0.924	–
THTD-Net [108]	0.553	3.188	0.425	–	0.924	–
ViNet-E [45]	0.549	3.134	0.409	–	0.922	–
STSANet [105]	0.539	3.082	0.411	–	0.920	–
DiffSal [91]	0.533	3.066	0.405	–	0.918	–

Comparison with SalFoM. SalFoM is the closest prior work: it was the first video saliency model built on a video foundation encoder, it reported the best DHF1K test results among the methods it compared with [10], and its released weights allowed it to be scored by the same evaluator on the same videos. Table 5.6 gives the paired bootstrap intervals. The ViT-g model is better than SalFoM on all six metrics as a single model and as an ensemble, with every interval excluding zero; the CC gain of the single model is +0.036 [+0.026, +0.048]. The ViT-L ensemble is also better on all six metrics, and the single ViT-L model on five, its NSS difference not being significant. The per-video differences in Figure 5.2 show that the gain is broad rather than driven

by a few videos: the ViT-g ensemble has the higher CC on 79 of the 100 videos. These intervals are against SalFoM as released; the checks below test them against a stronger version of the baseline.

Table 5.6: Paired bootstrap differences on DHF1K validation (100 videos, 10,000 resamples), with 95 % intervals; differences whose interval excludes zero are in bold. The intervals for all six metrics are given in Table A.1

Comparison	**ΔCC** [95 % CI]	**ΔNSS** [95 % CI]
ViT-L single − SalFoM	**+0.020** [+0.010, +0.031]	+0.023 [−0.050, +0.097]
ViT-L ensemble − SalFoM	**+0.027** [+0.017, +0.038]	**+0.091** [+0.022, +0.164]
ViT-g single − SalFoM	**+0.036** [+0.026, +0.048]	**+0.130** [+0.060, +0.205]
ViT-g ensemble − SalFoM	**+0.039** [+0.029, +0.050]	**+0.160** [+0.089, +0.233]
ViT-g single − ViT-L single	**+0.017** [+0.012, +0.022]	**+0.108** [+0.075, +0.142]
ViT-g ensemble − ViT-L ensemble	**+0.012** [+0.008, +0.016]	**+0.069** [+0.041, +0.097]

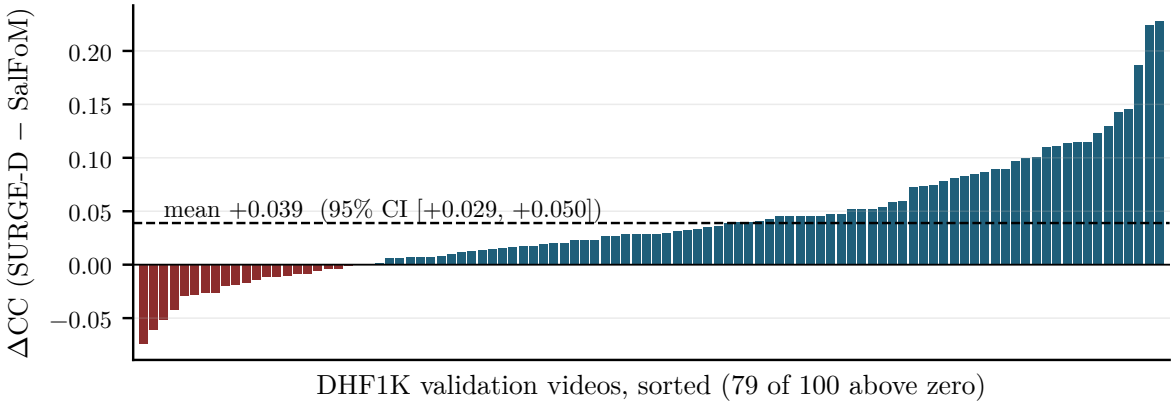


Figure 5.2: Per-video CC difference between the SURGE-D ViT-g ensemble and SalFoM on the 100 DHF1K validation videos, sorted. The dashed line is the mean difference; the interval is the paired bootstrap 95 % interval

Two qualifications apply. First, SURGE-D uses frames on both sides of the predicted frame (Section 4.2), whereas SalFoM predicts the last frame of its input clip; part of the difference could therefore come from access to future context, which SalFoM does not use. The causal control of Section 5.5.2 shows that it does not. Second, the comparison involves one re-scored baseline. The published validation scores in the lower block of Table 5.5 were computed with other code, which agrees with the present evaluator to within 0.0003 CC on the reference maps (Section 5.2); with that caveat, SURGE-D's validation CC is 0.046–0.049 above the highest of them. Reported validation scores also differ between papers: UniST [48], for instance, lists TMFI-Net at 0.524 rather than 0.554 and reports 0.541 for itself. Several other peer-reviewed models published since 2024 [104], [109], [110] report DHF1K results whose tables could not be accessed or that refer to the test set; test-set scores cannot be compared with the validation scores of this thesis.

Checks of the comparison. After the first draft of this thesis, further checks were specified in the research log and then run (rules marked † in Table 4.1; Table 5.7).

The first check concerns the training seed. Seeds 1 and 2 had been scored only on the holdout. Scored singly on the validation set, every seed of both encoders had a higher CC than SalFoM, with intervals above zero. The reported seed 0 was an average ViT-L seed (CC 0.5836, against a three-seed mean of 0.5838 with a standard deviation of 0.0034) and the best of the three ViT-g seeds (0.6003, against 0.5969 with a standard deviation of 0.0030).

The second check concerns the output blur. SURGE-D's blur is tuned on the holdout, whereas SalFoM's maps had been scored as released. Blurring SalFoM's maps with the same grid raised its CC from 0.564 to 0.574 at $\sigma = 16$ pixels. This value was chosen on the validation set itself, so it favours SalFoM. Against this stronger baseline, every ViT-g model and the ViT-L ensemble remain better: the ViT-g ensemble by +0.029 [+0.019, +0.039], and the ViT-g single model and ensemble on all six metrics (Table A.3). Single ViT-L models do not: for seeds 0 and 1 the interval includes zero (seed 0: +0.0095 [−0.0008, +0.0203]). At the encoder size of SalFoM, therefore, only the ensemble is shown to be more accurate.

The third check concerns the encoder. With the same decoder and the same frozen protocol, the UMT-L encoder of gate G1 reached a validation CC of 0.545. That is below SalFoM (−0.019 [−0.029, −0.010]) and 0.039 below SURGE-D with the V-JEPA 2 ViT-L encoder. A frozen encoder of SalFoM's own family therefore does not reach SalFoM with this decoder, and the advantage of SURGE-D comes from the V-JEPA 2 encoder rather than from freezing the encoder as such. The UMT-L model is a single seed, trained for ten epochs at 224 × 224 pixels.

Table 5.7: Checks of the comparison with SalFoM on DHF1K validation: CC and paired bootstrap differences (100 videos, 95 % intervals). "Best blur": SalFoM's maps blurred with $\sigma = 16$, the value of the grid that is best for SalFoM on this set, an upper bound for the baseline. Differences whose interval excludes zero are in bold

Model	CC	ΔCC, SalFoM released	ΔCC, SalFoM best blur
SURGE-D ViT-L, seed 0 (reported)	0.5836	**+0.020** [+0.010, +0.031]	+0.010 [−0.001, +0.020]
SURGE-D ViT-L, seed 1	0.5805	**+0.017** [+0.007, +0.027]	+0.006 [−0.003, +0.016]
SURGE-D ViT-L, seed 2	0.5872	**+0.023** [+0.013, +0.034]	**+0.013** [+0.003, +0.024]
SURGE-D ViT-L, ensemble	0.5910	**+0.027** [+0.017, +0.038]	**+0.017** [+0.007, +0.027]
SURGE-D ViT-g, seed 0 (reported)	0.6003	**+0.036** [+0.026, +0.048]	**+0.026** [+0.016, +0.037]
SURGE-D ViT-g, seed 1	0.5948	**+0.031** [+0.021, +0.042]	**+0.021** [+0.011, +0.031]
SURGE-D ViT-g, seed 2	0.5957	**+0.032** [+0.021, +0.043]	**+0.022** [+0.011, +0.033]
SURGE-D ViT-g, ensemble	0.6029	**+0.039** [+0.029, +0.050]	**+0.029** [+0.019, +0.039]
Frozen UMT-L, SURGE-D decoder, seed 0	0.5450	**−0.019** [−0.029, −0.010]	**−0.029** [−0.039, −0.020]

DHF1K test set. Maps for the 300 official test videos were produced with the ViT-g three-seed ensemble and checked for format by passing five validation videos through the same export, which reproduced their validation scores to within 10^{-4} CC. The scores on the test set are computed by the dataset authors on request and were not available when this thesis was written, so no test-set result is claimed.

5.5.2. Causal Prediction on DHF1K

The comparison above leaves one difference between the two models: SURGE-D sees frames after the predicted one and SalFoM does not. The causal variant of Section 4.2 removes this difference. It was trained with the ViT-g recipe for three seeds and tested under the causal-control rule of Table 4.1, which was written before it was run. After the component ablations, the same procedure was repeated with the ViT-L encoder, the size of SalFoM's. For all six models the best holdout epoch was the second of three, and for both encoders the holdout selected a blur of $\sigma = 8$ for the seed-0 model, which seeds 1 and 2 also use, and for the ensemble. Because a causal model needs one forward pass per output frame, the validation set was scored once on every tenth frame of each video (6,076 frames), and SalFoM was scored on the same frames.

Table 5.8 gives the results. The causal ViT-g ensemble reached a CC of 0.602, and its difference from SalFoM with its best blur was +0.029 [+0.019, +0.040]. Under the rule, SURGE-D is therefore more accurate than SalFoM under an identical evaluation protocol: the same videos and frames, the same evaluator, no frame after the predicted one for either model, and the baseline's blur set to the value that is best for it on the test set. Each causal ViT-g single model was also more accurate than SalFoM with its best blur, by 0.018 to 0.025 CC, and the ensemble was better on all six metrics (Table A.5).

With the ViT-L encoder, the causal ensemble reached a CC of 0.590, and its difference from SalFoM with its best blur was +0.017 [+0.007, +0.028], so the conclusion of the rule also holds at the encoder size of SalFoM. This ensemble was better on all six metrics, on sAUC only narrowly. As in the offline comparison (Table 5.7), the single ViT-L models were not reliably better: against SalFoM with its best blur, the interval of seed 0 lies above zero (+0.011 [+0.001, +0.022]), but those of seeds 1 and 2 include it. A single causal ViT-L model needs about two thirds of the computation per output frame of SalFoM's encoder (Section 5.5.5). At that cost, therefore, an advantage over SalFoM is not established; the ensemble, which shows one, needs about twice SalFoM's encoder computation per output frame.

The offline models were rescored on the same frames to measure what the later frames contribute, and with neither encoder did they add measurable accuracy. The causal ensemble's CC differed from the offline ensemble's by +0.0016 [−0.0015, +0.0046] with ViT-g and by +0.0012 [−0.0022, +0.0047] with ViT-L, and that of the causal

seed-0 model from its offline counterpart by +0.0003 [−0.0044, +0.0049] and +0.0030 [−0.0014, +0.0077], although each causal training clip supervised one frame instead of eight. The other metrics moved by small amounts in both directions (Table A.6). The advantage of SURGE-D over SalFoM therefore does not come from access to future frames. What the causal model gives up is efficiency, because it needs one forward pass per output frame (Section 5.5.5).

Table 5.8: Causal prediction on DHF1K validation, scored on every tenth frame (100 videos, 6,076 frames). The causal models and SalFoM use only the predicted frame and earlier frames; the offline models, for reference, also use later frames. Paired bootstrap differences in CC with 95 % intervals (10,000 resamples); differences whose interval excludes zero are in bold. All six metrics are compared in Tables A.5 and A.6

Model	CC	NSS	ΔCC, SalFoM released	ΔCC, SalFoM best blur
SalFoM, released	0.5622	3.302	–	–
SalFoM, best blur ($\sigma = 16$)	0.5729	3.244	–	–
Causal SURGE-D ViT-g ($\sigma = 8$)				
Seed 0	0.5978	3.425	**+0.036** [+0.024, +0.047]	**+0.025** [+0.014, +0.036]
Seed 1	0.5913	3.408	**+0.029** [+0.018, +0.040]	**+0.018** [+0.008, +0.029]
Seed 2	0.5925	3.414	**+0.030** [+0.019, +0.042]	**+0.020** [+0.008, +0.031]
Three-seed ensemble	0.6021	3.462	**+0.040** [+0.029, +0.051]	**+0.029** [+0.019, +0.040]
Offline SURGE-D ViT-g on the same frames ($\sigma = 8$)				
Seed 0	0.5976	3.427	**+0.035** [+0.025, +0.047]	**+0.025** [+0.014, +0.036]
Three-seed ensemble	0.6005	3.457	**+0.038** [+0.028, +0.049]	**+0.028** [+0.017, +0.038]
Causal SURGE-D ViT-L ($\sigma = 8$)				
Seed 0	0.5842	3.346	**+0.022** [+0.011, +0.033]	**+0.011** [+0.001, +0.022]
Seed 1	0.5767	3.312	**+0.014** [+0.004, +0.025]	+0.004 [−0.006, +0.014]
Seed 2	0.5792	3.346	**+0.017** [+0.006, +0.029]	+0.006 [−0.005, +0.018]
Three-seed ensemble	0.5900	3.390	**+0.028** [+0.018, +0.038]	**+0.017** [+0.007, +0.028]
Offline SURGE-D ViT-L on the same frames ($\sigma = 12$ for seed 0, $\sigma = 8$ for the ensemble)				
Seed 0	0.5812	3.321	**+0.019** [+0.009, +0.030]	+0.008 [−0.002, +0.019]
Three-seed ensemble	0.5888	3.390	**+0.027** [+0.016, +0.037]	**+0.016** [+0.006, +0.026]

5.5.3. Generalisation to Other Datasets

UCF Sports. UCF Sports differs from DHF1K in two ways: its gaze was recorded during an action-recognition task [82], and it is small (103 training videos). Without any UCF training, the DHF1K model reached a CC of 0.608 (Table 5.9), well below the fine-tuned models, which confirms that the task-driven gaze of this dataset differs from free viewing. After fine-tuning, the ViT-g ensemble reached CC 0.735 and the best NSS, SIM and AUC-J of the SURGE-D models.

On UCF Sports the encoder made a smaller difference than on DHF1K: the ViT-g ensemble was not significantly better than the ViT-L ensemble (ΔCC +0.0019 [−0.0071, +0.0125]), whereas the three-seed ViT-L ensemble was significantly better than its single member (+0.0080 [+0.0039, +0.0123]). Against published results, SURGE-D clearly exceeds SalFoM (by about 0.10 CC) and TMFI-Net, and it ties ViNet-A (+0.001).

Table 5.9: Results on the UCF Sports test set (47 videos, 2,929 frames). Published results were computed with other evaluation code

Model	CC↑	NSS↑	SIM↑	KLD↓	AUC-J↑	sAUC↑
SURGE-D ViT-L, zero-shot	0.6077	3.128	0.5033	1.113	0.9223	0.7453
SURGE-D ViT-L, fine-tuned	0.7252	3.916	0.5799	0.773	0.9406	0.8089
SURGE-D ViT-L, 3-seed ensemble	0.7331	4.028	0.5833	0.749	0.9430	0.8153
SURGE-D ViT-g, fine-tuned	0.7305	4.061	0.5831	0.747	0.9442	0.8145
SURGE-D ViT-g, 3-seed ensemble	0.7350	4.085	0.5857	**0.736**	**0.9449**	**0.8173**
Published ([45]; SalFoM from [10])						
ViNet-E [45]	**0.744**	**4.156**	**0.587**	–	0.941	–
ViNet-A [45]	0.734	4.108	0.586	–	0.940	–
THTD-Net [108]	0.711	3.840	0.565	–	0.933	–
TMFI-Net [107]	0.707	3.863	0.565	–	0.936	–
SalFoM [10]	0.631	3.543	0.516	–	0.928	–

Its CC is 0.009 below that of ViNet-E, which is a tie under the rule of Table 4.1; its NSS and SIM are also below those of ViNet-E, and its AUC-J is above. SURGE-D is therefore competitive on UCF Sports but not ahead of the best published model. Both ViNet variants use encoders pretrained for action recognition or action localisation [45], which may suit the task-driven gaze of this dataset.

Audio-visual datasets. Table 5.10 reports the six datasets of the STAViS benchmark [90], as fold means. The DHF1K model already transferred reasonably without any training (zero-shot row), and fine-tuning improved every dataset, most of all Coutrot2 (+0.17 CC), whose meetings differ most from DHF1K. With ViT-g, the ensemble was significantly better than the ViT-L ensemble on DIEM, Coutrot1, ETMD and AVAD (paired bootstrap over the pooled test videos), not significantly different on Coutrot2, and tied on SumMe.

Against the best peer-reviewed results that could be verified, the CC of the ViT-g ensemble is higher by more than the tie margin on four datasets (DIEM, Coutrot1, ETMD and SumMe, by 0.018 to 0.028 over the best result on each). It ties on AVAD, where it is +0.002 above the best audio-visual model, and it is lower on Coutrot2. The other metrics broadly agree with CC: the ensemble's NSS (DIEM 2.88, Coutrot1 3.30, ETMD 3.69, AVAD 4.31, SumMe 3.26) is above the published values of both reference models on these five datasets, while SIM is lower than DiffSal's on Coutrot1 and AVAD.

These comparisons have three limits. First, they cross evaluation code, and published evaluations of the same model on these data disagree by a few hundredths: STAViS on Coutrot2, for example, is reported at a CC of 0.734 [91] and of 0.652 [45]. Margins of 0.02 to 0.03 are therefore within the noise of such comparisons, and the tie margin of 0.01 is smaller than that noise. Second, they are restricted to peer-reviewed results. Results that had not appeared in peer-reviewed venues when this thesis was written

Table 5.10: CC on the six audio-visual datasets (fold means; DIEM has one split). The lower block lists the best peer-reviewed results whose tables could be verified, computed with other evaluation code. Visual-only: ViNet-E, or ViNet-A on Coutrot2 [45]; on SumMe, which those models do not report, the visual-only variant of the relevance-guided fusion model AVRSP of Yu et al. [116]. Audio-visual: AVRSP [116] on DIEM and Coutrot2, DiffSal [91] on the other four

Model	DIEM	Coutrot1	Coutrot2	ETMD	AVAD	SumMe
SURGE-D ViT-L, zero-shot	0.675	0.600	0.645	0.588	0.692	0.551
SURGE-D ViT-L, fine-tuned	0.706	0.632	0.816	0.662	0.718	0.581
SURGE-D ViT-L, 3-seed ensemble	0.714	0.641	0.807	0.670	0.729	**0.593**
SURGE-D ViT-g, fine-tuned	0.714	0.648	0.820	0.672	0.731	0.586
SURGE-D ViT-g, 3-seed ensemble	**0.729**	**0.657**	0.819	**0.676**	**0.740**	0.590
Best visual-only (published)	0.701	0.614	**0.862**	0.632	0.729	0.494
Best audio-visual (published)	0.685	0.638	0.837	0.652	0.738	0.572
ViT-g ens. − best visual-only	+0.028	+0.043	−0.043	+0.044	+0.011	+0.096
ViT-g ens. − best audio-visual	+0.044	+0.019	−0.018	+0.024	+0.002	+0.018

are not included, and some of them are higher than the peer-reviewed results on several of these datasets; against them, the margins on Coutrot1 and ETMD would also be ties. Third, they are restricted to results whose tables could be accessed: further recent audio-visual models, among them TM2SP [109], FOSal [128] and the model of Mao et al. [129], could not be checked and are therefore not represented. For AVRSP, as for SalFoM in Table 5.9, the publisher's version could not be accessed, and the numbers were taken from the authors' openly available version of the cited paper. The table should consequently not be read as establishing a state of the art on these datasets. Coutrot2 is discussed in Section 5.7.

Hollywood-2 with the Vig gaze data. Table 5.11 reports the 884 test clips of Hollywood-2 with the gaze of Vig et al. [84] (257,733 scored frames). These numbers are not comparable with published Hollywood-2 results (Section 4.3). With this many test clips, small differences are resolved precisely: ViT-g was better than ViT-L on all six metrics as a single model (ΔCC +0.0046 [+0.0032, +0.0060]) and on five as an ensemble (+0.0041 [+0.0030, +0.0052]). Ensembling raised CC, KLD, AUC-J and sAUC but significantly *lowered* NSS and SIM, for both encoders. With at most five observers per frame, the ground truth is sparse, and the averaging of three members smooths the peaks that NSS and SIM reward.

5.5.4. Qualitative Results

Figure 5.3 shows frames from five DHF1K validation videos, chosen by a rule fixed before any map was inspected: the videos at the 10th, 50th and 90th percentile of the ensemble's per-video CC, and the videos with the largest gain and the largest loss relative to SalFoM; the frame shown is the middle frame of each video.

Table 5.11: Results on Hollywood-2 with the gaze of Vig et al. (884 test clips). Not comparable with published Hollywood-2 results, which use other ground truth

Model	CC↑	NSS↑	SIM↑	KLD↓	AUC-J↑	sAUC↑
SURGE-D ViT-L, zero-shot	0.632	3.595	0.503	1.061	0.951	0.780
SURGE-D ViT-L, fine-tuned	0.6798	3.898	0.5178	0.925	0.9564	0.8167
SURGE-D ViT-L, 3-seed ensemble	0.6841	3.875	0.5073	0.918	0.9574	0.8192
SURGE-D ViT-g, fine-tuned	0.6844	**3.933**	**0.5214**	0.917	0.9569	0.8180
SURGE-D ViT-g, 3-seed ensemble	**0.6882**	3.907	0.5104	**0.911**	**0.9577**	**0.8199**



Figure 5.3: Qualitative examples on DHF1K validation. Rows: videos at the 10th, 50th and 90th percentile of SURGE-D's per-video CC, and the videos with the largest gain (0696) and loss (0663) relative to SalFoM; columns: middle frame with the observers' fixations, ground-truth density, SalFoM and the SURGE-D ViT-g ensemble, overlaid on the frame. Numbers give the CC of the shown frame. Video-level CC (SURGE-D / SalFoM): 0650 0.455/0.376, 0632 0.613/0.639, 0685 0.741/0.715, 0696 0.664/0.437, 0663 0.388/0.462

On these frames, the two models produce similar maps: both locate the horse rider, the surfer and the flying object that observers follow, and on the single frames shown SalFoM is in fact marginally closer to the ground truth in four of the five rows and equal in the fifth. Single frames are a poor guide to average differences that amount to a few hundredths of CC: for video 0696, for example, SURGE-D's CC over the whole video is 0.23 higher, yet the two maps at its middle frame are almost equal. The last row shows a characteristic failure. In a panel discussion, the observers' fixations concentrate on one person on the right, while SURGE-D spreads its prediction over several faces at the table. In conversation scenes, gaze is drawn to the faces of the people who are speaking [2], which a model that sees only the image cannot determine; the same limitation reappears on Coutrot2 below.

5.5.5. Computational Cost

Table 5.12 reports the parameters and operations of SURGE-D, counted with the PyTorch FLOP counter on one $16 \times 256 \times 256$ clip. Because one forward pass yields maps for 16 original frames, the cost per output frame is one sixteenth of the cost per clip.

Table 5.12: Parameters and operations of SURGE-D. TFLOPs count multiply–add pairs as two operations

	ViT-L	**ViT-g**
Encoder parameters (frozen)	303.9 M	1,012.2 M
Decoder parameters (trained)	6.99 M	7.39 M
Encoder TFLOPs per clip	1.656	5.088
Decoder TFLOPs per clip	0.166	0.168
TFLOPs per output frame	0.114	0.328

For comparison, the UMT-L encoder that SalFoM fine-tunes requires 2.87 TFLOPs for one $16 \times 224 \times 224$ clip (3,136 tokens; counted on a re-implementation loaded from the released checkpoint and, like the counts of SURGE-D, including the products with the attention matrices), and SalFoM runs it once per output frame in a sliding window [10]. Counting the encoder alone, SalFoM therefore needs about 25 times the computation per output frame of SURGE-D with ViT-L, and about 9 times that of SURGE-D with ViT-g.

Operation counts leave out memory traffic and the efficiency of an implementation, so the released SalFoM model and SURGE-D were also timed on the same GPU. Before the measurement, it was fixed that the statement of an efficiency advantage would be withdrawn if SalFoM needed less than twice the time per output frame of offline SURGE-D with ViT-g at either batch size. Each model ran 10 warm-up and 50 timed forward passes in bfloat16 on random input of its own size, already on the GPU, so

that frame decoding, resizing and the final blur are excluded for every model. The measurement was made three times. During the first two runs, short commands were started on the same machine; during the third, none was, and Table 5.13 reports it. The medians of the three runs differ by at most 5.1 %.

Table 5.13: Measured inference time on one RTX A5000, in milliseconds per forward pass and per output frame (median of 50 passes; third of three runs). The causal rows time the same forward pass as the offline rows but keep one map per pass

	Batch 1		Batch 8	
Model	Pass	Frame	Pass	Frame
SalFoM, released (one map per pass)	227.0	227.0	1619.4	202.4
SURGE-D ViT-L, offline (16 maps per pass)	88.5	5.5	409.2	3.2
SURGE-D ViT-L, causal (one map per pass)	88.5	88.5	409.2	51.2
SURGE-D ViT-g, offline (16 maps per pass)	165.5	10.3	988.3	7.7
SURGE-D ViT-g, causal (one map per pass)	165.5	165.5	988.3	123.5

In every run and at both batch sizes, SalFoM needed 22 to 26 times the time per output frame of offline SURGE-D with ViT-g, and 41 to 64 times that of offline SURGE-D with ViT-L, so the criterion is met by a wide margin. The ratios of time are larger than those of operations, and they reflect the released implementations: SalFoM's encoder forms its attention matrix explicitly over 3,136 tokens, whereas SURGE-D uses a fused attention kernel over 2,048 tokens. How much of the difference this explains was not measured. Peak memory at batch 8 was 14.2 GB for SalFoM, against 3.9 and 8.2 GB for SURGE-D with ViT-L and ViT-g. Offline, both SURGE-D models run faster than real time on this GPU.

The saving exists only in offline prediction, where one pass serves sixteen frames. A causal model, which may not see later frames, needs one pass per output frame, so its cost per output frame is the cost per clip in Table 5.12: 1.82 TFLOPs with ViT-L, about two thirds of SalFoM's encoder, and 5.26 TFLOPs with ViT-g, nearly twice as much. Measured, a causal ViT-L model took 0.25 to 0.39 times SalFoM's time per output frame, and a causal ViT-g model 0.61 to 0.73 times. Run one after another, the three models of the causal ViT-L ensemble need about 1.9 times SalFoM's encoder computation and 0.75 to 1.17 times its time. These causal comparisons are descriptive; the criterion concerned the offline model only.

5.6. Ablation Study

Table 5.14 collects the ablations. Each was evaluated on the holdout under the rules of Table 4.1. The four component ablations were first run with seed 0 and, in the follow-up to the first draft, repeated with seeds 1 and 2; each run is compared with the frozen ViT-L model of the same seed trained with the same three-epoch schedule. For the

comparison, the holdout scores of all fifteen models were recomputed with the evaluator of the thesis and reproduced the values recorded during training.

Table 5.14: Ablations on the DHF1K holdout (60 videos) unless marked "val". ΔCC is relative to the reference in each block; for the components, it is the mean over three seeds of the difference to the reference of the same seed, with its 95 % interval from a bootstrap over seeds and videos

Variant	CC	ΔCC	Outcome under the rule
Components (ViT-L, seeds 0 / 1 / 2; reference CC 0.5772 / 0.5801 / 0.5776)			
Without centre prior	0.5772 / 0.5792 / 0.5754	−0.0010 [−0.0026, +0.0005]	equivalent
Without auxiliary heads	0.5757 / 0.5790 / 0.5773	−0.0010 [−0.0023, +0.0003]	equivalent
With SALICON co-training	0.5801 / 0.5801 / 0.5775	+0.0009 [−0.0039, +0.0058]	inconclusive
Stride 4 (≈7.5 fps, 2 s window)	0.5803 / 0.5838 / 0.5795	+0.0029 [−0.0009, +0.0068]	inconclusive
Encoder adaptation (ViT-L, three seeds)			
LoRA, seed 0 / 1 / 2	0.5743 / 0.5791 / 0.5739	−0.0029 / −0.0010 / −0.0037	rejected (lost 3 of 3)
Encoder scale (seed 0; gate 0.5833)			
ViT-g instead of ViT-L	0.5866	+0.0083 vs ViT-L mean	adopted
Clip length (ViT-g, seed 0; gate 0.5904)			
32 frames instead of 16	0.5902	+0.0036 vs seed 0	not pursued (gate missed by 0.0002)
Training data (ViT-g, val, three-seed ensembles)			
600 instead of 540 training videos (val)	0.6044	+0.0015 [−0.0006, +0.0037]	not adopted

Encoder adaptation. LoRA was expected to help, because adapting a pretrained encoder to the target task usually improves on frozen features, and LoRA does so while training only a small fraction of the parameters that full fine-tuning would change [77]. Here it did not. Each LoRA model started from the frozen model of the same seed and lost to it on the holdout in all three seeds (mean −0.0025 CC), and its best epoch was always the first. With only 540 training videos, adapting 1.6 million encoder parameters appears to add more variance than signal; the same early overfitting that limits the decoder (Figure 5.1) would act on the adapted encoder as well. A second, compatible explanation is that adaptation degrades features that already suit the task. Kumar et al. [75] showed that fine-tuning can distort good pretrained features, although the harm they measured was clearest under distribution shift, whereas the holdout here comes from the training distribution; the explanation therefore remains a hypothesis. Full fine-tuning, as done by SalFoM, was not tested.

Components. With one seed, a difference of the size of the threshold cannot be told apart from seed variation, so the four component ablations were repeated with two further seeds and assessed with the three-seed rule of Table 4.1. Removing the learned centre prior or the auxiliary heads was equivalent to keeping them: the seed-averaged change in holdout CC was −0.0010 in both cases, and its 90 % interval lies within ±0.003. The centre prior's lack of effect suggests that the encoder features and the convolutional decoder can already represent the centre bias, for instance through the image borders, which the squashed input keeps at fixed positions; the explicit prior then adds nothing that can be measured. Co-training on 10,000 static SALICON images per epoch was inconclusive: its gain of +0.0028 with seed 0 did not recur with seeds 1 and 2, and its 95 % interval includes zero (−0.0039 to +0.0058) while its 90 % interval extends beyond the margin (−0.0031 to +0.0049). The static images, which helped CNN-era models [8], thus brought no measurable benefit here, although a small effect in either direction cannot be ruled out. A stride of four, which halves the cost per output frame, had reached exactly the single-seed threshold (+0.0031). It raised holdout CC with all three seeds (+0.0031, +0.0037 and +0.0020), but the mean of +0.0029 lies below the margin and its interval includes zero, so this outcome is inconclusive as well. Outside the three outcomes of the rule, its 90 % interval (−0.0003 to +0.0062) excludes a loss as large as the margin: the cheaper stride was at least not measurably worse.

Encoder scale. Replacing ViT-L by ViT-g, with every other setting unchanged, raised holdout CC for all three seeds (0.5866, 0.5867 and 0.5829 against 0.5772, 0.5801 and 0.5776) and validation CC by 0.017 for single models (Table 5.6). Scale was the largest single improvement found after the choice of encoder family.

Temporal context and training data. Doubling the clip to 32 frames (about 2.1 s) reached a holdout CC of 0.5902, 0.0002 below its gate, so it was not pursued; unlike the 16-frame runs, it was still improving at its last epoch, so a longer schedule might change this conclusion, but that would be a new experiment. With one seed and an unfinished learning curve, this result is inconclusive rather than negative. Training on all 600 DHF1K training videos instead of 540 did not change validation CC (+0.0015 [−0.0006, +0.0037] for the ensembles), but it moved other metrics in both directions (Table A.7): for the ensembles, SIM fell slightly; for seed 0, NSS and SIM rose and sAUC fell. Neither more context nor 11 % more videos improved CC measurably, but the clip-length test was not conclusive.

Training-set size. The decoder overfits within three epochs (Figure 5.1), and no addition to it brought an established gain. Both observations would follow if the 540 training videos, rather than the model, limited accuracy. Adding the 60 holdout videos

to the training set could not test this: the step is small, and it leaves no holdout on which to select the epoch. A learning curve was therefore measured in the second follow-up, under a rule written before the runs (Table 4.1). The frozen ViT-L model was trained on random subsets of 270 and of 135 of the 540 training videos, with three seeds. For each seed the subsets were drawn anew, the smaller one from within the larger. The number of clips sampled per video in each epoch was raised from 8 to 16 and to 32, so that every model took the same 1,620 optimisation steps, and all other settings were those of the main model. The models were scored on the 60-video holdout in the same way as the three-seed ablations, and each was compared with the 540-video model of the same seed.

Figure 5.4 and Table 5.15 give the result. Holdout CC fell with every halving of the training set, in every seed: from 0.5783 with 540 videos to 0.5729 with 270 and 0.5645 with 135 (means over seeds); NSS, SIM and KLD moved in the same direction. Under the rule, the primary comparison is nevertheless inconclusive. Halving the training set from 540 to 270 videos cost 0.0054 CC on average, which is more than the margin, but the 95 % interval of the difference includes zero (−0.0112 to +0.0011). The second halving, from 270 to 135 videos, was an effect (−0.0084 [−0.0164, −0.0004]), as was the difference between the two ends of the curve (−0.0138). The hypothesis that the amount of training data limits accuracy at 540 videos is therefore neither established nor refuted. Equivalence, the outcome that would have refuted it, is not supported either, because the 90 % interval of the primary difference (−0.0101 to −0.0001) extends well beyond the margin.

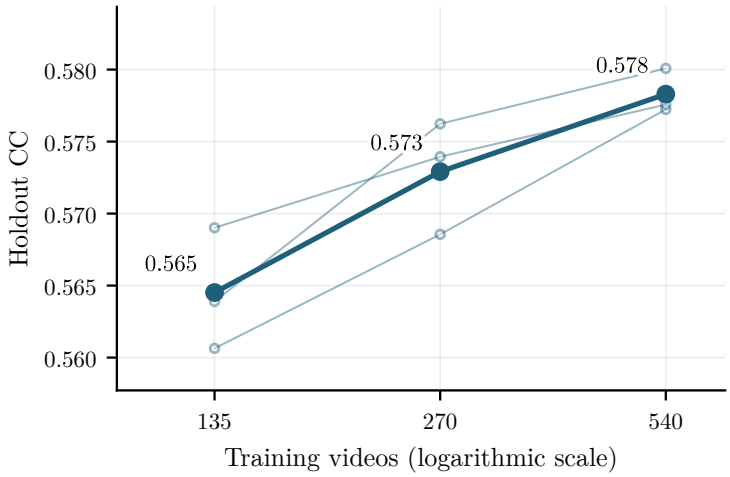


Figure 5.4: Holdout CC of the frozen ViT-L model against the number of training videos. Thin lines are the three seeds, whose subsets are nested, and the thick line with the printed values is their mean

Two further observations are descriptive and carry no outcome under the rule. First, on the logarithmic axis of Figure 5.4 the curve flattens only gradually. The step from 270 to 540 videos brought about two thirds of the gain of the step from 135 to 270,

Table 5.15: Learning curve over the number of training videos (frozen ViT-L, DHF1K holdout, seeds 0 / 1 / 2). The last four columns of the upper part are means over the seeds. The differences in the lower part are paired by seed and video, their intervals come from the bootstrap over seeds and videos, and the outcomes follow the three-seed rule of Table 4.1 with its margin of 0.003

Training videos	CC per seed	CC	NSS	SIM	KLD
540	0.5772 / 0.5801 / 0.5776	0.5783	3.323	0.445	1.227
270	0.5686 / 0.5762 / 0.5739	0.5729	3.286	0.442	1.238
135	0.5606 / 0.5639 / 0.5690	0.5645	3.227	0.425	1.263

Comparison	ΔCC	95 % interval	90 % interval	Outcome
270 − 540 (primary)	−0.0054	[−0.0112, +0.0011]	[−0.0101, −0.0001]	inconclusive
135 − 270	−0.0084	[−0.0164, −0.0004]	[−0.0150, −0.0018]	effect
135 − 540	−0.0138	[−0.0238, −0.0048]	[−0.0219, −0.0063]	effect

and its size, 0.005 CC, is larger than the change produced by any decoder component (at most 0.003) and of the same order as the gain from replacing ViT-L by ViT-g on the same holdout (+0.007 on average over the three seeds). Second, the gap between training and holdout accuracy widened as the training set shrank: the CC on the training clips during the last epoch rose from 0.683 with 540 videos to 0.702 with 270 and 0.718 with 135, while the holdout CC fell. Both observations are what a model limited by its training data would show, but neither establishes it.

The limits of this test were written down with its rule. It used three seeds and one encoder. Every model took the same number of optimisation steps, so each video of a smaller subset was sampled more often. The holdout that scores the models also selected the epoch of each run, in the same way for every subset size. In addition, the curve ends at 540 videos because DHF1K offers no further annotated training videos apart from the 60 of the holdout, so whether more data than the present training set would help cannot be read from it.

Ensembles and output blur. Averaging three seeds raised validation CC by 0.007 for ViT-L and by 0.003 for ViT-g (Table 5.5); the stronger encoder leaves less variance for the ensemble to remove. In paired comparisons with the seed-0 model, both gains have intervals above zero (+0.0074 [+0.0050, +0.0099] and +0.0025 [+0.0002, +0.0049]), and NSS, KLD and AUC-J improved as well, but sAUC fell slightly for both encoders (−0.0022 and −0.0029, intervals below zero; Table A.8). Figure 5.5 shows that the optimal blur is broad: holdout CC varied by at most 0.003 between $\sigma = 0$ and $\sigma = 12$ for every configuration. The ViT-L ensemble preferred less blur than its single model (8 against 12 pixels), consistent with the smoothing that averaging already provides.

Audio. The audio fusion module of Section 4.5.3 was trained on every fold of the six audio-visual datasets with the same recipe as the visual fine-tuning (ViT-L, seed 0).

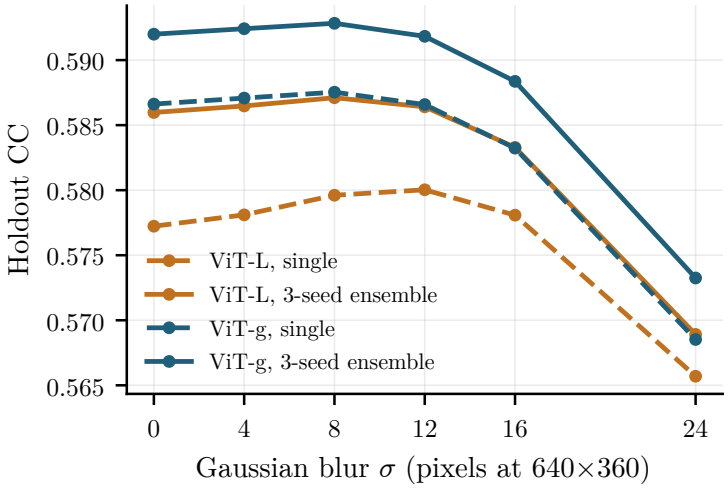


Figure 5.5: Holdout CC as a function of the Gaussian blur applied to the output, for single models and three-seed ensembles. The selected values are 12 (ViT-L single) and 8 (others)

Table 5.16 shows that it did not improve the visual model on any dataset: the largest change was +0.001 CC, far below the threshold of +0.010. The control is informative nonetheless. On Coutrot2, scoring the trained model with the wrong soundtrack cost 0.018 CC, so the model had learned to use the audio there, but with the correct soundtrack it only matched the visual model, which already finds the faces in these meeting scenes from the image alone. With five test videos per fold, this observation is suggestive rather than conclusive.

Table 5.16: Audio fusion on the audio-visual datasets (ViT-L, fold-mean CC). The rule requires audio − visual ≥ +0.010 and matched − rotated ≥ +0.005

Dataset	Visual	Audio, matched	Audio, rotated	Audio − visual	Matched − rotated
DIEM	0.706	0.706	0.706	+0.0008	+0.0002
Coutrot1	0.632	0.633	0.630	+0.0010	+0.0031
Coutrot2	0.816	0.816	0.797	−0.0003	+0.0183
ETMD	0.662	0.661	0.661	−0.0014	−0.0001
AVAD	0.718	0.718	0.717	+0.0002	+0.0013
SumMe	0.581	0.579	0.580	−0.0021	−0.0011

5.7. Discussion

This section interprets the results in terms of the research questions of Section 1.4.

RQ1: Can a frozen self-supervised video encoder with a light decoder predict fixations better than the strongest foundation-model approach? On DHF1K, yes for the larger encoder, also under an identical protocol, with qualifications. With ViT-g, every seed and the three-seed ensemble were more accurate than SalFoM, whose 300-million-parameter encoder is fully fine-tuned, under the same evaluator, and they

stayed more accurate when SalFoM's maps received the blur that is best for it on the test set; the single model and the ensemble were then better on all six metrics. With ViT-L, whose encoder is the size of SalFoM's, the ensemble was better, offline and causally, but single models were not reliably so once SalFoM's blur was tuned. Three features of the result matter for its interpretation. First, the gain comes from the encoder rather than from freezing as such. With the same decoder and frozen protocol, the UMT-L encoder of SalFoM's family scored below SalFoM and 0.039 below V-JEPA 2 ViT-L, so a frozen encoder suffices only when its features suit the task. The frozen V-JEPA 2 encoder also outperformed its LoRA-adapted version in the one adaptation setting that was tried; full fine-tuning was not tested. Second, the gain does not depend on frames after the predicted one. Restricted to past frames, as SalFoM is, the ensembles were still more accurate than SalFoM with its best blur on the same frames (+0.029 CC with ViT-g and +0.017 with ViT-L), and they were as accurate as the offline ensembles (Section 5.5.2); this causal comparison is the one made under an identical protocol. It gives up the efficiency of the offline model, however, since a causal model needs one forward pass per output frame. Third, the evidence is a validation-set comparison with one re-scored baseline; the official test-set evaluation is still outstanding.

RQ2: Does predictive surprise carry information about where people look? It does, but in the opposite direction to the hypothesis. On its own, the error of the released V-JEPA 2 predictor scored below both simple reference cues at both frame rates and was anti-correlated with fixations, with the trained mask token as with the untrained one, so the gate failed. Combined with the centre prior, frame difference and token atypicality, the error did add information about fixations: 0.069 and 0.062 bits per fixation at 4 fps in the two primary comparisons, about a fifth of what the two simple cues together add over the centre prior. Its fitted weight, however, was negative: observers looked less often where the error was high. Filling in masked tubes, a query that matches pre-training in the shape of its masks, gave the same picture as forecasting the second half of a clip, so the form of the query does not explain the sign. The question as posed is therefore answered for this predictor: its error carries information beyond the simple cues, but not the information that surprise-driven attention predicts. This is a statement about one released predictor and two ways of reading it out, not a refutation of surprise-based accounts of attention [13]. What the error measures that draws gaze away was not established, and a predictor trained to forecast could behave differently.

RQ3: Which design factors matter when the encoder is frozen? The encoder dominates: its pre-training (V-JEPA 2 over V-JEPA 2.1 and UMT-L) and its scale (ViT-g over ViT-L) produced the only gains that were large relative to seed noise

on DHF1K, although the comparison of pre-training rests on one seed per encoder, and V-JEPA 2.1, which trailed narrowly, was used below its training resolution. Seed ensembles added a smaller gain in CC, with intervals above zero for both encoders, at a slight cost in sAUC. With three seeds, removing the learned centre prior or deep supervision was equivalent to keeping them, within ±0.003 CC, and LoRA lost to the frozen encoder in every seed. Static-image co-training, 32-frame clips and 11 % more training videos produced no gain that could be established; for static images and longer clips, the evidence is inconclusive rather than negative. A sparser temporal stride, at half the cost per output frame, improved holdout CC with all three seeds without reaching the margin. One explanation is that, with a strong frozen encoder, the small number of training videos limits further progress: the decoder overfits within three epochs (Figure 5.1), and the additions that would increase its effective capacity do not help, whereas better features do. The learning curve of Section 5.6 tested this explanation and left it open. Holdout CC rose with each doubling of the training set in every seed, by 0.008 from 135 to 270 videos, which is an effect under the rule, and by 0.005 from 270 to 540 videos, where the interval includes zero. The amount of data matters between 135 and 270 videos; whether it still limits the model at 540 is neither established nor refuted.

RQ4: How well does the approach transfer, and does audio help? The model transferred well after brief fine-tuning, within the limits of comparisons across evaluation code. Fine-tuned for three epochs, its ensemble reached a higher CC than the best peer-reviewed results that could be verified on four of the six audio-visual datasets, tied on AVAD and was lower on Coutrot2. These margins are of the order of the disagreements between published evaluations of the same models, and the comparison depends on its restriction to peer-reviewed results. The model was competitive on UCF Sports, where action-oriented encoders retain an edge. The one clear failure, Coutrot2, has a plausible account: in recordings of conversations, observers look at the person who is speaking, which an image-only model cannot determine, as the last row of Figure 5.3 also illustrates. The audio experiment, with one seed and the ViT-L encoder, shows that adding the soundtrack in a simple way does not resolve this. The fusion module learned to depend on the audio on Coutrot2, but it did not improve on the visual model. This agrees with the observations of Jain et al. [102] and Agrawal et al. [117], who found that deep audio-visual saliency models tend to become insensitive to their audio input and that none of a range of audio modules and fusion techniques improved on visual-only models across ten datasets. Exploiting sound in this setting probably requires explicit speaker localisation rather than global feature modulation.

Relation to the original proposal. The final model differs from the proposed one in three respects, each decided by a test whose rule was written in advance: the encoder is V-JEPA 2 rather than V-JEPA 2.1, it is frozen rather than adapted with LoRA, and the surprise branch that gave the project its name is absent. The second follow-up did not overturn the last decision: the predictor's error carries information about fixations, but with the opposite sign to the one on which the branch was premised; whether a decoder could use it as a cue with a negative weight was not tested. The proposal's ranking of what drives performance, which placed the encoder first, was confirmed. Its expectation that LoRA would help, based on results with much larger training sets, was not.

5.8. Threats to Validity

Internal validity. All selection used holdout data, every test set was scored once per reported model, and the decision rules were written in advance in a dated log kept with the code, with the one exception of the tie margin for published results (Section 4.1); this protects against selection bias. The log is an internal record rather than an external registration, and the DHF1K validation set, although never used to select a reported model or setting, informed two decisions about the official test submission (Table A.4). Several residual weaknesses remain. The thresholds of the rules (for example the margin of 0.005 in the gates) were chosen by judgement, partly before any seed variation had been measured, rather than derived formally, and the tie margin of 0.01 for published results was fixed only after the first such comparison had been seen. Many design choices were fixed without a test (Section 4.1). On the audio-visual datasets, the best epoch and the output blur were selected on one to six holdout videos per fold. Finally, the SalFoM baseline was run by a program that is not part of the code of this thesis, and how its released weights were selected is not documented; if they were selected on the DHF1K validation set, the baseline's scores on that set are optimistic, which would make the reported differences conservative (Section 5.2).

Construct validity. The six metrics capture different aspects of a saliency map [17] and did not always agree; ensembles, for instance, improved CC but lowered NSS and SIM on the sparse Hollywood-2 (Vig gaze) data. The use of one map for all metrics, with the blur chosen by CC, favours CC slightly over the other metrics, and the choice of blur matters for the baseline as well: SalFoM's CC rose by 0.010 when its maps were blurred. The prediction error is a distance in the predictor's feature space rather than a change of belief, so it is one of several possible measures of surprise, and the masked-tube readout matches pre-training in the shape of its masks but hides a smaller share of the clip. The Hollywood-2 (Vig gaze) maps rest on at most five observers per

frame, and their reliability was not measured. The shuffled AUC depends on how its negatives are drawn and is comparable only within this thesis.

External validity and comparability. Comparisons with published numbers cross evaluation code and, in places, protocol details, such as the handling of frames without fixations or the exact fold composition. Published evaluations of the same model disagree by a few hundredths of CC, and the comparisons are restricted to peer-reviewed results, which on the audio-visual datasets leaves out some higher results. Only one baseline, SalFoM, was re-scored under the protocol of this thesis. SURGE-D's offline predictions use about 0.27 s or more of future context, which causal models do not have. On DHF1K, the causal control (Section 5.5.2) showed that this context adds no measurable accuracy, but the control was not repeated on the other datasets. The DHF1K test-set result and the standard Hollywood-2 ground truth were not available. Finally, the pre-training data of V-JEPA 2 include large collections of internet video, and DHF1K is drawn from YouTube; an overlap between pre-training and test videos cannot be ruled out. The same is true for every method that relies on a large pretrained video encoder, including SalFoM.

CHAPTER SIX

CONCLUSION AND FUTURE WORK

6.1. Conclusion

This thesis asked whether the representations of a self-supervised video encoder trained by feature prediction are sufficient for video saliency prediction without adapting the encoder, and whether the encoder's own prediction error indicates where people look. To answer these questions, it developed SURGE-D, which combines a frozen V-JEPA 2 encoder with a light decoder that reads four intermediate layers, rebuilds a multi-scale pyramid, decodes it with factorised spatio-temporal convolutions and predicts saliency maps for sixteen frames per forward pass. Each component that was tested was evaluated under a rule written in a dated log before the corresponding experiment, all models were scored once on their test sets by a single evaluator validated against reference code, and differences were assessed with paired statistics. The answers to the research questions are as follows.

RQ1. With the larger encoder, yes. On the DHF1K validation set, the frozen V-JEPA 2 ViT-g model reached a correlation coefficient of 0.600 as a single model and 0.603 as a three-seed ensemble, against 0.564 for SalFoM, the closest foundation-model approach, scored by the same evaluator, and was better on all six metrics with paired bootstrap intervals that exclude zero. Every seed stayed ahead when SalFoM's maps were given the blur that suits them best. Under an identical protocol, with both models restricted to past frames and scored on the same frames, the causal ViT-g ensemble was more accurate than SalFoM by 0.029 and as accurate as the offline model, so the advantage does not come from future context. With the ViT-L encoder, the size of SalFoM's, only the ensemble was reliably better, offline and under the identical protocol, where its margin was 0.017. The gain comes from the V-JEPA 2 encoder rather than from freezing as such, because a frozen UMT-L encoder with the same decoder scored below SalFoM. In offline use, SURGE-D with ViT-g needs about a ninth of SalFoM's encoder computation per output frame and, in the released implementations, less than a twentieth of its time; the causal variant needs nearly twice the computation but less time. These are validation-set results; the official test-set evaluation was pending when the thesis was written.

RQ2. The prediction error of the released V-JEPA 2 predictor carries information about where people look, but not in the way that the surprise hypothesis predicts. On

its own, it scored below frame difference and token atypicality at both frame rates tested and was anti-correlated with fixations, with bootstrap intervals over the 60 holdout videos that exclude zero. This held with the trained mask token as with the untrained one of the first run, and for a readout that fills in masked tubes as in pre-training. Combined with the centre prior and the two simple cues, the error added 0.069 and 0.062 bits per fixation in the two comparisons fixed in advance, with a negative weight: observers looked less often where the error was high. The finding concerns one released predictor and two ways of reading it out; it does not refute surprise-based accounts of attention.

RQ3. With a frozen encoder, the encoder itself dominated. Its pre-training (V-JEPA 2 over V-JEPA 2.1 and UMT-L) and its scale (ViT-g over ViT-L) produced the only large gains, and seed ensembles a small one. Over three seeds, removing the learned centre prior or deep supervision was equivalent to keeping them, within ±0.003 CC, and low-rank adaptation of the encoder lost in every seed. Static-image co-training, 32-frame clips and 11 % more training videos produced no gain that could be established, and a sparser temporal stride, which halves the cost per output frame, improved with every seed without reaching the margin. Whether the amount of training data, rather than the capacity of the decoder, limits further progress was tested with a learning curve and remains open: holdout CC rose with every doubling of the training set between 135 and 540 videos, but the last doubling, a gain of 0.005, was inconclusive under its rule.

RQ4. The approach transferred well after brief fine-tuning, within the limits of comparisons across evaluation code. It was competitive with the best published results on UCF Sports. On the audio-visual datasets, its ensemble reached a higher correlation than the best peer-reviewed results that could be verified on four of the six datasets, tied on a fifth and was lower on the sixth. These margins are of the same order as the disagreements between published evaluations of the same models, and the comparison is limited to peer-reviewed results, so it does not establish a state of the art. Its clear failure was on recordings of conversations, where gaze follows the speaker, which an image-only model cannot identify. Adding the soundtrack through a simple fusion module did not improve any dataset, in line with earlier findings that audio-visual saliency models tend not to exploit the audio.

6.2. Contributions

The thesis makes four contributions. First, it presents SURGE-D, a video saliency model with a frozen self-supervised encoder and a 7-million-parameter multi-frame decoder, which was more accurate than a fine-tuned foundation-model approach on the

DHF1K validation set, both offline, at a fraction of its encoder computation, and as a causal model under an identical evaluation protocol. Second, it reports a negative result, tested under rules written in advance: the error of the released V-JEPA 2 predictor is anti-correlated with fixations and, combined with simple motion and appearance cues, adds information only with a negative weight, so it does not behave as a surprise signal that attracts gaze. Along the way, it shows that the predictor of a widely used model-loading library reads by default a mask token that pre-training never updated. Third, it provides controlled evidence about which design choices matter once the encoder is frozen: the encoder's pre-training and scale dominate, while two components commonly included in saliency models, a learned centre prior and deep supervision, are equivalent to their absence over three seeds, and the low-rank adaptation of the encoder that was tested did not help. Fourth, it contributes an evaluation methodology, with a validated evaluator, decision rules fixed in advance, paired statistics and a reproducible ground truth for Hollywood-2 built from openly licensed gaze recordings, that makes its positive and negative findings verifiable.

6.3. Future Work

The limitations of the study point to several directions for future work.

Completing the evaluation. The DHF1K test-set scores of the prepared predictions and an evaluation on the standard Hollywood-2 ground truth would place the results on the benchmarks' official footing. Both depend on the dataset authors and were pending when the thesis was written.

Strengthening the ablations. With three seeds, two of the four component ablations remained inconclusive. A temporal stride of four, which halves the cost per output frame, improved holdout CC with every seed without reaching the margin, and is the most useful one to settle, with more seeds or on a larger evaluation set; static-image co-training is the other. The encoder comparison and the clip-length test used one seed each, and the 32-frame model was still improving at the end of its schedule, so it merits a longer training run under a new rule written in advance. The learning curve ended at the 540 training videos that DHF1K provides, and its last step was inconclusive; more seeds, or training videos from a second dataset, would show whether the amount of data limits progress.

Adapting the stronger encoder. Low-rank adaptation was tested only with ViT-L, and full fine-tuning not at all. Whether adaptation helps the larger ViT-g encoder, or helps once more training data are available, remains open. The V-JEPA 2.1 encoder

should also be compared at the resolution at which it was trained.

Speaker-aware audio-visual prediction. The failure on conversations calls for models that determine who is speaking, for example by explicit localisation of the active speaker, rather than modulating visual features with a global audio embedding.

Efficient causal prediction. A causal model needs one forward pass per output frame, so the saving that SURGE-D gains from predicting sixteen frames per pass does not carry over to streaming use: the causal ViT-g model needs nearly twice SalFoM's encoder computation per output frame, and the causal ViT-L model, which needs about two thirds of it, was reliably more accurate only as a three-seed ensemble. Cheaper causal designs, for instance with shorter windows or windows that forecast several frames ahead, remain to be tested, as does causal prediction on datasets other than DHF1K.

Revisiting surprise. The error of the released predictor carries information about fixations with the opposite sign to the hypothesis, and what it measures that draws gaze away is not known. A first step is to describe the regions where the error is high, for instance by the amount of texture, motion or object boundaries that they contain, and to test whether the negative weight survives when such properties are added to the simple cues. Predictors trained explicitly to extrapolate into the future, or surprise measured as a change of belief relative to the scene's own history, may behave differently, and the incremental-information analysis of this thesis could test them directly.

Evaluation practice. Scoring each metric with the map that is optimal for it under the model's prediction [94], and reporting paired statistics against re-scored baselines, would make comparisons in the field more informative.

REFERENCES

[1] P. K. Mital, T. J. Smith, R. L. Hill, and J. M. Henderson, "Clustering of gaze during dynamic scene viewing is predicted by motion," *Cognitive Computation*, vol. 3, no. 1, pp. 5–24, 2011, doi: 10.1007/s12559-010-9074-z.

[2] A. Coutrot and N. Guyader, "How saliency, faces, and sound influence gaze in dynamic social scenes," *Journal of Vision*, vol. 14, no. 8, 2014, Art. no. 5, doi: 10.1167/14.8.5.

[3] L. Itti and C. Koch, "Computational modelling of visual attention," *Nature Reviews Neuroscience*, vol. 2, no. 3, pp. 194–203, 2001, doi: 10.1038/35058500.

[4] A. Borji and L. Itti, "State-of-the-art in visual attention modeling," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 1, pp. 185–207, 2013, doi: 10.1109/tpami.2012.89.

[5] P. H. Tseng, R. Carmi, I. G. M. Cameron, D. P. Munoz, and L. Itti, "Quantifying center bias of observers in free viewing of dynamic natural scenes," *Journal of Vision*, vol. 9, no. 7, 2009, Art. no. 4, doi: 10.1167/9.7.4.

[6] A. Borji, "Saliency prediction in the deep learning era: Successes and limitations," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 43, no. 2, pp. 679–700, 2021, doi: 10.1109/tpami.2019.2935715.

[7] L. Itti, C. Koch, and E. Niebur, "A model of saliency-based visual attention for rapid scene analysis," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 20, no. 11, pp. 1254–1259, 1998, doi: 10.1109/34.730558.

[8] W. Wang, J. Shen, F. Guo, M.-M. Cheng, and A. Borji, "Revisiting video saliency: A large-scale benchmark and a new model," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2018, pp. 4894–4903, doi: 10.1109/cvpr.2018.00514.

[9] W. Wang, J. Shen, J. Xie, M.-M. Cheng, H. Ling, and A. Borji, "Revisiting video saliency prediction in the deep learning era," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 43, no. 1, pp. 220–237, 2021, doi: 10.1109/tpami.2019.2924417.

[10] M. Moradi, M. Moradi, F. Rundo, C. Spampinato, A. Borji, and S. Palazzo, "SalFoM: Dynamic saliency prediction with video foundation models," in *Proceedings of the 27th International Conference on Pattern Recognition (ICPR)*, ser. Lecture Notes in Computer Science, vol. 15322, Springer Nature Switzerland, 2025, pp. 33–48, doi: 10.1007/978-3-031-78312-8_3.

[11] M. Assran et al., "Self-supervised learning from images with a joint-embedding predictive architecture," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2023, pp. 15 619–15 629, doi: 10.1109/cvpr52729.2023.01499.

[12] A. Bardes et al., "Revisiting feature prediction for learning visual representations from video," *Transactions on Machine Learning Research*, 2024.

[13] L. Itti and P. Baldi, "Bayesian surprise attracts human attention," *Vision Research*, vol. 49, no. 10, pp. 1295–1306, 2009, doi: 10.1016/j.visres.2008.09.007.

[14] L. Mur-Labadia et al., "V-JEPA 2.1: Unlocking dense features in video self-supervised learning," in *Computer Vision – ECCV 2026*, ser. Lecture Notes in Computer Science, Springer Nature Switzerland, 2026, pp. 671–689, doi: 10.1007/978-3-032-37369-4_36.

[15] A. Paszke et al., "PyTorch: An imperative style, high-performance deep learning library," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.

[16] T. Wolf et al., "Transformers: State-of-the-art natural language processing," in *Proceedings of the Conference on Empirical Methods in Natural Language Processing: System Demonstrations (EMNLP)*, Association for Computational Linguistics, 2020, pp. 38–45, doi: 10.18653/v1/2020.emnlp-demos.6.

[17] Z. Bylinskii, T. Judd, A. Oliva, A. Torralba, and F. Durand, "What do different evaluation metrics tell us about saliency models?" *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 41, no. 3, pp. 740–757, 2019, doi: 10.1109/tpami.2018.2815601.

[18] D. D. Salvucci and J. H. Goldberg, "Identifying fixations and saccades in eye-tracking protocols," in *Proceedings of the Symposium on Eye Tracking Research and Applications (ETRA)*, ACM Press, 2000, pp. 71–78, doi: 10.1145/355017.355028.

[19] O. Le Meur and T. Baccino, "Methods for comparing scanpaths and saliency maps: Strengths and weaknesses," *Behavior Research Methods*, vol. 45, no. 1, pp. 251–266, 2013, doi: 10.3758/s13428-012-0226-9.

[20] N. Wilming, T. Betz, T. C. Kietzmann, and P. König, "Measures and limits of models of fixation selection," *PLoS ONE*, vol. 6, no. 9, e24038, 2011, doi: 10.1371/journal.pone.0024038.

[21] M. Kümmerer, T. S. A. Wallis, and M. Bethge, "Information-theoretic model comparison unifies saliency metrics," *Proceedings of the National Academy of Sciences*, vol. 112, no. 52, pp. 16 054–16 059, 2015, doi: 10.1073/pnas.1510393112.

[22] A. M. Treisman and G. Gelade, "A feature-integration theory of attention," *Cognitive Psychology*, vol. 12, no. 1, pp. 97–136, 1980, doi: 10.1016/0010-0285(80)90005-5.

[23] C. Koch and S. Ullman, "Shifts in selective visual attention: Towards the underlying neural circuitry," *Human Neurobiology*, vol. 4, no. 4, pp. 219–227, 1985, PMID: 3836989.

[24] J. Harel, C. Koch, and P. Perona, "Graph-based visual saliency," in *Advances in Neural Information Processing Systems 19*, The MIT Press, 2007, pp. 545–552, doi: 10.7551/mitpress/7503.003.0073.

[25] L. Zhang, M. H. Tong, T. K. Marks, H. Shan, and G. W. Cottrell, "SUN: A bayesian framework for saliency using natural statistics," *Journal of Vision*, vol. 8, no. 7, 2008, Art. no. 32, doi: 10.1167/8.7.32.

[26] T. Judd, K. Ehinger, F. Durand, and A. Torralba, "Learning to predict where humans look," in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, IEEE, 2009, pp. 2106–2113, doi: 10.1109/iccv.2009.5459462.

[27] M. S. Castelhano, M. L. Mack, and J. M. Henderson, "Viewing task influences eye movement control during active scene perception," *Journal of Vision*, vol. 9, no. 3, 2009, Art. no. 6, doi: 10.1167/9.3.6.

[28] A. Borji and L. Itti, "Defending yarbus: Eye movements reveal observers' task," *Journal of Vision*, vol. 14, no. 3, 2014, Art. no. 29, doi: 10.1167/14.3.29.

[29] R. Carmi and L. Itti, "Visual causes versus correlates of attentional selection in dynamic scenes," *Vision Research*, vol. 46, no. 26, pp. 4333–4345, 2006, doi: 10.1016/j.visres.2006.08.019.

[30] M. Dorr, T. Martinetz, K. R. Gegenfurtner, and E. Barth, "Variability of eye movements when viewing dynamic natural scenes," *Journal of Vision*, vol. 10, no. 10, 2010, Art. no. 28, doi: 10.1167/10.10.28.

[31] T. J. Smith and P. K. Mital, "Attentional synchrony and the influence of viewing task on gaze behavior in static and dynamic scenes," *Journal of Vision*, vol. 13, no. 8, 2013, Art. no. 16, doi: 10.1167/13.8.16.

[32] B. W. Tatler, "The central fixation bias in scene viewing: Selecting an optimal viewing position independently of motor biases and image feature distributions," *Journal of Vision*, vol. 7, no. 14, 2007, Art. no. 4, doi: 10.1167/7.14.4.

[33] A. Coutrot, N. Guyader, G. Ionescu, and A. Caplier, "Influence of soundtrack on eye movements during video exploration," *Journal of Eye Movement Research*, vol. 5, no. 4, 2012, doi: 10.16910/jemr.5.4.2.

[34] L. Itti and P. Baldi, "A principled approach to detecting surprising events in video," in *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2005, pp. 631–637, doi: 10.1109/cvpr.2005.40.

[35] R. P. N. Rao and D. H. Ballard, "Predictive coding in the visual cortex: A functional interpretation of some extra-classical receptive-field effects," *Nature Neuroscience*, vol. 2, no. 1, pp. 79–87, 1999, doi: 10.1038/4580.

[36] M. Jiang, S. Huang, J. Duan, and Q. Zhao, "SALICON: Saliency in context," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2015, pp. 1072–1080, doi: 10.1109/cvpr.2015.7298710.

[37] J. Pan, E. Sayrol, X. Giro-I-Nieto, K. McGuinness, and N. E. OConnor, "Shallow and deep convolutional networks for saliency prediction," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2016, pp. 598–606, doi: 10.1109/cvpr.2016.71.

[38] X. Huang, C. Shen, X. Boix, and Q. Zhao, "SALICON: Reducing the semantic gap in saliency prediction by adapting deep neural networks," in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, IEEE, 2015, pp. 262–270, doi: 10.1109/iccv.2015.38.

[39] S. S. S. Kruthiventi, K. Ayush, and R. V. Babu, "DeepFix: A fully convolutional neural network for predicting human eye fixations," *IEEE Transactions on Image Processing*, vol. 26, no. 9, pp. 4446–4456, 2017, doi: 10.1109/tip.2017.2710620.

[40] M. Cornia, L. Baraldi, G. Serra, and R. Cucchiara, "Predicting human eye fixations via an LSTM-Based saliency attentive model," *IEEE Transactions on Image Processing*, vol. 27, no. 10, pp. 5142–5154, 2018, doi: 10.1109/tip.2018.2851672.

[41] J. Lou, H. Lin, D. Marshall, D. Saupe, and H. Liu, "TranSalNet: Towards perceptually relevant visual saliency prediction," *Neurocomputing*, vol. 494, pp. 455–467, 2022, doi: 10.1016/j.neucom.2022.04.080.

[42] M. Kümmerer, T. S. Wallis, L. A. Gatys, and M. Bethge, "Understanding low- and high-level contributions to fixation prediction," in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, IEEE, 2017, pp. 4799–4808, doi: 10.1109/iccv.2017.513.

[43] A. Linardos, M. Kümmerer, O. Press, and M. Bethge, "DeepGaze IIE: Calibrated prediction in and out-of-domain for state-of-the-art saliency modeling," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2021, pp. 12 899–12 908, doi: 10.1109/iccv48922.2021.01268.

[44] S. Jetley, N. Murray, and E. Vig, "End-to-end saliency mapping via probability distribution prediction," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2016, pp. 5753–5761, doi: 10.1109/cvpr.2016.620.

[45] R. Girmaji, S. Jain, B. Beri, S. Bansal, and V. Gandhi, "Minimalistic video saliency prediction via efficient decoder & spatio temporal action cues," in *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, IEEE, 2025, pp. 1–5, doi: 10.1109/icassp49660.2025.10888852.

[46] W. Wang, Q. Lai, H. Fu, J. Shen, H. Ling, and R. Yang, "Salient object detection in the deep learning era: An in-depth survey," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 6, pp. 3239–3259, 2022, doi: 10.1109/tpami.2021.3051099.

[47] D.-P. Fan, W. Wang, M.-M. Cheng, and J. Shen, "Shifting more attention to video salient object detection," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2019, pp. 8546–8556, doi: 10.1109/cvpr.2019.00875.

[48] J. Xiong et al., "Toward unifying saliency transformer for video saliency prediction and detection," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 35, no. 8, pp. 7339–7353, 2025, doi: 10.1109/tcsvt.2025.3544031.

[49] D. Tran, L. Bourdev, R. Fergus, L. Torresani, and M. Paluri, "Learning spatiotemporal features with 3d convolutional networks," in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, IEEE, 2015, pp. 4489–4497, doi: 10.1109/iccv.2015.510.

[50] J. Carreira and A. Zisserman, "Quo vadis, action recognition? a new model and the kinetics dataset," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2017, pp. 4724–4733, doi: 10.1109/cvpr.2017.502.

[51] S. Xie, C. Sun, J. Huang, Z. Tu, and K. Murphy, "Rethinking spatiotemporal feature learning: Speed-accuracy trade-offs in video classification," in *Computer Vision – ECCV 2018*, ser. Lecture Notes in Computer Science, Springer International Publishing, 2018, pp. 318–335, doi: 10.1007/978-3-030-01267-0_19.

[52] D. Tran, H. Wang, L. Torresani, J. Ray, Y. LeCun, and M. Paluri, "A closer look at spatiotemporal convolutions for action recognition," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2018, pp. 6450–6459, doi: 10.1109/cvpr.2018.00675.

[53] K. Simonyan and A. Zisserman, "Two-stream convolutional networks for action recognition in videos," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 27, 2014.

[54] C. Feichtenhofer, H. Fan, J. Malik, and K. He, "SlowFast networks for video recognition," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2019, pp. 6201–6210, doi: 10.1109/iccv.2019.00630.

[55] A. Vaswani et al., "Attention is all you need," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.

[56] A. Dosovitskiy et al., "An image is worth 16x16 words: Transformers for image recognition at scale," in *International Conference on Learning Representations (ICLR)*, 2021.

[57] A. Arnab, M. Dehghani, G. Heigold, C. Sun, M. Lucic, and C. Schmid, "ViViT: A video vision transformer," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2021, pp. 6816–6826, doi: 10.1109/iccv48922.2021.00676.

[58] G. Bertasius, H. Wang, and L. Torresani, "Is space-time attention all you need for video understanding?" In *Proceedings of the International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 139, 2021.

[59] Z. Liu et al., "Video swin transformer," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2022, pp. 3192–3201, doi: 10.1109/cvpr52688.2022.00320.

[60] T.-Y. Lin, P. Dollar, R. Girshick, K. He, B. Hariharan, and S. Belongie, "Feature pyramid networks for object detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2017, pp. 936–944, doi: 10.1109/cvpr.2017.106.

[61] R. Ranftl, A. Bochkovskiy, and V. Koltun, "Vision transformers for dense prediction," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2021, pp. 12 159–12 168, doi: 10.1109/iccv48922.2021.01196.

[62] Y. Li, H. Mao, R. Girshick, and K. He, "Exploring plain vision transformer backbones for object detection," in *Computer Vision – ECCV 2022*, ser. Lecture Notes in Computer Science, Springer Nature Switzerland, 2022, pp. 280–296, doi: 10.1007/978-3-031-20077-9_17.

[63] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, "A simple framework for contrastive learning of visual representations," in *Proceedings of the International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 119, 2020.

[64] K. He, H. Fan, Y. Wu, S. Xie, and R. Girshick, "Momentum contrast for unsupervised visual representation learning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2020, pp. 9726–9735, doi: 10.1109/cvpr42600.2020.00975.

[65] J.-B. Grill et al., "Bootstrap your own latent: A new approach to self-supervised learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020.

[66] M. Caron et al., "Emerging properties in self-supervised vision transformers," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2021, pp. 9630–9640, doi: 10.1109/iccv48922.2021.00951.

[67] A. Radford et al., "Learning transferable visual models from natural language supervision," in *Proceedings of the International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 139, 2021.

[68] K. He, X. Chen, S. Xie, Y. Li, P. Dollar, and R. Girshick, "Masked autoencoders are scalable vision learners," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2022, pp. 15 979–15 988, doi: 10.1109/cvpr52688.2022.01553.

[69] Z. Tong, Y. Song, J. Wang, and L. Wang, "VideoMAE: Masked autoencoders are data-efficient learners for self-supervised video pre-training," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022.

[70] K. Li et al., "Unmasked teacher: Towards training-efficient video foundation models," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2023, pp. 19 891–19 903, doi: 10.1109/iccv51070.2023.01826.

[71] Y. Wang et al., "InternVideo2: Scaling foundation models for multimodal video understanding," in *Computer Vision – ECCV 2024*, ser. Lecture Notes in Computer Science, Springer Nature Switzerland, 2025, pp. 396–416, doi: 10.1007/978-3-031-73013-9_23.

[72] M. Oquab et al., "DINOv2: Learning robust visual features without supervision," *Transactions on Machine Learning Research*, 2024.

[73] D. Ha and J. Schmidhuber, "Recurrent world models facilitate policy evolution," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018.

[74] J. Yosinski, J. Clune, Y. Bengio, and H. Lipson, "How transferable are features in deep neural networks?" In *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 27, 2014.

[75] A. Kumar, A. Raghunathan, R. Jones, T. Ma, and P. Liang, "Fine-tuning can distort pretrained features and underperform out-of-distribution," in *International Conference on Learning Representations (ICLR)*, 2022.

[76] N. Houlsby et al., "Parameter-efficient transfer learning for NLP," in *Proceedings of the International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 97, 2019.

[77] E. J. Hu et al., "LoRA: Low-rank adaptation of large language models," in *International Conference on Learning Representations (ICLR)*, 2022.

[78] E. Perez, F. Strub, H. De Vries, V. Dumoulin, and A. Courville, "FiLM: Visual reasoning with a general conditioning layer," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 32, Association for the Advancement of Artificial Intelligence (AAAI), 2018, doi: 10.1609/aaai.v32i1.11671.

[79] B. Lakshminarayanan, A. Pritzel, and C. Blundell, "Simple and scalable predictive uncertainty estimation using deep ensembles," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.

[80] L. Jiang, M. Xu, T. Liu, M. Qiao, and Z. Wang, "DeepVS: A deep learning based video saliency prediction approach," in *Computer Vision – ECCV 2018*, ser. Lecture Notes in Computer Science, Springer International Publishing, 2018, pp. 625–642, doi: 10.1007/978-3-030-01264-9_37.

[81] M. Marszalek, I. Laptev, and C. Schmid, "Actions in context," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2009, pp. 2929–2936, doi: 10.1109/cvpr.2009.5206557.

[82] S. Mathe and C. Sminchisescu, "Actions in the eye: Dynamic gaze datasets and learnt saliency models for visual recognition," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 37, no. 7, pp. 1408–1424, 2015, doi: 10.1109/tpami.2014.2366154.

[83] M. D. Rodriguez, J. Ahmed, and M. Shah, "Action MACH a spatio-temporal maximum average correlation height filter for action recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2008, pp. 1–8, doi: 10.1109/cvpr.2008.4587727.

[84] E. Vig, M. Dorr, and D. Cox, "Space-variant descriptor sampling for action recognition based on saliency and eye movements," in *Computer Vision – ECCV 2012*, ser. Lecture Notes in Computer Science, Springer Berlin Heidelberg, 2012, pp. 84–97, doi: 10.1007/978-3-642-33786-4_7.

[85] A. Coutrot and N. Guyader, "An efficient audiovisual saliency model to predict eye positions when looking at conversations," in *Proceedings of the 23rd European Signal Processing Conference (EUSIPCO)*, IEEE, 2015, pp. 1531–1535, doi: 10.1109/eusipco.2015.7362640.

[86] P. Koutras and P. Maragos, "A perceptually based spatio-temporal computational framework for visual saliency estimation," *Signal Processing: Image Communication*, vol. 38, pp. 15–31, 2015, doi: 10.1016/j.image.2015.08.004.

[87] A. Tsiami, P. Koutras, A. Katsamanis, A. Vatakis, and P. Maragos, "A behaviorally inspired fusion approach for computational audiovisual saliency modeling," *Signal Processing: Image Communication*, vol. 76, pp. 186–200, 2019, doi: 10.1016/j.image.2019.05.001.

[88] X. Min, G. Zhai, K. Gu, and X. Yang, "Fixation prediction through multimodal analysis," *ACM Transactions on Multimedia Computing, Communications, and Applications*, vol. 13, no. 1, pp. 1–23, 2017, doi: 10.1145/2996463.

[89] M. Gygli, H. Grabner, H. Riemenschneider, and L. Van Gool, "Creating summaries from user videos," in *Computer Vision – ECCV 2014*, ser. Lecture Notes in Computer Science, Springer International Publishing, 2014, pp. 505–520, doi: 10.1007/978-3-319-10584-0_33.

[90] A. Tsiami, P. Koutras, and P. Maragos, "STAViS: Spatio-temporal AudioVisual saliency network," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2020, pp. 4765–4775, doi: 10.1109/cvpr42600.2020.00482.

[91] J. Xiong, P. Zhang, T. You, C. Li, W. Huang, and Y. Zha, "DiffSal: Joint audio and video learning for diffusion saliency prediction," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2024, pp. 27 263–27 273, doi: 10.1109/cvpr52733.2024.02575.

[92] R. J. Peters, A. Iyer, L. Itti, and C. Koch, "Components of bottom-up gaze allocation in natural images," *Vision Research*, vol. 45, no. 18, pp. 2397–2416, 2005, doi: 10.1016/j.visres.2005.03.019.

[93] A. Borji, D. N. Sihite, and L. Itti, "Quantitative analysis of human-model agreement in visual saliency modeling: A comparative study," *IEEE Transactions on Image Processing*, vol. 22, no. 1, pp. 55–69, 2013, doi: 10.1109/tip.2012.2210727.

[94] M. Kümmerer, T. S. A. Wallis, and M. Bethge, "Saliency benchmarking made easy: Separating models, maps and metrics," in *Computer Vision – ECCV 2018*, ser. Lecture Notes in Computer Science, Springer International Publishing, 2018, pp. 798–814, doi: 10.1007/978-3-030-01270-0_47.

[95] C. Bak, A. Kocak, E. Erdem, and A. Erdem, "Spatio-temporal saliency networks for dynamic saliency prediction," *IEEE Transactions on Multimedia*, vol. 20, no. 7, pp. 1688–1698, 2018, doi: 10.1109/tmm.2017.2777665.

[96] Q. Lai, W. Wang, H. Sun, and J. Shen, "Video saliency prediction using spatiotemporal residual attentive networks," *IEEE Transactions on Image Processing*, vol. 29, pp. 1113–1126, 2020, doi: 10.1109/tip.2019.2936112.

[97] S. Gorji and J. J. Clark, "Going from image to video saliency: Augmenting image salience with dynamic attentional push," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2018, pp. 7501–7511, doi: 10.1109/cvpr.2018.00783.

[98] P. Linardos, E. Mohedano, J. J. Nieto, N. E. O'Connor, X. Giró-i-Nieto, and K. McGuinness, "Simple vs complex temporal recurrences for video saliency prediction," in *Proceedings of the British Machine Vision Conference (BMVC)*, 2019.

[99] R. Droste, J. Jiao, and J. A. Noble, "Unified image and video saliency modeling," in *Computer Vision – ECCV 2020*, ser. Lecture Notes in Computer Science, Springer International Publishing, 2020, pp. 419–435, doi: 10.1007/978-3-030-58558-7_25.

[100] K. Min and J. Corso, "TASED-Net: Temporally-aggregating spatial encoder-decoder network for video saliency detection," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, 2019, pp. 2394–2403, doi: 10.1109/iccv.2019.00248.

[101] G. Bellitto, F. Proietto Salanitri, S. Palazzo, F. Rundo, D. Giordano, and C. Spampinato, "Hierarchical domain-adapted feature learning for video saliency prediction," *International Journal of Computer Vision*, vol. 129, no. 12, pp. 3216–3232, 2021, doi: 10.1007/s11263-021-01519-y.

[102] S. Jain, P. Yarlagadda, S. Jyoti, S. Karthik, R. Subramanian, and V. Gandhi, "ViNet: Pushing the limits of visual modality for audio-visual saliency prediction," in *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, IEEE, 2021, pp. 3520–3527, doi: 10.1109/iros51168.2021.9635989.

[103] F. Hu et al., "TinyHD: Efficient video saliency prediction with heterogeneous decoders using hierarchical maps distillation," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, IEEE, 2023, pp. 2050–2059, doi: 10.1109/wacv56688.2023.00209.

[104] Y. Zhang, T. Zhang, C. Wu, and R. Tao, "Multi-scale spatiotemporal feature fusion network for video saliency prediction," *IEEE Transactions on Multimedia*, vol. 26, pp. 4183–4193, 2024, doi: 10.1109/tmm.2023.3321394.

[105] Z. Wang et al., "Spatio-temporal self-attention network for video saliency prediction," *IEEE Transactions on Multimedia*, vol. 25, pp. 1161–1174, 2023, doi: 10.1109/tmm.2021.3139743.

[106] C. Ma, H. Sun, Y. Rao, J. Zhou, and J. Lu, "Video saliency forecasting transformer," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 32, no. 10, pp. 6850–6862, 2022, doi: 10.1109/tcsvt.2022.3172971.

[107] X. Zhou et al., "Transformer-based multi-scale feature integration network for video saliency prediction," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 33, no. 12, pp. 7696–7707, 2023, doi: 10.1109/tcsvt.2023.3278410.

[108] M. Moradi, S. Palazzo, and C. Spampinato, "Transformer-based video saliency prediction with high temporal dimension decoding," in *Proceedings of the International Joint Conference on Computer Vision, Imaging and Computer Graphics Theory and Applications (VISAPP)*, SCITEPRESS - Science and Technology Publications, 2024, pp. 616–623, doi: 10.5220/0012422800003660.

[109] C. Li and S. Liu, "TM2SP: A transformer-based multi-level spatiotemporal feature pyramid network for video saliency prediction," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 35, no. 6, pp. 5236–5250, 2025, doi: 10.1109/tcsvt.2025.3529473.

[110] M. Moradi et al., "Learning long- and short-term dynamics for human attention prediction using large video models," *Computer Vision and Image Understanding*, vol. 268, p. 104 740, 2026, doi: 10.1016/j.cviu.2026.104740.

[111] M. Moradi, M. Moradi, C. Spampinato, A. Borji, and S. Palazzo, "Knowledge distillation meets video foundation models: A video saliency prediction case study," *Journal of Visual Communication and Image Representation*, vol. 115, p. 104 706, 2026, doi: 10.1016/j.jvcir.2026.104706.

[112] M. Wu et al., "Learning explainable video saliency prediction from reasoning-annotated eye-tracking supervision," in *Proceedings of the ACM International Conference on Multimedia (ACM MM)*, ACM, 2026, pp. 11 613–11 622, doi: 10.1145/3767308.3835244.

[113] J. Xiong, G. Wang, P. Zhang, W. Huang, Y. Zha, and G. Zhai, "CASP-Net: Rethinking video saliency prediction from an audio-visual consistency perceptual perspective," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2023, pp. 6441–6450, doi: 10.1109/cvpr52729.2023.00623.

[114] D. Zhu et al., "From discrete representation to continuous modeling: A novel audio-visual saliency prediction model with implicit neural representations," *IEEE Transactions on Emerging Topics in Computational Intelligence*, vol. 8, no. 6, pp. 4059–4074, 2024, doi: 10.1109/tetci.2024.3386619.

[115] N. Zhang, M. Xiong, D. Zhu, K. Zhu, G. Zhai, and X. Yang, "Audio-visual saliency prediction model with implicit neural representation," *ACM Transactions on Multimedia Computing, Communications, and Applications*, vol. 21, no. 4, pp. 1–23, 2025, doi: 10.1145/3698881.

[116] L. Yu, X. Sun, P. Gao, and M. Gabbouj, "Relevance-guided audio visual fusion for video saliency prediction," *IEEE Transactions on Emerging Topics in Computational Intelligence*, vol. 10, no. 4, pp. 3079–3091, 2026, doi: 10.1109/tetci.2026.3683092.

[117] R. Agrawal, S. Jyoti, R. Girmaji, S. Sivaprasad, and V. Gandhi, "Does audio help in deep audio-visual saliency prediction models?" In *Proceedings of the International Conference on Multimodal Interaction (ICMI)*, ACM, 2022, pp. 48–56, doi: 10.1145/3536221.3556625.

[118] W. Lotter, G. Kreiman, and D. Cox, "Deep predictive coding networks for video prediction and unsupervised learning," in *International Conference on Learning Representations (ICLR)*, 2017.

[119] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "ImageNet: A large-scale hierarchical image database," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, 2009, pp. 248–255, doi: 10.1109/cvpr.2009.5206848.

[120] Y. Gong, Y.-A. Chung, and J. Glass, "AST: Audio spectrogram transformer," in *Proceedings of Interspeech*, ISCA, 2021, pp. 571–575, doi: 10.21437/interspeech.2021-698.

[121] J. F. Gemmeke et al., "Audio set: An ontology and human-labeled dataset for audio events," in *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, IEEE, 2017, pp. 776–780, doi: 10.1109/icassp.2017.7952261.

[122] Y. Wu and K. He, "Group normalization," in *Computer Vision – ECCV 2018*, ser. Lecture Notes in Computer Science, Springer International Publishing, 2018, pp. 3–19, doi: 10.1007/978-3-030-01261-8_1.

[123] C.-Y. Lee, S. Xie, P. Gallagher, Z. Zhang, and Z. Tu, "Deeply-supervised nets," in *Proceedings of the International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 38, 2015.

[124] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *International Conference on Learning Representations (ICLR)*, 2019.

[125] I. Loshchilov and F. Hutter, "SGDR: Stochastic gradient descent with warm restarts," in *International Conference on Learning Representations (ICLR)*, 2017.

[126] B. Efron, "Bootstrap methods: Another look at the jackknife," *The Annals of Statistics*, vol. 7, no. 1, pp. 1–26, 1979, doi: 10.1214/aos/1176344552.

[127] C. R. Harris et al., "Array programming with NumPy," *Nature*, vol. 585, no. 7825, pp. 357–362, 2020, doi: 10.1038/s41586-020-2649-2.

[128] L. Yu, W. He, J. Xiao, and H. Bai, "FOSal: Fine-grained semantic auditory guidance and omni-scale visual perception for audio-visual saliency prediction," *IEEE Transactions on Multimedia*, pp. 1–12, 2026, doi: 10.1109/tmm.2026.3729395.

[129] A. Mao, J. Yan, Y. Fang, and C. Cai, "Audio-visual saliency prediction based on joint adversarial learning and co-attention mechanism," *Pattern Recognition*, vol. 179, p. 113 548, 2026, doi: 10.1016/j.patcog.2026.113548.

APPENDIX A

ADDITIONAL RESULTS

This appendix gives the complete statistics behind the summaries of Chapter 5. All values were produced by the evaluator of Section 4.8.

A.1. Paired Bootstrap on the DHF1K Validation Set

Tables A.1 and A.2 extend Table 5.6 to all six metrics. Each cell gives the mean per-video difference and, below it, its 95 % percentile interval over 10,000 bootstrap resamples of the 100 validation videos. For KLD, lower is better, so negative differences favour the first model.

Table A.1: Paired bootstrap differences between SURGE-D and SalFoM on DHF1K validation, all six metrics

Metric	ViT-L single	ViT-L ensemble	ViT-g single	ViT-g ensemble
CC	+0.0197 [+0.0095, +0.0305]	+0.0271 [+0.0171, +0.0376]	+0.0364 [+0.0259, +0.0475]	+0.0390 [+0.0286, +0.0498]
NSS	+0.0225 [−0.0499, +0.0967]	+0.0910 [+0.0218, +0.1637]	+0.1304 [+0.0603, +0.2047]	+0.1597 [+0.0891, +0.2333]
SIM	+0.0362 [+0.0291, +0.0436]	+0.0356 [+0.0285, +0.0430]	+0.0401 [+0.0328, +0.0476]	+0.0444 [+0.0371, +0.0520]
KLD	−0.0747 [−0.1021, −0.0486]	−0.1101 [−0.1359, −0.0858]	−0.1270 [−0.1527, −0.1022]	−0.1369 [−0.1634, −0.1119]
AUC-J	+0.0040 [+0.0022, +0.0060]	+0.0060 [+0.0042, +0.0079]	+0.0065 [+0.0048, +0.0084]	+0.0071 [+0.0053, +0.0090]
sAUC	+0.0230 [+0.0166, +0.0298]	+0.0208 [+0.0143, +0.0277]	+0.0304 [+0.0238, +0.0373]	+0.0275 [+0.0209, +0.0345]

A.2. Checks Added After the First Draft

The checks of this section were specified in the research log (section 12) before they were run (Table 4.1, rules marked †). Table A.3 repeats the comparison of Table A.1 against SalFoM with the blur that is best for SalFoM on the validation set ($\sigma = 16$), an upper bound for the baseline. Table A.4 lists every use of the DHF1K validation set, the test set of this thesis. Tables A.5 and A.6 give all six metrics for the causal models of Section 5.5.2, against SalFoM and against the offline models. Table A.7 gives all six metrics for the models trained on 600 instead of 540 DHF1K videos, and Tables A.9 and A.10 the intervals of gate G2 with the untrained and the trained mask token. Table A.11 gives, for comparison with Table 5.4, the information that each cue

Table A.2: Paired bootstrap differences between the ViT-g and ViT-L encoders on DHF1K validation, all six metrics

Metric	Single models	Three-seed ensembles
CC	+0.0167 [+0.0116, +0.0222]	+0.0118 [+0.0077, +0.0163]
NSS	+0.1079 [+0.0746, +0.1423]	+0.0687 [+0.0407, +0.0974]
SIM	+0.0039 [+0.0000, +0.0077]	+0.0088 [+0.0058, +0.0119]
KLD	−0.0523 [−0.0669, −0.0386]	−0.0268 [−0.0381, −0.0158]
AUC-J	+0.0025 [+0.0016, +0.0035]	+0.0011 [+0.0004, +0.0020]
sAUC	+0.0074 [+0.0041, +0.0107]	+0.0067 [+0.0038, +0.0095]

adds on its own to the centre prior. Table A.8, a descriptive comparison that was added afterwards and computed from the existing validation scores, compares each three-seed ensemble with its seed-0 member.

Table A.3: Paired bootstrap differences between SURGE-D and SalFoM with its best blur ($\sigma = 16$) on DHF1K validation, all six metrics (100 videos, 10,000 resamples)

Metric	ViT-L single	ViT-L ensemble	ViT-g single	ViT-g ensemble
CC	+0.0095 [−0.0008, +0.0203]	+0.0169 [+0.0069, +0.0273]	+0.0262 [+0.0160, +0.0371]	+0.0287 [+0.0186, +0.0394]
NSS	+0.0842 [+0.0149, +0.1574]	+0.1528 [+0.0839, +0.2253]	+0.1921 [+0.1252, +0.2636]	+0.2214 [+0.1534, +0.2938]
SIM	+0.0483 [+0.0406, +0.0563]	+0.0477 [+0.0399, +0.0558]	+0.0522 [+0.0445, +0.0601]	+0.0565 [+0.0485, +0.0645]
KLD	−0.0577 [−0.0860, −0.0304]	−0.0932 [−0.1191, −0.0683]	−0.1100 [−0.1357, −0.0855]	−0.1199 [−0.1461, −0.0949]
AUC-J	+0.0006 [−0.0011, +0.0024]	+0.0025 [+0.0009, +0.0042]	+0.0031 [+0.0015, +0.0048]	+0.0037 [+0.0020, +0.0054]
sAUC	+0.0078 [+0.0015, +0.0147]	+0.0056 [−0.0008, +0.0125]	+0.0152 [+0.0087, +0.0223]	+0.0123 [+0.0058, +0.0196]

A.3. All Metrics on the Audio-Visual Datasets

Table A.12 reports all six metrics of the SURGE-D ViT-g three-seed ensemble on the six audio-visual datasets, as fold means. The output blur was tuned on each fold's holdout; the selected values lay between 8 and 16 pixels.

Table A.4: Every use of the DHF1K validation set, the test set of this thesis. No reported model or setting was selected on it

Model	**Purpose**	**Decision informed**
SalFoM, released	evaluator validation (Section 5.2); baseline	none
SalFoM, blur grid (σ from 0 to 24)	upper bound for the baseline	none
SURGE-D ViT-L, seed 0 and ensemble	RQ1	which maps were first prepared for the official test set
SURGE-D ViT-g, seed 0 and ensemble	RQ1; RQ3 (scale)	which maps were prepared for the official test set, replacing the ViT-L maps
SURGE-D ViT-g on 600 videos, seed 0 and ensemble	RQ3 (amount of data)	whether these models replace the 540-video models for the official test set (they did not)
SURGE-D, seeds 1 and 2 of both encoders	seed check	none
Frozen UMT-L with the SURGE-D decoder	encoder check	none
Causal SURGE-D ViT-g and ViT-L, three seeds and ensemble each (every tenth frame)	RQ1 (causal control)	none
SalFoM, released and best blur, and offline SURGE-D ViT-g and ViT-L, on the same frames	causal control	none

Table A.5: Paired bootstrap differences between causal SURGE-D and SalFoM with its best blur ($\sigma = 16$) on DHF1K validation, every tenth frame, all six metrics (100 videos, 10,000 resamples)

Metric	Seed 0	Seed 1	Seed 2	Ensemble
Causal SURGE-D ViT-g				
CC	+0.0249 [+0.0142, +0.0362]	+0.0184 [+0.0080, +0.0294]	+0.0196 [+0.0082, +0.0314]	+0.0292 [+0.0188, +0.0400]
NSS	+0.1807 [+0.1101, +0.2540]	+0.1640 [+0.0966, +0.2352]	+0.1700 [+0.0984, +0.2462]	+0.2177 [+0.1506, +0.2894]
SIM	+0.0560 [+0.0478, +0.0642]	+0.0427 [+0.0349, +0.0505]	+0.0576 [+0.0489, +0.0664]	+0.0558 [+0.0479, +0.0638]
KLD	−0.1132 [−0.1397, −0.0867]	−0.0998 [−0.1253, −0.0750]	−0.0886 [−0.1190, −0.0580]	−0.1275 [−0.1532, −0.1023]
AUC-J	+0.0036 [+0.0019, +0.0055]	+0.0036 [+0.0019, +0.0053]	+0.0025 [+0.0006, +0.0046]	+0.0047 [+0.0031, +0.0065]
sAUC	+0.0096 [+0.0032, +0.0164]	+0.0108 [+0.0038, +0.0185]	+0.0032 [−0.0050, +0.0115]	+0.0116 [+0.0048, +0.0189]
Causal SURGE-D ViT-L				
CC	+0.0113 [+0.0009, +0.0221]	+0.0038 [−0.0062, +0.0142]	+0.0063 [−0.0049, +0.0180]	+0.0171 [+0.0071, +0.0275]
NSS	+0.1022 [+0.0313, +0.1751]	+0.0680 [+0.0022, +0.1374]	+0.1022 [+0.0300, +0.1776]	+0.1464 [+0.0793, +0.2172]
SIM	+0.0502 [+0.0420, +0.0586]	+0.0324 [+0.0249, +0.0402]	+0.0497 [+0.0401, +0.0594]	+0.0486 [+0.0406, +0.0568]
KLD	−0.0712 [−0.1000, −0.0424]	−0.0606 [−0.0861, −0.0357]	−0.0481 [−0.0808, −0.0156]	−0.0942 [−0.1204, −0.0685]
AUC-J	+0.0014 [−0.0005, +0.0034]	+0.0011 [−0.0005, +0.0028]	+0.0000 [−0.0019, +0.0021]	+0.0028 [+0.0011, +0.0045]
sAUC	+0.0059 [−0.0003, +0.0126]	+0.0021 [−0.0047, +0.0094]	−0.0014 [−0.0091, +0.0067]	+0.0066 [+0.0002, +0.0135]

Table A.6: Paired bootstrap differences between causal and offline SURGE-D on the same DHF1K validation frames (every tenth frame), all six metrics (100 videos, 10,000 resamples). For KLD, negative differences favour the causal model

	ViT-g		ViT-L	
Metric	Seed 0	Ensemble	Seed 0	Ensemble
CC	+0.0003 [−0.0044, +0.0049]	+0.0016 [−0.0015, +0.0046]	+0.0030 [−0.0014, +0.0077]	+0.0012 [−0.0022, +0.0047]
NSS	−0.0022 [−0.0318, +0.0272]	+0.0042 [−0.0155, +0.0226]	+0.0249 [−0.0041, +0.0560]	+0.0007 [−0.0203, +0.0225]
SIM	+0.0043 [+0.0011, +0.0076]	−0.0002 [−0.0025, +0.0020]	+0.0025 [−0.0007, +0.0060]	+0.0014 [−0.0009, +0.0039]
KLD	−0.0074 [−0.0192, +0.0051]	−0.0099 [−0.0179, −0.0017]	−0.0172 [−0.0293, −0.0052]	−0.0025 [−0.0120, +0.0069]
AUC-J	+0.0006 [−0.0003, +0.0015]	+0.0010 [+0.0004, +0.0016]	+0.0007 [−0.0001, +0.0016]	+0.0001 [−0.0006, +0.0008]
sAUC	−0.0049 [−0.0084, −0.0014]	+0.0000 [−0.0023, +0.0023]	−0.0011 [−0.0041, +0.0019]	+0.0014 [−0.0009, +0.0040]

Table A.7: Training on all 600 DHF1K training videos instead of 540 (ViT-g): paired bootstrap differences on DHF1K validation, all six metrics (100 videos, 10,000 resamples)

Metric	Seed 0	Three-seed ensemble
CC	−0.0001 [−0.0035, +0.0034]	+0.0015 [−0.0006, +0.0037]
NSS	+0.0386 [+0.0131, +0.0656]	+0.0140 [−0.0002, +0.0284]
SIM	+0.0063 [+0.0038, +0.0089]	−0.0017 [−0.0032, −0.0001]
KLD	+0.0031 [−0.0075, +0.0136]	−0.0039 [−0.0101, +0.0026]
AUC-J	−0.0001 [−0.0007, +0.0006]	+0.0001 [−0.0005, +0.0006]
sAUC	−0.0062 [−0.0089, −0.0034]	+0.0017 [−0.0000, +0.0033]

Table A.8: Three-seed ensemble minus its seed-0 member on DHF1K validation: paired bootstrap differences, all six metrics (100 videos, 10,000 resamples). For KLD, negative differences favour the ensemble

Metric	ViT-L	ViT-g
CC	+0.0074 [+0.0050, +0.0099]	+0.0025 [+0.0002, +0.0049]
NSS	+0.0686 [+0.0514, +0.0871]	+0.0293 [+0.0141, +0.0454]
SIM	−0.0006 [−0.0024, +0.0013]	+0.0043 [+0.0025, +0.0063]
KLD	−0.0354 [−0.0440, −0.0273]	−0.0099 [−0.0164, −0.0035]
AUC-J	+0.0019 [+0.0014, +0.0025]	+0.0006 [+0.0002, +0.0010]
sAUC	−0.0022 [−0.0040, −0.0005]	−0.0029 [−0.0046, −0.0012]

Table A.9: Gate G2 with the untrained mask token, as in its first run, with 95 % bootstrap intervals over the 60 holdout videos, computed from the stored per-video scores. The masked-tube error was added in the second follow-up. The reference maps do not depend on the mask token. The lower block gives paired differences

	NSS		CC	
Map	4 fps	15 fps	4 fps	15 fps
Raw surprise	−0.190 [−0.332, −0.045]	−0.502 [−0.620, −0.385]	−0.043 [−0.069, −0.015]	−0.099 [−0.122, −0.077]
Contextual surprise	+0.116 [−0.043, +0.275]	−0.366 [−0.505, −0.225]	+0.018 [−0.011, +0.047]	−0.072 [−0.098, −0.046]
Masked-tube error	−0.858 [−0.946, −0.767]	−0.819 [−0.912, −0.721]	−0.170 [−0.186, −0.153]	−0.161 [−0.178, −0.143]
Token atypicality	+0.454 [+0.315, +0.599]	+0.650 [+0.532, +0.770]	+0.081 [+0.054, +0.108]	+0.123 [+0.101, +0.145]
Frame difference	+0.810 [+0.677, +0.954]	+0.803 [+0.678, +0.938]	+0.149 [+0.126, +0.174]	+0.148 [+0.126, +0.172]
Empirical centre prior	+1.464 [+1.270, +1.664]	+1.479 [+1.289, +1.676]	+0.289 [+0.253, +0.325]	+0.292 [+0.257, +0.327]
Raw surprise − frame difference	−1.000 [−1.183, −0.817]	−1.305 [−1.478, −1.137]	−0.192 [−0.225, −0.158]	−0.248 [−0.279, −0.217]
Raw surprise − atypicality	−0.644 [−0.768, −0.523]	−1.152 [−1.278, −1.028]	−0.124 [−0.147, −0.100]	−0.222 [−0.246, −0.200]
Contextual − frame difference	−0.694 [−0.908, −0.486]	−1.169 [−1.382, −0.965]	−0.131 [−0.170, −0.093]	−0.220 [−0.259, −0.183]
Contextual − atypicality	−0.338 [−0.514, −0.166]	−1.015 [−1.193, −0.837]	−0.063 [−0.097, −0.030]	−0.195 [−0.228, −0.162]
Masked-tube − frame difference	−1.667 [−1.854, −1.487]	−1.622 [−1.800, −1.445]	−0.319 [−0.350, −0.288]	−0.310 [−0.340, −0.279]
Masked-tube − atypicality	−1.312 [−1.465, −1.163]	−1.468 [−1.609, −1.335]	−0.250 [−0.278, −0.224]	−0.284 [−0.309, −0.261]

Table A.10: Gate G2 re-scored with the trained mask token, with 95 % bootstrap intervals over the 60 holdout videos. The reference maps are those of Table A.9. The lower block gives paired differences

	NSS		CC	
Map	4 fps	15 fps	4 fps	15 fps
Raw surprise	−0.665 [−0.763, −0.566]	−0.877 [−0.958, −0.796]	−0.129 [−0.146, −0.110]	−0.168 [−0.182, −0.154]
Contextual surprise	−0.239 [−0.373, −0.107]	−0.660 [−0.770, −0.548]	−0.045 [−0.070, −0.020]	−0.127 [−0.147, −0.106]
Masked-tube error	−0.958 [−1.046, −0.873]	−0.942 [−1.026, −0.858]	−0.186 [−0.201, −0.172]	−0.182 [−0.196, −0.167]
Raw surprise − frame difference	−1.475 [−1.675, −1.283]	−1.680 [−1.855, −1.507]	−0.278 [−0.312, −0.244]	−0.316 [−0.347, −0.286]
Raw surprise − atypicality	−1.119 [−1.285, −0.954]	−1.526 [−1.679, −1.374]	−0.209 [−0.239, −0.180]	−0.291 [−0.318, −0.264]
Contextual − frame difference	−1.049 [−1.285, −0.828]	−1.463 [−1.670, −1.268]	−0.194 [−0.236, −0.154]	−0.275 [−0.312, −0.240]
Contextual − atypicality	−0.693 [−0.899, −0.495]	−1.310 [−1.494, −1.123]	−0.126 [−0.164, −0.088]	−0.250 [−0.283, −0.215]
Masked-tube − frame difference	−1.768 [−1.967, −1.580]	−1.745 [−1.928, −1.567]	−0.335 [−0.368, −0.304]	−0.330 [−0.361, −0.300]
Masked-tube − atypicality	−1.412 [−1.590, −1.234]	−1.591 [−1.751, −1.434]	−0.267 [−0.299, −0.235]	−0.305 [−0.333, −0.278]

Table A.11: Information gain of each cue on its own over the fitted centre prior, in bits per fixation, cross-fitted as in Equation (4.11), with 95 % bootstrap intervals over the 60 holdout videos (descriptive). A positive gain can come from a negative weight. Maps without a token named are computed with the untrained token

Cue	4 fps	15 fps
Frame difference	+0.1925 [+0.1248, +0.2713]	+0.1872 [+0.1232, +0.2590]
Token atypicality	+0.1765 [+0.0933, +0.2633]	+0.1749 [+0.0981, +0.2559]
Raw surprise	+0.0360 [+0.0012, +0.0720]	−0.0010 [−0.0055, +0.0032]
Contextual surprise	+0.0247 [−0.0123, +0.0615]	+0.0023 [−0.0068, +0.0113]
Masked-tube error	+0.0191 [−0.0076, +0.0466]	+0.0246 [−0.0040, +0.0534]
Raw surprise, trained token	+0.0293 [+0.0020, +0.0579]	+0.0946 [+0.0540, +0.1367]
Contextual surprise, trained token	+0.0048 [−0.0075, +0.0179]	+0.0654 [+0.0276, +0.1064]
Masked-tube error, trained token	+0.0891 [+0.0421, +0.1382]	+0.1056 [+0.0609, +0.1510]

Table A.12: SURGE-D ViT-g three-seed ensemble on the audio-visual datasets (fold means; DIEM has one split)

Dataset	CC↑	NSS↑	SIM↑	KLD↓	AUC-J↑	sAUC↑
DIEM	0.729	2.878	0.584	0.664	0.923	0.735
Coutrot1	0.657	3.302	0.504	0.947	0.911	0.774
Coutrot2	0.819	6.673	0.601	0.683	0.966	0.759
ETMD	0.676	3.687	0.510	0.930	0.950	0.765
AVAD	0.740	4.306	0.565	0.795	0.942	0.782
SumMe	0.590	3.261	0.448	1.148	0.927	0.763

A.4. Encoder Comparison on the Audio-Visual Datasets

Table A.13 gives the paired bootstrap comparison of the ViT-g and ViT-L ensembles on the audio-visual datasets. Because every video of these datasets is tested in exactly one fold, the test videos of all folds were pooled for the bootstrap.

Table A.13: CC difference between the ViT-g and ViT-L three-seed ensembles on the audio-visual datasets (pooled test videos, 10,000 resamples)

Dataset	Test videos	ΔCC [95 % CI]	Reading
DIEM	17	+0.014 [+0.003, +0.027]	ViT-g better
Coutrot1	60	+0.017 [+0.005, +0.031]	ViT-g better
Coutrot2	15	+0.012 [−0.007, +0.032]	no significant difference
ETMD	12	+0.006 [+0.002, +0.010]	ViT-g better
AVAD	45	+0.011 [+0.005, +0.019]	ViT-g better
SumMe	23	−0.003 [−0.010, +0.004]	no significant difference

APPENDIX B

REPRODUCIBILITY AND SAMPLE CODE

The implementation was written for this thesis and does not import code from other saliency projects. This appendix lists its main parts, the points at which the work departed from its own procedure, the commands that produce the main results, and the forward pass of the decoder.

B.1. Organisation of the Code

Table B.1 lists the main files. The research log records every decision rule with its date, written before the experiment it governs, and every result with the folder it was read from.

B.2. Corrections and Deviations

The procedure of Section 4.1 requires every rule to be written before its experiment and to be applied unchanged. This section collects the points at which the work departed from that procedure, or at which an error in an earlier step was corrected, so that they need not be searched for in the log.

The mask token of the predictor. Gate G2 was first run with the mask token that the model-loading library selects by default, a zero vector that pre-training never updated (Section 4.5.1). The error was found in the second follow-up, while the masked-tube readout was being prepared. Before any score of the corrected maps against gaze data existed, the log recorded that the trained token is the primary readout and what the thesis would report for either outcome of the re-scored gate. The re-scored gate failed as its first run had (Section 5.4), so the decision not to train the surprise module stands. The numbers of the first run are kept in Tables A.9 and 5.3 as its record. The diagnostic on eleven clips without gaze data (Section 5.4) used the untrained token and was not repeated. SURGE-D uses only the encoder, so no model, score or comparison of SURGE-D is affected.

A test of the analysis programs. After those decisions had been recorded, the programs that compute the prediction-error maps and their incremental information were tested end to end on four of the 60 holdout videos. The test printed scores of these four videos against gaze data, including scores with the trained token, and they

Table B.1: Main files of the implementation

File	Content
`surge/backbones.py`	frozen encoders (V-JEPA 2 ViT-L and ViT-g, V-JEPA 2.1, UMT-L) and their taps; predictor error for gate G2
`surge/model.py`	feature pyramid, (2+1)D decoder, auxiliary heads, centre prior, surprise and audio modules
`surge/data.py`	dataset readers, clip sampling, targets, inference windows
`surge/losses.py`	training loss (Equation (4.17))
`surge/metrics.py`	CC, NSS, SIM, KLD, AUC-J and sAUC
`surge/infer.py`	whole-video inference (Algorithm 4.1) and scoring
`surge/lora.py`	low-rank adaptation (Equation (4.12))
`scripts/train.py`	training with holdout selection and clean stop/resume
`scripts/evaluate.py`	scoring of models, ensembles or saved maps
`scripts/tune_sigma.py`	selection of the output blur on the holdout
`scripts/compare.py`	paired bootstrap over videos
`scripts/g2_surprise.py`	gate G2
`scripts/prep_hollywood_gaze.py`	Hollywood-2 (Vig gaze) maps
`scripts/prep_stavis_audio.py`	audio features for the audio experiment
`scripts/thesis_qualitative.py`	qualitative examples and cost measurements
`scripts/g2_intervals.py`	bootstrap intervals of gate G2
`scripts/ablation_seeds.py`	component ablations with three seeds
`scripts/review_tables.py`	numbers of the checks added after the first draft
`scripts/g2_maps.py`	prediction-error maps with both mask tokens, the masked-tube readout and the checks of the visible-token encoder (Section 4.5.1)
`scripts/g2_incremental.py`	incremental information (Equation (4.11)) and the re-scored gate G2
`scripts/make_lc_splits.py`, `scripts/learning_curve.py`	training subsets and analysis of the learning curve (Section 5.6)
`scripts/time_inference.py`	measured inference time of SalFoM and SURGE-D (Section 5.5.5)
`scripts/count_umt_flops.py`	operation count of SalFoM's encoder (Section 5.5.5)
`scripts/queue_review_a.sh to _e.sh`	the runs of the checks added after the first draft, in order
`configs/`	configurations and all split and holdout lists
`research/log.md`	dated decision rules and results

were seen before the analysis of all 60 videos had finished. No rule, margin, primary comparison or line of the analysis programs was changed afterwards, and the four videos are part of the full analysis.

Rules and analyses added after results. The tie margin of 0.01 for comparisons with published numbers was fixed after the first such comparison had been seen (Section 4.1). The bootstrap intervals of the first run of gate G2 were specified and computed after its outcome was known. The comparison of each ensemble with its seed-0 member (Table A.8) was added after the scores that it uses. None of the three decided a model or a setting.

The operation count of SalFoM's encoder. The first draft gave 1.90 TFLOPs for SalFoM's encoder. That count omitted the products with the attention matrices, which the counts of SURGE-D include; with them, it is 2.87 TFLOPs (Section 5.5.5). The correction raises SalFoM's encoder computation per output frame from about 17 to about 25 times that of SURGE-D with ViT-L and from about 6 to about 9 times that with ViT-g, and it makes a causal ViT-L model cheaper than SalFoM's encoder rather than comparable to it.

The measured times. The protocol of the timing required that nothing else be started on the machine while it ran. Short commands were started during the first run, and two brief file searches during the second; neither used the GPU. The measurement was therefore made a third time under the same rule, which it met, and the third run is reported (Table 5.13). The criterion was applied to all three runs, and their medians differ by at most 5.1 %.

Uses of the validation set. The DHF1K validation set is the test set of this thesis. Table A.4 lists every use of it, including the two decisions about the test-set submission that it informed.

B.3. Commands for the Main Results

The main DHF1K model with the ViT-g encoder (seed 0) is trained, its blur is tuned on the holdout, and it is scored once on the validation set with:

```
python scripts/train.py --config configs/main_b1.yaml --set \
    run_dir=runs/vitg_s0 train.seed=0 train.epochs=3 \
    train.warmup_epochs=0.25 model.backbone=vjepa2_g "data.size=[256,256]"
python scripts/tune_sigma.py --run runs/vitg_s0 \
    --out runs/vitg/sigma_single.json
```

```
python scripts/evaluate.py --run runs/vitg_s0 --split val --metrics all \
    --sigma 8 --out runs/vitg/val_single_s0.json
```

Passing `--run` three times (seeds 0, 1 and 2) to the last two commands evaluates the three-seed ensemble. The paired bootstrap against another evaluation file is computed with `python scripts/compare.py A.json B.json`.

The causal variant (Section 5.5.2) is trained and scored on every tenth frame with:

```
python scripts/train.py --config configs/main_b1.yaml --set \
    run_dir=runs/causal_g_s0 train.seed=0 train.epochs=3 \
    train.warmup_epochs=0.25 model.backbone=vjepa2_g \
    "data.size=[256,256]" data.target=last eval.frame_step=10
python scripts/evaluate.py --run runs/causal_g_s0 --split val \
    --metrics all --sigma 8 --frame-step 10 \
    --out runs/review/causal_g_val_s0.json
```

`scripts/evaluate.py` and `scripts/tune_sigma.py` run a model causally whenever its configuration names this target. The ViT-L variant is trained with the same command, in which `vjepa2_l` replaces `vjepa2_g`. For the comparison of Section 5.5.1, SalFoM's saved maps are blurred by passing `--sigma` together with `--maps`.

The analyses of the second follow-up (the learning curve of Section 5.6, the prediction-error maps and their incremental information in Section 5.4, and the measured times of Section 5.5.5) are run with:

```
bash scripts/queue_review_d.sh      # six training runs and their analysis
python scripts/g2_maps.py --out runs/g2/maps
python scripts/g2_incremental.py --out runs/review/g2_incremental.json
python -I scripts/time_inference.py --out runs/review/timing_run3.json
```

The last command needs SalFoM's released code and weights, which are not part of the repository of the thesis.

B.4. Forward Pass of the Decoder

The listing below is the forward method of the decoder in `surge/model.py`, which implements Equations (4.3), (4.5) and (4.7) (the softmax over pixels is applied in the loss and at inference). The two branches used only by the audio and surprise modules are abbreviated.

```
def forward(self, taps, surprise=None, T_out=None, audio=None):
    p4, p8, p16a, p16b = self.fpn(taps)            # SimpleFPN-3D
    if self.audio:                                  # audio experiment only
        ...
```

```
if self.use_scm:                                # surprise module (not used)
    ...
d16 = self.r16(self.fuse16(torch.cat([p16a, p16b], 1)))
d8 = self.r8(p8 + F.interpolate(d16, size=p8.shape[2:], mode="nearest"))
d4 = self.r4(p4 + F.interpolate(d8, size=p4.shape[2:], mode="nearest"))
out = {"logits": self.head(d4)[:, 0]}           # B, T', 64, 64
if self.aux and self.training:
    out["aux16"] = self.aux16(d16)[:, 0]
    out["aux8"] = self.aux8(d8)[:, 0]
for k in list(out):
    x = out[k]
    if k != "logits":                           # aux maps to 64 x 64
        x = F.interpolate(x[:, None], size=out["logits"].shape[1:],
                          mode="trilinear", align_corners=False)[:, 0]
    if T_out is not None and x.shape[1] != T_out:   # tubelets -> frames
        x = F.interpolate(x[:, None], size=(T_out, *x.shape[2:]),
                          mode="trilinear", align_corners=False)[:, 0]
    if self.prior is not None:                  # + log G_theta
        x = x + self.prior(x.shape[-2], x.shape[-1], x.device)
    out[k] = x
return out
```